



PhD-FSTC-2015-39
Faculté des Sciences, de la Technologie et de la Communication

DISSERTATION

Defense held on 15/09/2015 in Luxembourg
to obtain the degree of

DOCTEUR DE L'UNIVERSITÉ DU LUXEMBOURG EN INFORMATIQUE

by

Alejandro CORREA BAHNSEN
Born on March 18, 1985 in Bogota, Colombia

EXAMPLE-DEPENDENT COST-SENSITIVE CLASSIFICATION

APPLICATIONS IN FINANCIAL RISK MODELING
AND MARKETING ANALYTICS

DISSERTATION DEFENSE COMMITTEE

Dr BjÖRN OTTERSTEN, dissertation supervisor
Professor, Université du Luxembourg

Dr DJAMILA AOUADA, dissertation co-supervisor
Research Scientist, Université du Luxembourg

Dr YVES LE TRAON, Chairman
Professor, Université du Luxembourg

Dr BART DE MOOR, Vice Chairman
Professor, Katholieke Universiteit Leuven

Dr GIANLUCA BONTEMPI
Professor, Université Libre de Bruxelles

ABSTRACT

Several real-world binary classification problems are example-dependent cost-sensitive in nature, where the costs due to misclassification vary between examples and not only within classes. However, standard binary classification methods do not take these costs into account, and assume a constant cost of misclassification errors. This approach is not realistic in many real-world applications. For example in credit card fraud detection, failing to detect a fraudulent transaction may have an economical impact from a few to thousands of Euros, depending on the particular transaction and card holder. In churn modeling, a model is used for predicting which customers are more likely to abandon a service provider. In this context, failing to identify a profitable or unprofitable churning has a significant different economic result. Similarly, in direct marketing, wrongly predicting that a customer will not accept an offer when in fact he will, may have different financial impact, as not all customers generate the same profit. Lastly, in credit scoring, accepting loans from bad customers does not have the same economical loss, since customers have different credit lines, therefore, different profit.

Accordingly, the goal of this thesis is to provide an in-depth analysis of example-dependent cost-sensitive classification. We analyze four real-world classification problems, namely, credit card fraud detection, credit scoring, churn modeling and direct marketing. For each problem, we propose an example-dependent cost-sensitive evaluation measure.

We propose four example-dependent cost-sensitive methods; the first method is a cost-sensitive Bayes minimum risk classifier which consists in quantifying tradeoffs between various decisions using probabilities and the costs that accompany such decisions. Second, we propose a cost-sensitive logistic regression technique. This algorithm is based on a new logistic regression cost function; one that takes into account the real costs due to misclassification and correct classification. Subsequently, we propose a cost-sensitive decision trees algorithm which is based on incorporating the different example-dependent costs into a new cost-based impurity measure and a new cost-based pruning criteria. Lastly, we define an example-dependent cost-sensitive framework for ensembles of decision-trees. It is based on training example-dependent cost-sensitive decision trees using four different random inducer methods and then blending them using three different combination approaches. Moreover, we present the library *CostCla* developed as part of the the-

sis. This library is an open-source implementation of all the algorithms covered in this manuscript.

Finally, the experimental results show the importance of using the real example-dependent financial costs associated with real-world applications. We found that there are significant differences in the results when evaluating a model using a traditional cost-insensitive measure such as accuracy or F1Score, than when using the financial savings. Moreover, the results show that the proposed algorithms have better results for all databases, in the sense of higher savings.

To my wife Alejandra
it's a privilege to share my life with you.

ACKNOWLEDGMENTS

First and foremost, I would like to sincerely thank my advisor Prof. Dr. Björn Ottersten for the continuous support during my Ph.D study and research, for his motivation, enthusiasm, and for the freedom I was granted throughout these years.

My sincere thanks also goes to my co-advisor Dr. Djamila Aouada for her support, guidance, and encouragement not only from the research point of view, but for life in general.

Deepest gratitude are also due to all members of the jury for their interest in this work and for taking the time to evaluate this dissertation.

Special thanks go to the fraud management team at CETREL a SIX Company. In particular to Guy Weber and Philippe Davin, for all the time they invest in helping me understand all the particularities of the credit card fraud detection problem. Moreover, for their insights during the different steps of my research.

I would also like to convey thanks for supporting this Ph.D project to the AFR Grant Scheme (Aides la Formation-Recherche), managed by the National Research Fund of Luxembourg (FNR).

I cannot finish without saying how grateful I am with my wife Alejandra and son Pablo. They have always supported and encouraged me to do my best in all matters of life.

CONTENTS

NOTATIONS	xiii
ACRONYMS	xv
LIST OF TABLES	xvi
LIST OF FIGURES	xix
1 INTRODUCTION	1
1.1 Motivation and scope	1
1.2 Contributions	3
1.3 Outline	4
1.4 Publications	5
I EXAMPLE-DEPENDENT COST-SENSITIVE CLASSIFICATION	7
2 BASICS ON CLASSIFICATION	9
2.1 Introduction	9
2.2 Traditional evaluation measures	11
2.2.1 Brier score	14
3 COST-SENSITIVE CLASSIFICATION	15
3.1 Introduction	15
3.2 Class-dependent cost-sensitive classification	16
3.3 Example-dependent cost-sensitive classification	17
3.3.1 Example-dependent evaluation measures	18
3.3.2 Binary classification cost characteristic	21
3.3.3 State-of-the-art methods	22
II REAL-WORLD EXAMPLE-DEPENDENT APPLICATIONS	23
4 FINANCIAL RISK MODELING	25
4.1 Credit card fraud detection	25
4.1.1 Credit card fraud detection evaluation	27
4.1.2 Feature engineering for fraud detection	28
4.1.3 Database	35
4.2 Credit scoring	35
4.2.1 Financial evaluation of a credit scorecard	36
4.2.2 Databases	39
4.3 Summary of the datasets	40
5 MARKETING ANALYTICS	43
5.1 Churn modeling	43
5.1.1 Flow analysis of a churn campaign	44
5.1.2 Propose evaluation measure of a churn campaign	46

5.1.3	Churn modeling database	49
5.2	Direct marketing	50
5.3	Summary of the datasets	51
III PROPOSED EXAMPLE-DEPENDENT COST-SENSITIVE METHODS		53
6	BAYES MINIMUM RISK	55
6.1	Bayes minimum risk model	55
6.2	Calibration of probabilities	56
6.2.1	Calibration due to a change in base rates	56
6.2.2	Calibration using the ROC convex hull	56
6.3	Experiments	58
7	COST-SENSITIVE LOGISTIC REGRESSION	63
7.1	Logistic regression	63
7.2	Cost-sensitive logistic regression	64
7.2.1	Implicit costs of the logistic regression	64
7.2.2	Cost-sensitive logistic regression cost function	65
7.3	Experiments	66
8	COST-SENSITIVE DECISION TREES	69
8.1	Introduction	69
8.2	Decision trees	70
8.2.1	Construction of classification trees	70
8.2.2	Decision tree algorithms	75
8.3	Example-Dependent Cost-sensitive Decision Trees	75
8.3.1	Cost-sensitive impurity measures	75
8.3.2	Cost-sensitive pruning	76
8.4	Experiments	77
9	ENSEMBLES OF COST-SENSITIVE DECISION TREES	83
9.1	Introduction	83
9.2	Ensemble methods	84
9.2.1	Theoretical performance of an ensemble	85
9.3	Ensembles of cost-sensitive decision trees	86
9.3.1	Random inducers	86
9.3.2	Combination methods	87
9.3.3	Algorithms	89
9.4	Theoretical analysis of the ECSDT	90
9.5	Experiments	93
10	CONCLUSIONS	103
10.1	Future research directions	104
A	<i>costcla</i> : A COST-SENSITIVE CLASSIFICATION LIBRARY IN python	107
A.1	Library overview	107

A.2	Usage	108
A.3	Installation	110
BIBLIOGRAPHY		111

NOTATIONS

\mathcal{S}	Set of examples
$\mathcal{S}_{\text{train}}$	Training set
$\mathcal{S}_{\text{test}}$	Testing set
\mathcal{S}_0	Set of negative examples
\mathcal{S}_1	Set of positive examples
\mathbf{y}	Vector of class labels
\mathbf{x}_i	Feature vector of example i
k	Number of features
N	Number of examples in \mathcal{S}
N_0	Number of negative examples in \mathcal{S}_0
N_1	Number of positive examples in \mathcal{S}_1
y_i	Class label of example i
π_0	Percentage of negative examples in \mathcal{S}
π_1	Percentage of positive examples in \mathcal{S}
$f(\mathcal{S})$	A classification algorithm
$f_a(\mathcal{S})$	Function that classifies every example as a
\mathbf{c}	Predicted class labels
c_i	Predicted class for example i
\hat{p}_i	Predicted probability of positive for example i
t	Probability threshold
TP	Number of true positives examples
TN	Number of true negatives examples
FP	Number of false positives examples
FN	Number of false negatives examples
C_{FP}	Cost of a false positive
C_{FN}	Cost of a false negative
C_{TP_i}	Cost of a true positive for example i
C_{TN_i}	Cost of a true negative for example i
C_{FP_i}	Cost of a false positive for example i
C_{FN_i}	Cost of a false negative for example i
\mathbf{x}_i^*	Augmented feature vector of example i

$\mathbf{1}_c(z)$	Indicator function that takes the value of one if $z \in c$ and zero if $z \notin c$
$ \mathcal{S} $	Cardinality of a set \mathcal{S}
b_i	Classification problem cost characteristic
μ_b	Mean of the classification cost characteristic
σ_b	Standard deviation of the classification cost characteristic
C_a	Administrative cost of contacting a customer
Amt_i	Amount of transaction i
t_p	Number of hours to aggregate transactions
x_i^{id}	Customer who made transaction i
x_i^{time}	Time of transaction i
int_r	Annual interest rate charged to a customer
L_{gd}	Credit scoring loss given default
r_i	Expected profit of customer i
CL_i	Credit line of customer i
l_i	Term of the loan of customer i
int_{cf}	Cost of funds of a financial institution
debt_i	Outstanding debt of customer i
γ	Probability of accepting an offer
CLV	Average customer lifetime value
CLV_i	Customer lifetime value of example i
C_{o_i}	Cost of making an offer to customer i
Int_i	Expected profit generated by customer i
θ	Parameters of the logistic regression
h_θ	Logistic regression hypothesis
$J(\theta)$	Logistic regression cost function
$J^c(\theta)$	Cost-sensitive logistic regression cost function
\mathcal{V}	Splitting rule of a decision tree
I	Impurity measure of a decision tree
Tree	A decision tree
Branch	A branch of a decision tree

ACRONYMS

CS	Cost-sensitive classification
CI	Cost-insensitive classification
CPS	Cost-proportionate sampling
CST	Cost-sensitive training
LR	Logistic regression
DT	Decision tree
RF	Random forest
BMR	Bayes minimum risk
CSLR	Cost-sensitive logistic regression
CSDT	Cost-sensitive decision tree
ECSDT	Ensemble of cost-sensitive decision tree
CSB	Cost-sensitive bagging
CSP	Cost-sensitive pasting
CSRF	Cost-sensitive random forest
CSRP	Cost-sensitive random patches
t	Training
u	Under-sampling
r	Cost-proportionate rejection-sampling
o	Cost-proportionate over-sampling
mv	Majority voting
wv	Cost-sensitive weighted voting
s	Cost-sensitive stacking

LIST OF TABLES

Table 2.1	Classification confusion matrix	12
Table 3.1	Classification cost matrix	18
Table 3.2	Simplified classification cost matrix [Elkan, 2001] .	18
Table 4.1	Credit card fraud cost matrix	28
Table 4.2	Summary of typical raw credit card fraud detection features	29
Table 4.3	Example calculation of aggregated features. For all aggregated features $t_p = 24$	31
Table 4.4	Example calculation of periodic features.	34
Table 4.5	Credit scoring example-dependent cost matrix . .	37
Table 4.6	Credit scoring model parameters	40
Table 4.7	Summary of the financial datasets, where N is the number of examples and π_1 is the percentage of positive examples.	41
Table 5.1	Proposed churn modeling example-dependent cost matrix	47
Table 5.2	Direct marketing example-dependent cost matrix .	50
Table 5.3	Summary of the marketing datasets, where N is the number of examples and π_1 is the percentage of positive examples.	51
Table 6.1	Results of the algorithms measured by savings . .	59
Table 6.2	Results of the algorithms measured by Brier score	61
Table 7.1	Logistic regression cost matrix	65
Table 7.2	Results of the algorithms measured by savings . .	68
Table 7.3	Results of the algorithms measured by F_1 Score . .	68
Table 8.1	Results on the three datasets of the cost-sensitive and standard decision tree, without pruning (notp), with error based pruning (errp), and with cost-sensitive pruning technique (costp). Estimated using the different training sets: training, under-sampling, cost-proportionate rejection-sampling and cost-proportionate over-sampling	78

Table 8.2	Training time and tree size of the different cost-sensitive and standard decision tree, estimated using the different training sets: training, under-sampling, cost-proportionate rejection-sampling and cost-proportionate over-sampling, for the three databases.	80
Table 9.1	Results of the CI and BMR algorithms measured by savings	94
Table 9.2	Results of the CST and ECSDT algorithms measured by savings	95
Table 9.3	Savings Friedman ranking and average percentage of best result	96
Table 9.4	Savings ranks of best algorithm of each family by database	97
Table 9.5	Results of the CI and BMR algorithms measured by F_1 Score	101
Table 9.6	Results of the CST and ECSDT algorithms measured by F_1 Score	102
Table A.1	Results of the different algorithms using the <i>Cost-Cla</i> library	109

LIST OF FIGURES

Figure 1.1	Different example-dependent cost-sensitive algorithms grouped according to the stage in a classification system where they are used.	2
Figure 2.1	Classification process	10
Figure 2.2	Example of a classification algorithm. Using a set of examples from two classes, a classification algorithm is learned in order to separate between the positives and the negatives.	11
Figure 2.3	Example of a classification algorithm. Using a set of examples from two classes, a classification algorithm is learned in order to separate between the positives and the negatives.	13
Figure 3.1	Class-dependent cost-sensitive classification algorithm. Since the cost of misclassifying positives and negatives is different, the algorithm focus on maximizing the correct classification of the positives.	17
Figure 3.2	Example with example-dependent costs. Examples with the highest cost in darker colors, and the ones with the lowest cost are in lighter colors.	20
Figure 3.3	Example-dependent cost-sensitive classification algorithm. The algorithm focus first on correctly classify the dark examples, as the cost of misclassification in this cases is several times more expensive than the other cases.	20
Figure 4.1	Analysis of the time of a transaction using a 24 hour clock. The arithmetic mean of the transactions time (dashed line) do not accurately represents the actual times distribution.	32
Figure 4.2	Fitted von Mises distribution including the periodic mean (dashed line) and the probability distribution (purple area).	33
Figure 4.3	Expected time of a transaction (green area). Using the confidence interval, a transaction can be flag normal or suspicious, depending whether or not the time of the transaction is within the confidence interval.	34

Figure 4.4	Credit scoring sensitivity versus specificity thresholding procedure.	36
Figure 5.1	Flow analysis of a churn campaign [Verbraken, 2012]	45
Figure 5.2	Financial impact of the different decisions, i.e., False positives, false negatives, true positives and true negatives	46
Figure 5.3	Acceptance rate (γ) of the best offer for each customer profile. As expected, the higher the churn rate the lower the acceptance rate, as it is more difficult to make a good offer to a customer which is more likely to defect.	50
Figure 6.1	Estimation of calibrated probabilities using the ROC convex hull.	57
Figure 6.2	Comparison of the average savings of the algorithms versus the highest savings by family of classifiers. When the probabilities are calibrated there is a significant increase in savings.	60
Figure 6.3	Comparison of the average Brier score of the algorithms versus the lowest Brier score by family of classifiers. Overall, the models that are calibrated are indeed the ones with the best Brier score.	62
Figure 7.1	Sigmoid function	64
Figure 7.2	Comparison of the average savings and F_1Score of the algorithms versus the the best model. The models that perform the best measured by F_1 Score are not the best in terms of savings.	67
Figure 8.1	Impurity measures for a binary classification, as a function of the proportion of positive examples in the set (Cross-entropy is scaled) [Hastie et al., 2009].	71
Figure 8.2	Results of the DT and the CSDT. For both algorithms, the results are calculated with and without both types of pruning criteria.	79
Figure 8.3	Average savings on the three datasets of the different cost-sensitive and standard decision tree, estimated using the different training sets: training, under-sampling, cost-proportionate rejection-sampling and cost-proportionate over-sampling.	79
Figure 8.4	Average tree size (a) and training time (b), of the different cost-sensitive and standard decision tree, estimated using the different training sets.	81

Figure 9.1	Main reasons regarding why ensemble methods perform better than single models: statistical, computational and representational [Dietterich, 2000].	85
Figure 9.2	Visual representation of the random inducers algorithms.	87
Figure 9.3	Comparison of the savings of the algorithms versus the highest savings in each database. The CSRP – wt is very close to the best result in all the databases. Additionally, even though the LR – BMR is the best algorithm in two databases, the performance in the other three is very poor. . .	97
Figure 9.4	Comparison of the results by family of classifiers. The ECSDT family has the best performance measured either by Friedman ranking or average percentage of best model.	98
Figure 9.5	Comparison of the Friedman ranking within the ECSDT family. Overall, the random inducer method that provides the best results is the CSRP. Moreover, the best combination method compared by Friedman ranking is the cost-sensitive weighted voting.	99
Figure 9.6	Comparison of the Friedman ranking of the savings and F1Score sorted by F1Score ranking. The best two algorithms according to their Friedman rank of F1Score are indeed the best ones measured by the Friedman rank of the savings. However, this relation does not consistently hold for the other algorithms as the correlation between the rankings is just 65.10%.	100

INTRODUCTION

1.1 MOTIVATION AND SCOPE

Classification, in the context of machine learning, deals with the problem of predicting the class of a set of examples given their features. Traditionally, classification methods aim at minimizing the misclassification of examples in which an example is misclassified if the predicted class is different from the true class. Such a traditional framework assumes that all misclassification errors carry the same cost. This is not the case in many real-world applications. Methods that use different misclassification costs are known as cost-sensitive classifiers. Typical cost-sensitive approaches assume a constant cost for each type of error, in the sense that, the cost depends on the class and is the same among examples [Elkan, 2001; Kim et al., 2012].

This class-dependent approach is not realistic in many real-world applications. For example in credit card fraud detection, failing to detect a fraudulent transaction may have an economical impact from a few to thousands of Euros, depending on the particular transaction and card holder [Ngai et al., 2011]. In churn modeling, a model is used for predicting which customers are more likely to abandon a service provider. In this context, failing to identify a profitable or unprofitable churning has a significant different economic result [Verbraken et al., 2013]. Similarly, in direct marketing, wrongly predicting that a customer will not accept an offer when in fact he will, may have different financial impact, as not all customers generate the same profit [Zadrozny et al., 2003]. Lastly, in credit scoring, accepting loans from bad customers does not have the same economical loss, since customers have different credit lines, therefore, different profits [Verbraken et al., 2014].

Methods that use different misclassification costs are known as cost-sensitive classifiers. In particular, we are interested in methods that are example-dependent cost-sensitive, in the sense that the costs vary among examples and not only among classes [Elkan, 2001]. However, the literature on example-dependent cost-sensitive methods is limited, mostly because there is a lack of publicly available datasets that fit the problem [Aodha and Brostow, 2013]. Example-dependent cost-sensitive classification methods can be grouped according to the step where the costs are introduced into the system. Either the costs are introduced prior to the training of the algorithm, after the training or during train-

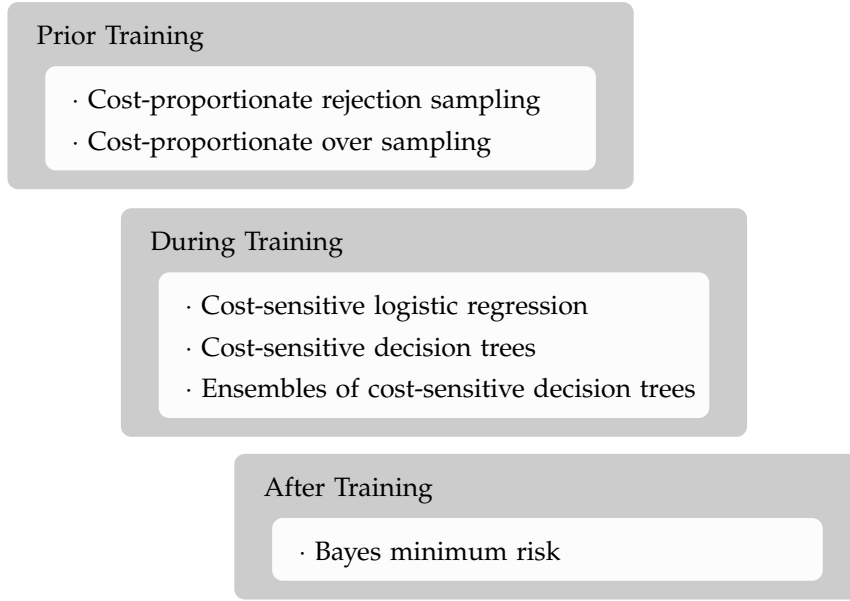


Figure 1.1: Different example-dependent cost-sensitive algorithms grouped according to the stage in a classification system where they are used.

ing [Wang, 2013]. In Figure 1.1, the different algorithms are grouped according to the stage in a classification system where they are used.

The first set of methods that were proposed to deal with cost-sensitivity consist in re-weighting the training examples based on their costs, either by cost-proportionate rejection-sampling [Zadrozny et al., 2003], or cost-proportionate over-sampling [Elkan, 2001]. The rejection-sampling approach consists in randomly selecting examples from a training set, and accepting each example with probability equal to the normalized misclassification cost of the example. On the other hand, the over-sampling method consists in creating a new set, by making n copies of each example, where n is related to the normalized misclassification cost of the example. These methods however, fail to introduce the example-dependent cost to the training of the different algorithms, and only rely on modifying the prior distribution of the training data.

The focus of this thesis is to investigate and define different example-dependent cost-sensitive classification algorithms, that not only focus on modifying the distribution of the training data but also introduce the different real financial costs during the training of the algorithms. We summarize the contributions of this thesis in the following section.

1.2 CONTRIBUTIONS

This dissertation summarizes several contributions to the field of example-dependent cost-sensitive machine learning.

(a) **Cost-sensitive financial evaluation measure**

Binary classification algorithms are normally evaluated using cost-insensitive evaluation measures, such as misclassification rate or F_1 Score. However, these measures may not be the most appropriate evaluation criteria when evaluating real-world cost-sensitive problems, because they tacitly assume that misclassification errors carry the same cost. In this thesis, we propose a new savings example-dependent cost-sensitive evaluation measure. The savings take into account the actual financial impact of the different misclassification errors. This work was published in [Correa Bahnsen et al., 2013].

(b) **Cost-sensitive credit card fraud detection**

Preventing credit card fraud is a classical example of a cost-sensitive problems, as the cost of a false negative is significantly different than the cost of a false positive. We discuss the particularities of credit card fraud detection and propose a financial evaluation measure that takes into account the economical costs associated with credit card fraud. Moreover, we propose to create a new set of features based on analyzing the periodic behavior of the time of a transaction using a modelling with a von Mises distribution. This work was published in [Correa Bahnsen et al., 2013, 2014b] and is under review in Correa Bahnsen et al. [2015d,e].

(c) **Example-dependent cost-sensitive real-world problems**

We analyze and propose financial evaluation measures for other real-world applications, namely, credit card fraud detection, churn modeling and direct marketing. The different analyses were presented in [Correa Bahnsen et al., 2014b,a, 2015b].

(d) **Cost-sensitive Bayes minimum risk**

In this thesis we propose a direct cost approach to make the classification decision based on the expected costs. This method is an extension of Bayes minimum risk, and consists in quantifying tradeoffs between various decisions using probabilities and the costs that accompany such decisions. This model was published in [Correa Bahnsen et al., 2013, 2014b].

(e) **Cost-sensitive logistic regression and decision trees**

The Bayes minimum risk method only introduces the costs after

the training of an algorithm, leaving opportunities to investigate the potential impact of algorithms that take into account the real financial example-dependent costs during the training of an algorithm. We propose a new cost-sensitive logistic regression. The method consists in introducing example-dependent costs into a logistic regression, by changing the objective function of the model to one that is cost-sensitive. We then apply a similar approach by introducing the costs in a cost-sensitive decision tree. The method is based on a new splitting criteria which is cost-sensitive, used during the tree construction. These methods were published in [Correa Bahnsen et al., 2014a, 2015a].

(f) **Ensembles of cost-sensitive decision trees**

Based on the cost-sensitive decision tree algorithm, we expand the cost-sensitive decision trees by creating a framework for an ensemble of cost-sensitive decision trees. This new method, step on the advantages of ensemble learning in order to create a more robust model. This work was submitted for review [Correa Bahnsen et al., 2015c].

(g) **CostCla: A cost-sensitive classification library**

The algorithms developed as part of this thesis are publicly available as part of the open-source *CostSensitiveClassification*¹ library.

1.3 OUTLINE

Part I of this manuscript is focused on giving the general concepts of classification and cost-sensitive classification. In particular, in Chapter 2, we give a background on classification. Then, in Chapter 3, we present the cost-sensitive problem and define the difference between cost-insensitive, class-dependent cost-sensitive and example-dependent cost-sensitive classification problems. Lastly, we give an introduction of the different evaluation measures used throughout this thesis.

Part II is dedicated to explaining the particularities of the four real-world classification problems that are the focus of this thesis, in particular, credit card fraud detection, credit scoring, churn modeling and direct marketing. In general, we show why each of the applications is example-dependent cost-sensitive, and we elaborate a framework for the analysis of each problem. This part is organized in two chapters. First, in Chapter 4, we discuss the applications within financial risk management. Second, in Chapter 5, we analyze marketing analytics applications.

¹ <https://github.com/albahnsen/CostSensitiveClassification>

Part III is focused on introducing our proposed example-dependent cost-sensitive methods. First, in Chapter 6, we present the Bayes minimum risk method. Then, we introduce the cost-sensitive logistic regression algorithm in Chapter 7. Afterwards, in Chapter 8, we show and discuss the cost-sensitive decision trees algorithm. Lastly, in Chapter 9, we present our framework for ensembles of cost-sensitive decision trees.

Chapter 10 concludes the thesis, and elaborates on possible lines for future research. Lastly, in the Appendix A, we present the library *CostCla* developed as part of the thesis. This library is an open-source implementation of all the algorithms covered in this manuscript.

1.4 PUBLICATIONS

Publications from this work are:

- (a) [Correa Bahnsen et al., 2013] *Cost Sensitive Credit Card Fraud Detection Using Bayes Minimum Risk*, Alejandro Correa Bahnsen, Aleksandar Stojanovic, Djamila Aouada and Björn Ottersten. In Proceedings of IEEE International Conference on Machine Learning and Applications, 2013.
- (b) [Correa Bahnsen et al., 2014b] *Improving Credit Card Fraud Detection with Calibrated Probabilities*, Alejandro Correa Bahnsen, Aleksandar Stojanovic, Djamila Aouada and Björn Ottersten. In Proceedings of SIAM International Conference on Data Mining, 2014.
- (c) [Correa Bahnsen et al., 2014a] *Example-Dependent Cost-Sensitive Logistic Regression for Credit Scoring*, Alejandro Correa Bahnsen, Djamila Aouada and Björn Ottersten. In Proceedings of IEEE International Conference on Machine Learning and Applications, 2014.
- (d) [Correa Bahnsen et al., 2015a] *Example-Dependent Cost-Sensitive Decision Trees*, Alejandro Correa Bahnsen, Djamila Aouada and Björn Ottersten. In Expert Systems with Applications, 42(19):6609-6619, 2015.
- (e) [Correa Bahnsen et al., 2015b] *A novel cost-sensitive framework for customer churn predictive modeling*, Alejandro Correa Bahnsen, Djamila Aouada and Björn Ottersten. In Decision Analytics, 2(1):5, 2015.
- (f) [Correa Bahnsen et al., 2015c] *Ensembles of Example-Dependent Cost-Sensitive Decision Trees*, Alejandro Correa Bahnsen, Djamila Aouada and Björn Ottersten. Submitted to IEEE Transactions on Knowledge and Data Engineering, 2015.

- (g) [[Correa Bahnsen et al., 2015d](#)] *Feature Engineering in Credit Card Fraud Detection*, Alejandro Correa Bahnsen, Djamila Aouada and Björn Ottersten. Submitted to Expert Systems with Applications, 2015.
- (h) [[Correa Bahnsen et al., 2015e](#)] *Detecting Credit Card Fraud using Periodic Features*, Alejandro Correa Bahnsen, Djamila Aouada and Björn Ottersten. Submitted to the IEEE International Conference on Machine Learning and Applications, 2015.

Part I

EXAMPLE-DEPENDENT COST-SENSITIVE CLASSIFICATION

BASICS ON CLASSIFICATION

OUTLINE

In this chapter, we introduce classification models from a machine learning perspective. First, in Section 2.1, we give a self contained introduction to classification, including the most-common algorithms, and the main applications of classification models. Then, in Section 2.2, we present the different evaluation measures that are normally used for analyzing the performance of classification methods.

2.1 INTRODUCTION

In machine learning, classification refers to the attempt of identifying to which of a set of classes a new example belongs, based on learning from examples whose class membership is known. A classification task begins with a training set in which the class of a set of examples is known. For example, a classification model that predicts credit card fraud is developed by analyzing many observed credit transactions over a period of time. The class in this case is a variable which indicates for each example whether or not the transaction was or not a fraud. Also, the features, are the transaction attributes like place, amount and time of the transaction.

Then, during the training process, a classification algorithm finds the patterns and relationships between the values of the features and the values of the target class. Different algorithms use different methods and techniques to estimate these relationships. Afterwards, these relationships are summarized in a model that is able to make predictions on new sets of data.

In general, there are two types of classification models: binary and multi-class. In binary classification problems, the objective is to classify examples between two classes, usually referred to as the negative and positive classes. On the other hand, multi-class problems are not bound to two classes but instead aim to classify examples among a number of classes. In this work we focus on binary classification problems.

Binary classification algorithms are widely used across a variety of domains. For example in the medical field, models have been used for making predictions about tumors, probability of a disease, probability of selecting the right drug for a particular patient, and estimating the

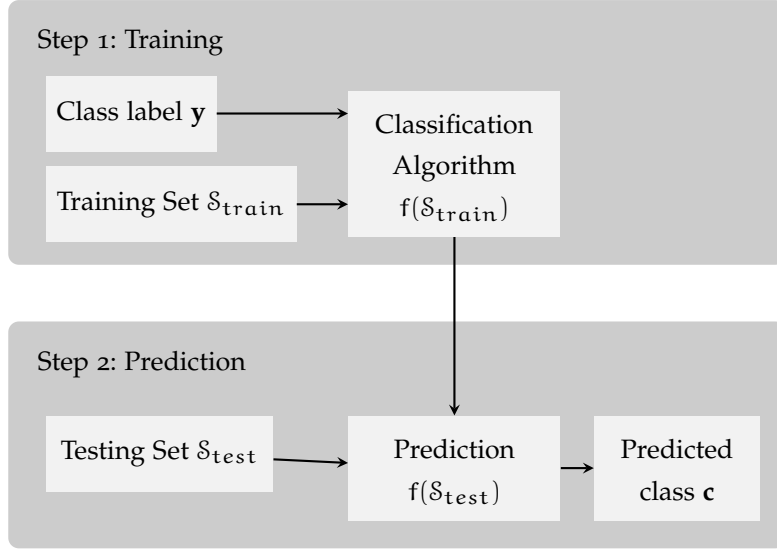


Figure 2.1: Classification process

probability of relapsing, among others [Herland et al., 2014]. In the financial sector, classification models have been successfully applied for fraud detection, credit scoring, portfolio management and algorithmic trading. Also, in marketing, several models are being currently used for churn modeling, customer targeting, behavior prediction and direct marketing [Baesens, 2014]. Additionally, classification algorithms are used in many other emerging applications such as terrorism prevention, malware detection, computer security, energy consumption prediction, spam classification, and others [Kriegel et al., 2007].

Formally, a binary classification algorithm deals with the problem of predicting the class y_i of a set S of examples or instances i , given their k features $x_i \in \mathbb{R}^k$. The objective is to construct a function $f(S)$ that makes a prediction c_i of the class of each example i from S using its feature vector x_i , where $|S| = N$. Moreover, some algorithms allow to not only estimate the prediction, but also its confidence, in the form of the probability \hat{p}_i of belonging to the positive class, i.e. $c_i = 1$. The way for finding from \hat{p}_i to c_i is simply by defining a probability threshold t , and applying the following formula

$$c_i = \begin{cases} 0 & \text{if } \hat{p}_i \leq t \\ 1 & \text{otherwise,} \end{cases} \quad (2.1)$$

Usually $t = \frac{1}{2}$ [Hastie et al., 2009]. However, if $t \neq \frac{1}{2}$, the function that generates the predicted class labels $c = [c_i]$ is denoted as f^t .

In Figure 2.1, the process of training and prediction in a classification algorithm are summarized. First, during the training phase, using a training set S_{train} , an algorithm is trained to predict y , where $y = [y_i]$.




Figure 2.2: Example of a classification algorithm. Using a set of examples from two classes, a classification algorithm is learned in order to separate between the positives and the negatives.

Then the algorithm is used to estimate the classes c of a set of testing examples S_{test} .

There exists several algorithms that can be used for classification tasks. In general a classification algorithm is learned with the objective of finding patterns that separate between the different classes [Hastie et al., 2009]. In order to clarify this intuition, in Figure 2.2 an example of a classification algorithm is shown. Let us consider a set of examples, where the red points represent the positive examples and the blue ones the negative examples. The objective of a classifier is to find the best way to separate between the positive and negative examples. Then, the output of a classifier learned using the set of training examples is shown as the dashed black line. It is observed that this classifier is able to separate almost all the examples using a linear classifier. However, not all examples are correctly classified. In particular, there are four negative examples that were predicted as positive, and five positive examples that were predicted as negative. In the next section, we present the standard methods for evaluating the performance of a classification algorithm.

2.2 TRADITIONAL EVALUATION MEASURES

When evaluating the performance of a classification algorithm, the first thing to do is to check the number of examples that were misclassified, since the true class of the examples is known. Therefore, evaluating

	Actual Positive $y_i = 1$	Actual Negative $y_i = 0$
Predicted Positive $c_i = 1$	True Positive (TP)	False Positive (FP)
Predicted Negative $c_i = 0$	False Negative (FN)	True Negative (TN)

Table 2.1: Classification confusion matrix

the error of a model is as simple as counting the number of times an example is misclassified divided by the number of examples

$$\text{Err}(f(\mathcal{S})) = 1 - \frac{1}{N} \sum_{i=1}^N \mathbf{1}_{y_i}(c_i), \quad (2.2)$$

where $\mathbf{1}_q(z)$ is an indicator function that is calculated as:

$$\mathbf{1}_q(z) = \begin{cases} 1 & \text{if } z = q \\ 0 & \text{if } z \neq q. \end{cases} \quad (2.3)$$

Moreover the accuracy is defined as the percentage of times the algorithm made the correct prediction

$$\text{Acc}(f(\mathcal{S})) = 1 - \text{Err}(f(\mathcal{S})). \quad (2.4)$$

However, just knowing these statistics is not enough to make decisions, as in many applications it is important to know where the errors are coming from. In particular, the misclassified examples may belong only to one class, which may give interesting insights about the problem. A way to observe the different errors is by looking at the confusion matrix, as shown in Table 2.1. Afterwards, using the cost matrix several statistics are extracted. In particular:

$$\text{Recall} = \frac{\text{TP}}{\text{TP} + \text{FN}} \quad (2.5)$$

$$\text{Precision} = \frac{\text{TP}}{\text{TP} + \text{FP}} \quad (2.6)$$

$$\text{F}_1\text{Score} = 2 \frac{\text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}} \quad (2.7)$$




Figure 2.3: Example of a classification algorithm. Using a set of examples from two classes, a classification algorithm is learned in order to separate between the positives and the negatives.

As an illustrative example, the different statistics are calculated for the example presented in Section 2.1. First, the confusion matrix is calculated as follows:

	Actual Positive $y_i = 1$	Actual Negative $y_i = 0$
Predicted Positive $c_i = 1$	36	4
Predicted Negative $c_i = 0$	5	68

Then using the confusion matrix, the different statistics are calculated as: Error = 11.11%, Recall = 87.8%, Precision = 90% and F_1 Score = 88.8%.

There are, however, several instances that are misclassified, that is because the simple linear classifier that was used in this example may not be good enough to separate between the positive and negative classes. In order to make a comparison, using the same example, a new algorithm is learned. This time the algorithm made the correct prediction more often as shown in Figure 2.3. Afterwards, the confusion matrix is calculated as follows:

	Actual Positive $y_i = 1$	Actual Negative $y_i = 0$
Predicted Positive $c_i = 1$	37	2
Predicted Negative $c_i = 0$	4	70

Then the different statistics are calculated as: Error = 5.3%, Recall = 90.2%, Precision = 94.9% and F₁Score = 92.5%. It is observed that in this case the FP are reduced more than the FN, this leads to a higher increase in precision than in recall. There is not a single rule regarding which one is more important than the other, it depends on the application. For example in applications with a high false negative cost such as failing to identify a tumor in a medical exam, the recall should be the priority, even if that implies having a significant number of false positives. On the other hand, In applications such as spam detection, predicting a normal email as spam it may have a large impact on the customer, therefore, in this example is better to allow some false negatives and focus on the false positives.

It is not always straightforward to define the right tradeoff between false positives and false negatives. The best approximation to solve that, is to focus on the actual costs incurred by the different decisions. This is usually solved using cost-sensitive classification methods.

2.2.1 Brier score

Traditional evaluation measures of binary classification problems, such as Accuracy and F₁Score, provide a way to analyze the performance of a model. However, when using the classifier output as a basis for decision making, there is a need of a measure that takes into account not only the misclassification of a classifier predicted class c , but also the quality of the estimated probabilities \hat{p} [Cohen and Goldszmidt, 2004]. The most appropriate is the Brier score [Brier, 1950]. The Brier score belongs to the class of so-called proper scores which are used in evaluating the subjective probability assessment of the prediction [DeGroot and Fienberg, 1983]. The Brier score is the average squared difference between the estimated probability and the true class label. It is defined as:

$$BS(f(S)) = \frac{1}{N} \sum_{i=1}^N (\hat{p}_i - y_i)^2. \quad (2.8)$$

The main justification of this score is based on decision theoretic considerations, in the sense that, a forecaster should pay a price proportional to the confidence with which it asserts its decision.

OUTLINE

In this chapter, we introduce the problem of cost-sensitive classification. Standard classification models aim at minimizing the misclassification of examples, in which an example is misclassified if the predicted class is different from the true class. However, this is not the case in many real-world applications. In this chapter, first, we introduce cost-sensitive classification in Section 3.1. Then, in Section 3.2, we present the problem of class-dependent cost-sensitive classification. Then, in Section 3.3, we present the general framework of example-dependent cost-sensitive classification. Within this section, we first introduce a method for defining the type of cost-sensitivity of a problem. Afterwards, we present the different cost-sensitive performance evaluation measures. Lastly, we present state-of-the-art example-dependent cost-sensitive methods, namely, cost-proportionate rejection-sampling and cost-proportionate over-sampling.

3.1 INTRODUCTION

Classification methods are used to predict the class of different examples given their features. Standard methods aim at maximizing the accuracy of the predictions, in which an example is correctly classified if the predicted class is the same as the true class. This traditional approach assumes that all correctly classified and misclassified examples carry the same cost. This, however, is not the case in many real-world applications. Methods that use different misclassification costs are known as cost-sensitive classifiers. Typical cost-sensitive approaches assume a constant cost for each type of error, in the sense that, the cost depends on the class and is the same among examples [Elkan, 2001; Kim et al., 2012]. Nevertheless, this class-dependent approach is not realistic in many real-world applications.

For example in credit card fraud detection, failing to detect a fraudulent transaction may have an economical impact from a few to thousands of Euros, depending on the particular transaction and card holder [Sahin et al., 2013]. In churn modeling, a model is used for predicting which customers are more likely to abandon a service provider. In this context, failing to identify a profitable or unprofitable churning customer has a significant different financial impact [Glady et al., 2009]. Similarly, in direct

marketing, wrongly predicting that a customer will not accept an offer when in fact he will, has a different impact than the other way around [Zadrozny et al., 2003]. Also in credit scoring, where declining good customers has a non constant impact since not all customers generate the same profit [Verbraken et al., 2014]. Lastly, in the case of intrusion detection, classifying a benign connection as malicious has a different cost than when a malicious connection is accepted [Ma et al., 2011].

In order to deal with these specific types of cost-sensitive problems, called example-dependent cost-sensitive, some methods have been proposed recently. However, the literature on example-dependent cost-sensitive methods is limited, mostly because there is a lack of publicly available datasets that fit the problem [Aodha and Brostow, 2013]. Standard solutions consist in modifying the training set by re-weighting the examples proportionately to the misclassification costs [Elkan, 2001; Zadrozny et al., 2003].

3.2 CLASS-DEPENDENT COST-SENSITIVE CLASSIFICATION

The literature in cost-sensitive classification is mostly focused in class-dependent problems [Elkan, 2001], where the cost of misclassification is associated with the class. Usually, the cost of misclassifying a positive example is denoted by C_{FN} and the one of misclassifying a negative example is denoted by C_{FP} . Conceptually, $C_{FN} \geq 0$ and $C_{FP} \geq 0$; moreover, they are normally defined such that $C_{FN} + C_{FP} = 2$ [Flach et al., 2011], as when $C_{FN} = C_{FP} = 1$ represents the case of cost-insensitive classification. Using the previous notation, a class-dependent cost measure is defined as [Wang et al., 2014]:

$$\text{Cost}_{cd}(f(S)) = C_{FP} \cdot FP + C_{FN} \cdot FN. \quad (3.1)$$

Over the past decades, various algorithms have been proposed for class-dependent cost-sensitive classification in literature. Several authors have used modifications of the decision trees that take into account the different class-dependent costs [Draper et al., 1994; Ting, 2002; Ling et al., 2004; Li et al., 2005; Kretowski and Grześ, 2006; Vadera, 2010; Lomax and Vadera, 2013]. Similarly, applications of bagging and boosting algorithms have been used for cost-sensitive classification [Nesbitt, 2010; Street, 2008; Masnadi-shirazi and Vasconcelos, 2011; Fan et al., 1999]. Recently, various variations to support vector machines have also been used for this problem [Li et al., 2010; Masnadi-shirazi, 2010]. Lastly, online learning algorithms have also been used for cost-sensitive tasks [Wang et al., 2014].

Following the example shown in Section 2.1, we now assume that misclassifying a negative example has a cost of $C_{FN} = 0.2$ and for a pos-




Figure 3.1: Class-dependent cost-sensitive classification algorithm. Since the cost of misclassifying positives and negatives is different, the algorithm focus on maximizing the correct classification of the positives.

itive example, the cost is $C_{FP} = 1.8$. Under this scenario, misclassifying a positive example has a much higher cost than misclassifying a negative one. Taking that into account, in Figure 3.1, we show an algorithm (Algorithm3), that focused on maximizing the correct classification of the positives. In the following table, we compare the results of the standard measures and the class-dependent cost, of Algorithm3 and the classification algorithms presented in Figure 2.2 (Algorithm1) and Figure 2.3 (Algorithm2):

Algorithm	Error	Recall	Precision	F ₁ Score	Cost _{cd}
Algorithm1	11.11%	87.8%	90%	88.8%	9.8
Algorithm2	5.3%	90.2%	94.9%	92.5%	7.6
Algorithm3	7.97%	92.68%	86.36%	89.41%	6.6

It is found, that by focusing on the positives, Algorithm3 arises to a lower cost, even though the traditional metrics are worse for Algorithm3 than for Algorithm2. In conclusion, it is of highly importance to take into account the cost when evaluating and training a classification model.

3.3 EXAMPLE-DEPENDENT COST-SENSITIVE CLASSIFICATION

The class-dependent framework introduced in the previous section is highly restrictive, as assuming that the different costs are constant between classes is not realistic in many real world applications. In fraud detection, fraudulent transactions can have a financial impact from hun-

	Actual Positive $y_i = 1$	Actual Negative $y_i = 0$
Predicted Positive $c_i = 1$	C_{TP_i}	C_{FP_i}
Predicted Negative $c_i = 0$	C_{FN_i}	C_{TN_i}

Table 3.1: Classification cost matrix

Negative	$C_{FN_i}^* = \frac{(C_{FN_i} - C_{TN_i})}{(C_{FP_i} - C_{TN_i})}$
Positive	$C_{TP_i}^* = \frac{(C_{TP_i} - C_{TN_i})}{(C_{FP_i} - C_{TN_i})}$

Table 3.2: Simplified classification cost matrix [Elkan, 2001]

dreds or thousands of Euros [Sahin et al., 2013]. In this context, the example-dependent costs can be represented using a 2x2 cost matrix [Elkan, 2001], that introduces the costs associated with two types of correct classification, cost of true positives (C_{TP_i}), cost of true negatives (C_{TN_i}), and the two types of misclassification errors, cost of false positives (C_{FP_i}), cost of false negatives (C_{FN_i}), as defined in Table 3.1.

Conceptually, the cost of correct classification should always be lower than the cost of misclassification. These are referred to as the “reasonableness” conditions [Elkan, 2001], and are defined as $C_{FP_i} > C_{TN_i}$ and $C_{FN_i} > C_{TP_i}$. Taking into account the “reasonableness” conditions, a simpler cost matrix with only one degree of freedom has been defined in [Elkan, 2001], by scaling and shifting the initial cost matrix. The simpler cost-matrix is shown in Table 3.2.

3.3.1 Example-dependent evaluation measures

Common cost-insensitive evaluation measures, such as misclassification rate or F_1 Score, assume the same cost for the different misclassification errors. Using these measures is not suitable for example-dependent cost-sensitive binary classification problems. Indeed, two classifiers with equal misclassification rates but different numbers of false positives and false negatives do not have the same impact on cost since $C_{FP_i} \neq C_{FN_i}$; therefore, there is a need for a measure that takes into account the actual costs $\{C_{TP_i}, C_{FP_i}, C_{FN_i}, C_{TN_i}\}$ of each example i , as introduced in Section 3.3.

Let \mathcal{S} be a set of N examples i , $N = |\mathcal{S}|$, where each example is represented by the augmented feature vector $\mathbf{x}_i^* = [\mathbf{x}_i, C_{TP_i}, C_{FP_i}, C_{FN_i}, C_{TN_i}]$

and labeled using the class label $y_i \in \{0, 1\}$. A classifier f which generates the predicted label c_i for each element i is trained using the set \mathcal{S} . Then the cost of using f on \mathcal{S} is calculated by

$$\text{Cost}(f(\mathcal{S})) = \sum_{i=1}^N \text{Cost}(f(\mathbf{x}_i^*)), \quad (3.2)$$

where

$$\begin{aligned} \text{Cost}(f(\mathbf{x}_i^*)) = & y_i(c_i C_{TP_i} + (1 - c_i) C_{FN_i}) + \\ & (1 - y_i)(c_i C_{FP_i} + (1 - c_i) C_{TN_i}). \end{aligned} \quad (3.3)$$

However, the total cost may not be easy to interpret. In [Whitrow et al., 2008], a *normalized* cost measure was proposed, by dividing the total cost by the theoretical maximum cost, which is the cost of misclassifying every example. The *normalized* cost is calculated using

$$\text{Cost}_n(f(\mathcal{S})) = \frac{\text{Cost}(f(\mathcal{S}))}{\sum_{i=1}^N C_{FN_i} \cdot \mathbf{1}_0(y_i) + C_{FP_i} \cdot \mathbf{1}_1(y_i)}. \quad (3.4)$$

We propose similar approach in [Correa Bahnsen et al., 2014a], where the savings of using an algorithm are defined as the cost of the algorithm versus the cost of using no algorithm at all. To do that, the cost of the costless class is defined as

$$\text{Cost}_l(\mathcal{S}) = \min\{\text{Cost}(f_0(\mathcal{S})), \text{Cost}(f_1(\mathcal{S}))\}, \quad (3.5)$$

where

$$f_a(\mathcal{S}) = \{a\}, \text{ with } a \in \{0, 1\}. \quad (3.6)$$

The cost improvement can be expressed as the cost savings as compared with $\text{Cost}_l(\mathcal{S})$.

$$\text{Savings}(f(\mathcal{S})) = \frac{\text{Cost}_l(\mathcal{S}) - \text{Cost}(f(\mathcal{S}))}{\text{Cost}_l(\mathcal{S})}. \quad (3.7)$$

In order to illustrate this concept, we use the same example shown in Section 2.1. However, now we add the variation in the costs by showing the examples with the highest cost darker, and the ones with the lowest cost lighter. The new example is shown in Figure 3.2. Moreover, in the following table we summarize the different example-dependent costs:

Class	Cost light	Cost normal	Cost dark
Negative	0.1	0.5	5.0
Positive	1.0	2.0	10.0




Figure 3.2: Example with example-dependent costs. Examples with the highest cost in darker colors, and the ones with the lowest cost are in lighter colors.




Figure 3.3: Example-dependent cost-sensitive classification algorithm. The algorithm focus first on correctly classify the dark examples, as the cost of misclassification in this cases is several times more expensive than the other cases.

Taking into account the example-dependent costs, in Figure 3.3 (Algorithm4), we show a new algorithm that gives a higher importance on correctly classify the dark examples, as the cost of misclassification in this cases is several times more expensive than the other cases. Furthermore, in the following table, we compare the results of the standard measures and the example-dependent savings of Algorithm4 and the

classification algorithms presented in Figure 2.2 (Algorithm1), Figure 2.3 (Algorithm2) and Figure 3.1 (Algorithm3):

Algorithm	Error	Recall	Precision	F ₁ Score	Savings
Algorithm1	11.11%	87.8%	90%	88.8%	46.86%
Algorithm2	5.3%	90.2%	94.9%	92.5%	68.36%
Algorithm3	7.97%	92.68%	86.36%	89.41%	48.07%
Algorithm4	6.19%	92.68%	90.48%	91.56%	87.42%

With these examples, we illustrate the impact that the costs have on the algorithms. Moreover, we highlight the importance of using the different costs when evaluating the different models. It is worth mentioning how distinct are the results if the costs are ignored and an algorithm is trained and evaluated not taking into account the different costs present in most real-world applications.

3.3.2 Binary classification cost characteristic

A classification problem is said to be cost-insensitive if costs of both errors are equal. It is class-dependent cost-sensitive if the costs are different but constant. Finally we talk about an example-dependent cost-sensitive classification problem if the cost matrix is not constant for all the examples.

However, the definition above is not general enough. There are many cases when the cost matrix is not constant and still the problem is cost-insensitive or class-dependent cost-sensitive. For example, if the costs of correct classification are zero, $C_{TP_i} = C_{TN_i} = 0$, and the costs of misclassification are $C_{FP_i} = a_0 \cdot z_i$ and $C_{FN_i} = a_1 \cdot z_i$, where a_0 and a_1 , are constants and z_i is a random variable. This is an example of a cost matrix that is not constant. However, $C_{FN_i}^*$ and $C_{TP_i}^*$ are constant, i.e. $C_{FN_i}^* = (a_1 \cdot z_i)/(a_0 \cdot z_i) = a_1/a_0$ and $C_{TP_i}^* = 0 \forall i$. In this case the problem is cost-insensitive if $a_0 = a_1$, or class-dependent cost-sensitive if $a_0 \neq a_1$, even given the fact that the cost matrix is not constant.

Nevertheless, using only the simpler cost matrix is not enough to define when a problem is example-dependent cost-sensitive. To achieve this, we defined the following cost characteristic for a given binary classification problem as:

$$b_i = C_{FN_i}^* - C_{TP_i}^*, \quad (3.8)$$

and define its mean and standard deviation as μ_b and σ_b , respectively.

Using μ_b and σ_b , we analyze different binary classification problems. A binary classification problem is defined according to the following conditions:

μ_b	σ_b	Type of classification problem
1	0	cost-insensitive
$\neq 1$	0	class-dependent cost-sensitive
	$\neq 0$	example-dependent cost-sensitive

3.3.3 State-of-the-art methods

As mentioned earlier, taking into account the different costs associated with each example, some methods have been proposed to make classifiers example-dependent cost-sensitive. These methods may be grouped in two categories. Methods based on changing the class distribution of the training data, which are known as cost-proportionate sampling methods; and direct cost methods [Wang, 2013].

A standard method to introduce example-dependent costs into classification algorithms is to re-weight the training examples based on their costs, either by cost-proportionate rejection-sampling [Zadrozny et al., 2003], or over-sampling [Elkan, 2001]. The rejection-sampling approach consists in selecting a random subset \mathcal{S}_r by randomly selecting examples from \mathcal{S} , and accepting each example i with probability $w_i / \max_{1, \dots, N} \{w_i\}$, where w_i is defined as the expected misclassification error of example i :

$$w_i = y_i \cdot C_{FN_i} + (1 - y_i) \cdot C_{FP_i}. \quad (3.9)$$

Lastly, the over-sampling method consists in creating a new set \mathcal{S}_o , by making w_i copies of each example i . However, cost-proportionate over-sampling increases the training since $|\mathcal{S}_o| \gg |\mathcal{S}|$, and it also may result in over-fitting [Drummond and Holte, 2003]. Furthermore, none of these methods uses the full cost matrix but only the misclassification costs.

The second approach consists in using the predicted probability \hat{p}_i , estimated using a given classifier f , and modify the threshold t such that the savings are maximized. This method is called cost-sensitive thresholding [Sheng and Ling, 2006]. The idea behind this approach is to adaptively modify the probability threshold of an algorithm f^t in order to maximize the savings $\text{Savings}(f^t(\mathcal{S}))$ of the algorithm f^t on a given set \mathcal{S} . The threshold is calculated using the following equation

$$t_{\text{thresholding}} = \arg \max_t \text{Savings}(f^t(\mathcal{S})). \quad (3.10)$$

Part II

REAL-WORLD EXAMPLE-DEPENDENT APPLICATIONS

OUTLINE

In this chapter, we present two different real-world example-dependent cost-sensitive problems related to financial risk modeling, namely, credit card fraud detection and credit scoring. Both problems are example-dependent cost-sensitive, as failing to identify a fraudulent transaction may have a financial impact ranging from tens to thousands of Euros. Similarly, in credit scoring, approving a customer that later does not pay his debt has a significant impact on a bank profit. First, in Section 4.1, we introduce the credit card fraud detection problem and propose new periodic features. Lastly, in Section 4.2, we present the credit scoring problem and propose a new financial evaluation measure for this problem.

4.1 CREDIT CARD FRAUD DETECTION

The use of credit and debit cards has increased significantly in the last years, unfortunately so has the fraud. Because of that, billions of Euros are lost every year. According to the European Central Bank [European Central Bank, 2014], during 2012 the total level of fraud reached 1.33 billion Euros in the Single Euro Payments Area, which represents an increase of 14.8% compared with 2011. Moreover, payments across non traditional channels (mobile, internet, ...) accounted for 60% of the fraud, whereas it was 46% in 2008. This opens new challenges as new fraud patterns emerge, and current fraud detection systems are not being successful in preventing fraud. Furthermore, fraudsters constantly change their strategies to avoid being detected, something that makes traditional fraud detection tools such as expert rules inadequate.

The use of machine learning in fraud detection has been an interesting topic in recent years. Several detection systems based on machine learning techniques have been successfully used for this problem [Bhattacharyya et al., 2011]. When constructing a credit card fraud detection model, there are several problems that have an important impact during the training phase: Skewness of the data, cost-sensitivity of the application, short time response of the system, dimensionality of the search space and how to preprocess the features [Bolton et al., 2002; Bachmayer, 2008; Whitrow et al., 2008; Pozzolo et al., 2014a; Van Vlas-

[selaer et al., 2015](#)]. We are interested in addressing the cost-sensitivity and the features preprocessing issues.

Credit card fraud detection is by definition a cost-sensitive problem, in the sense that the cost due to a false positive is different than the cost of a false negative. When predicting a transaction as fraudulent, when in fact it is not a fraud, there is an administrative cost that is incurred by the financial institution. On the other hand, when failing to detect a fraud, the amount of that transaction is lost [[Hand et al., 2007](#)]. Moreover, it is not enough to assume a constant cost difference between false positives and false negatives, as the amount of the transactions varies quite significantly; therefore, its financial impact is not constant but depends on each transaction.

When constructing a credit card fraud detection model, it is very important to use those features that will help the algorithm make the best decision. Typical models only use raw transactional features, such as time, amount, place of the transaction. However, these approaches do not take into account the spending behavior of the customer, which is expected to help discover fraud patterns [[Bachmayer, 2008](#)]. A standard way to include these behavioral spending patterns was proposed in [[Whitrow et al., 2008](#)], where Whitrow et al. proposed a transaction aggregation strategy in order to take into account a customer spending behavior. The derivation of the aggregated features consists in grouping the transactions made during the last given number of hours, first by card or account number, then by transaction type, merchant group, country or other, followed by calculating the number of transactions or the total amount spent on those transactions.

In this section, we first present a new cost-based measure to evaluate credit card fraud detection models, taking into account the different financial costs incurred by the fraud detection process. Afterwards, we propose an expanded version of the transaction aggregation strategy, by incorporating a combination criteria when grouping transactions, i.e., instead of aggregating only by card holder and transaction type, we combine it with country or merchant group. This allows to have a much richer feature space. Moreover, we are interested in analyzing the impact of adding the time of a transaction. The logic behind it, is that a customer is expected to make transactions at similar hours. We, hence, propose a new method for creating features based on the periodic behavior of a transaction time, using the von Mises distribution [[Fisher, 1996](#)]. In particular, these new time features should estimate if the time of a new transaction is within the confidence interval of the previous transaction time. Furthermore, we present the real credit card fraud dataset provided by a large European card processing company used for the experiments in this thesis.

4.1.1 Credit card fraud detection evaluation

A credit card fraud detection algorithm consists in identifying those transactions with a high probability of being fraud, based on historical fraud patterns. The use of machine learning in fraud detection has been an interesting topic in recent years. Different detection systems that are based on machine learning techniques have been successfully used for this problem, in particular: neural networks [Maes et al., 2002], Bayesian learning [Maes et al., 2002], artificial immune systems [Bachmayer, 2008], association rules [Sánchez et al., 2009], hybrid models [Krivko, 2010], support vector machines [Bhattacharyya et al., 2011], peer group analysis [Weston et al., 2008], online learning [Pozzolo et al., 2014a], discriminant analysis [Mahmoudi and Duman, 2015] and social network analysis [Van Vlasselaer et al., 2015].

Most of these studies compare their proposed algorithm with a benchmark algorithms and then make the comparison using a standard binary classification measure, such as misclassification error, receiver operating characteristic (ROC), Kolmogorov-Smirnov (KS) or F₁Score statistics [Bolton et al., 2002; Hand et al., 2007; Pozzolo et al., 2014a]. However, these measures may be not the most appropriate evaluation criteria when evaluating fraud detection models, because they tacitly assume that misclassification errors carry the same cost, similarly with the correct classified transactions. This assumption does not hold in practice, since wrongly predicting a fraudulent transaction as legitimate carries a significantly different financial cost than the inverse case. Furthermore, the accuracy measure also assumes that the class distribution among transactions is constant and balanced [Provost et al., 1998], and typically the distributions of a fraud detection dataset are skewed, with a percentage of frauds ranging from 0.005% to 0.5% [Bachmayer, 2008; Bhattacharyya et al., 2011].

In order to take into account the different costs of fraud detection during the evaluation of an algorithm, we may use the cost matrix defined in Table 3.1. Hand et al. [Hand et al., 2007] proposed a cost matrix, where in the case of false positive the associated cost is the administrative cost $C_{FP_i} = C_a$ related to analyzing the transaction and contacting the card holder. This cost is the same assigned to a true positive $C_{TP_i} = C_a$, because in this case, the card holder will have to be contacted. However, in the case of a false negative, in which a fraud is not detected, the cost is defined to be a hundred times larger, i.e. $C_{FN_i} = 100C_a$. This same approach was also used in [Bachmayer, 2008].

Nevertheless, in practice, losses due to a specific fraud range from few to thousands of Euros, which means that assuming constant cost for false negatives is unrealistic. In order to address this limitation, we pro-

	Actual Positive $y_i = 1$	Actual Negative $y_i = 0$
Predicted Positive $c_i = 1$	$C_{TP_i} = C_a$	$C_{FP_i} = C_a$
Predicted Negative $c_i = 0$	$C_{FN_i} = Amt_i$	$C_{TN_i} = 0$

Table 4.1: Credit card fraud cost matrix

pose a cost matrix that takes into account the actual example-dependent financial costs. Our cost matrix defines the cost of a false negative to be the amount $C_{FN_i} = Amt_i$ of the transaction i . We argue that this cost matrix is a better representation of the actual costs, since when a fraud is not detected, the losses of that particular fraud correspond to the stolen amount. The costs are summarized in Table 4.1.

Afterwards, using (3.3) a cost measure for fraud detection is calculated as:

$$\text{Cost}(f(S)) = \sum_{i=1}^N y_i(1 - c_i)Amt_i + c_i C_a, \quad (4.1)$$

then, the savings of an algorithm are calculated using (3.7).

4.1.2 Feature engineering for fraud detection

When constructing a credit card fraud detection algorithm, the initial set of features (raw features) include information regarding individual transactions. It is observed throughout the literature, that regardless of the study, the set of raw features is quite similar. This is because the data collected during a credit card transaction must comply with international financial reporting standards [American Institute of CPAs, 2011]. In Table 4.2, the typical credit card fraud detection raw features are summarized.

4.1.2.1 Customer spending patterns

Several studies use only the raw features in carrying their analysis [Brause et al., 1999; Minegishi and Niimi, 2011; Panigrahi et al., 2009; Sánchez et al., 2009]. However, as noted in [Bolton and Hand, 2001], a single transaction information is not sufficient to detect a fraudulent transaction, since using only the raw features leaves behind important information such as the consumer spending behavior, which is usually used by commercial fraud detection systems [Whitrow et al., 2008].

Attribute name	Description
Transaction ID	Transaction identification number
Time	Date and time of the transaction
Account number	Identification number of the customer
Card number	Identification of the credit card
Transaction type	ie. Internet, ATM, POS, ...
Entry mode	ie. Chip and pin, magnetic stripe, ...
Amount	Amount of the transaction in Euros
Merchant code	Identification of the merchant type
Merchant group	Merchant group identification
Country	Country of trx
Country 2	Country of residence
Type of card	ie. Visa debit, Mastercard, American Express...
Gender	Gender of the card holder
Age	Card holder age
Bank	Issuer bank of the card

Table 4.2: Summary of typical raw credit card fraud detection features

To deal with this, in [Bachmayer, 2008], a new set of features were proposed such that the information of the last transaction made with the same credit card is also used to make a prediction. The objective, is to be able to detect very dissimilar continuous transactions within the purchases of a customer. The new set of features include: time since the last transaction, previous amount of the transaction, previous country of the transaction. Nevertheless, these features do not take into account consumer behavior other than the last transaction made by a client, this leads to having an incomplete profile of customers.

A more compressive way to take into account a customer spending behavior is to derive some features using a transaction aggregation strategy. This methodology was initially proposed in [Whitrow et al., 2008]. The derivation of the aggregation features consists in grouping the transactions made during the last given number of hours, first by card or account number, then by transaction type, merchant group, country or other, followed by calculating the number of transactions or the total amount spent on those transactions. This methodology has been used by a number of studies [Bhattacharyya et al., 2011; Weston et al., 2008; Tasoulis and Adams, 2008; Correa Bahnsen et al., 2013; Sahin et al., 2013; Correa Bahnsen et al., 2014b; Pozzolo et al., 2014a].

When aggregating a customer transactions, there is an important question on how much to accumulate, in the sense that the marginal value of new information may diminish as time passes. Whitrow et al. [2008] discuss that aggregating 101 transactions is not likely to be

more informative than aggregating 100 transactions. Indeed, when time passes, information lose their value, in the sense that a customer spending patterns are not expected to remain constant over the years. In particular, Whitrow et al. define a fixed time frame to be 24, 60 or 168 hours.

The process of aggregating features consists in selecting those transactions that were made in the previous t_p hours, for each transaction i in the dataset \mathcal{S} ,

$$\mathcal{S}_{agg} \equiv \text{TRX}_{agg}(\mathcal{S}, i, t_p) = \left\{ x_l^{amt} \mid \left(x_l^{id} = x_i^{id} \right) \wedge \left(\text{hours}(x_i^{time}, x_l^{time}) < t_p \right) \right\}_{l=1}^N, \quad (4.2)$$

where $N = |\mathcal{S}|$, $|\cdot|$ being the cardinality of a set, x_i^{time} is the time of transaction i , x_i^{amt} is the amount of transaction i , x_i^{id} the customer identification number of transaction i , and $\text{hours}(t_1, t_2)$ is a function that calculates the number of hours between the times t_1 and t_2 . Afterwards the feature number of transactions and amount of transactions in the last t_p hours are calculated as:

$$x_i^{a1} = |\mathcal{S}_{agg}|, \quad (4.3)$$

and

$$x_i^{a2} = \sum_{x^{amt} \in \mathcal{S}_{agg}} x^{amt}, \quad (4.4)$$

respectively.

We note that this aggregation is not enough, in the sense that the combination of different features is not being taken into account. For example, it is not only interesting to see the total transactions, but also group them following a certain criteria, such as: transactions made in the last t_p hours, in the same country and of the same transaction type. For calculating such features, first we expand (4.2) as follows

$$\mathcal{S}_{agg2} \equiv \text{TRX}_{agg}(\mathcal{S}, i, t_p, \text{cond}_1, \text{cond}_2) = \left\{ x_l^{amt} \mid \left(x_l^{id} = x_i^{id} \right) \wedge \left(\text{hours}(x_i^{time}, x_l^{time}) < t_p \right) \wedge \left(x_l^{\text{cond}_1} = x_i^{\text{cond}_1} \right) \wedge \left(x_l^{\text{cond}_2} = x_i^{\text{cond}_2} \right) \right\}_{l=1}^N. \quad (4.5)$$

where, cond_1 and cond_2 , could be either of the features of a transaction listed in Table 4.2. Then, the features are calculated as:

$$x_i^{a3} = |\mathcal{S}_{agg2}|, \quad (4.6)$$

Raw features					Aggregated features			
Id	Time	Type	Country	Amt	x_i^{a1}	x_i^{a2}	x_i^{a3}	x_i^{a4}
1	01/01 18:20	POS	LUX	250	0	0	0	0
2	01/01 20:45	POS	LUX	400	1	250	1	250
3	01/01 22:40	ATM	LUX	250	2	650	0	0
4	02/01 00:50	POS	GER	50	3	900	0	0
5	02/01 19:18	POS	GER	100	3	700	1	50
6	02/01 23:45	POS	GER	150	2	150	2	150
7	03/01 06:00	POS	LUX	10	3	400	0	0

Table 4.3: Example calculation of aggregated features. For all aggregated features $t_p = 24$.

and

$$x_i^{a4} = \sum_{x^{amt} \in \mathcal{S}_{agg2}} x^{amt}. \quad (4.7)$$

To further clarify how the aggregated features are calculated we show an example. Consider a set of transactions made by a client between the first and third of January of 2015, as shown in Table 4.3. Then we estimate the aggregated features (x_i^{a1} , x_i^{a2} , x_i^{a3} and x_i^{a4}) by setting $t_p = 24$ hours. The different aggregated features give us different information of the customer spending behavior. Moreover, the total number of aggregated features can grow quite quickly, as t_p can have several values, and the combination of combination criteria can be quite large as well. In this thesis, we use a total of 280 aggregated features. In particular we set the different values of t_p to: 1, 3, 6, 12, 18, 24, 72 and 168 hours. Then calculate the aggregated features using (4.2), and also using (4.5) with the following grouping criteria: country, type of transaction, entry mode, merchant code and merchant group.

4.1.2.2 Time features

When using the aggregated features, there is still some information that is not completely captured by those features. In particular we are interested in analyzing the time of the transaction. The logic behind this, is that a customer is expected to make transactions at similar hours. The issue when dealing with the time of the transaction, specifically, when analyzing a feature such as the mean of transactions time, is that it is easy to make the mistake of using the arithmetic mean. Indeed, the arithmetic mean is not a correct way to average time because, as shown in Figure 4.1, it does not take into account the periodic behavior of the time feature. For example, the arithmetic mean of transaction time of

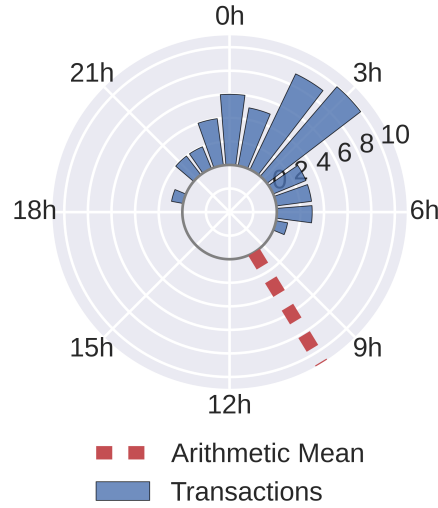


Figure 4.1: Analysis of the time of a transaction using a 24 hour clock. The arithmetic mean of the transactions time (dashed line) do not accurately represents the actual times distribution.

four transactions made at 2:00, 3:00, 22:00 and 23:00 is 12:30, which is counter intuitive since no transaction was made close to that time.

We propose to overcome this limitation by modeling the time of the transaction as a periodic variable, in particular using the von Mises distribution [Fisher, 1996]. The von Mises distribution, also known as the periodic normal distribution, is a distribution of a wrapped normal distributed variable across a circle. The von Mises probability distribution of a set of examples $D = \{t_1, t_2, \dots, t_N\}$ for a given angle t_x is given by

$$f(t_x | \mu_{vM}, \sigma_{vM}) = \frac{e^{\frac{1}{\sigma_{vM}} \cos(t_x - \mu_{vM})}}{2\pi I_0\left(\frac{1}{\sigma_{vM}}\right)} \quad (4.8)$$

where $I_0(\kappa)$ is the modified Bessel function of order 0, and μ_{vM} , σ_{vM} are the periodic mean and periodic standard deviation, respectively, and are defined as

$$\mu_{vM}(D) = 2 \cdot \arctan\left(\frac{\phi}{\left(\sqrt{\psi^2 + \phi^2} + \psi\right)}\right), \quad (4.9)$$

and

$$\sigma_{vM}(D) = \sqrt{\ln\left(\frac{1}{\left(\frac{\phi}{N}\right)^2 + \left(\frac{\psi}{N}\right)^2}\right)}, \quad (4.10)$$




Figure 4.2: Fitted von Mises distribution including the periodic mean (dashed line) and the probability distribution (purple area).

respectively. Where $\phi = \sum_{t_j \in D} \sin(t_j)$ and $\psi = \sum_{t_j \in D} \cos(t_j)$ [Bishop, 2006].

In particular we are interested in calculating a confidence interval (CI) for the time of a transaction. For doing that, initially we select a set of transactions made by the same client in the last t_p hours,

$$S_{\text{per}} \equiv \text{TRX}_{\text{vM}}(S, i, t_p) = \left\{ x_l^{\text{time}} \mid \left(x_l^{\text{id}} = x_i^{\text{id}} \right) \wedge \left(\text{hours}(x_i^{\text{time}}, x_l^{\text{time}}) < t_p \right) \right\}_{l=1}^N. \quad (4.11)$$

Afterwards, the probability distribution function of the time of the set of transactions is calculated as:

$$x_i^{\text{time}} \sim \text{vonmises} \left(\mu_{\text{vM}}(S_{\text{per}}), \frac{1}{\sigma_{\text{vM}}(S_{\text{per}})} \right). \quad (4.12)$$

In Figure 4.2, the von Mises distribution calculation for the earlier example is shown. It is observed that the arithmetic mean is quite different from the periodic mean, the latter being a more realistic representation of the actual transactional times. Then, using the estimated distribution, a new set of features can be extracted, ie., a binary feature (x_i^{p1}) if a new transaction time is within the confidence interval range with probability α . An example is presented in Figure 4.3. Furthermore, other features can be calculated, as the confidence interval range can be calculated for several values of α , and also the time period can have an arbitrary size.

Additionally, following the same example presented in Table 4.3, we calculate a feature x_i^{p1} , as a binary feature that takes the value of one if




Figure 4.3: Expected time of a transaction (green area). Using the confidence interval, a transaction can be flag normal or suspicious, depending whether or not the time of the transaction is within the confidence interval.

Raw features		Arithmetic	Periodic features		
Id	Time	mean	mean	Confidence interval	x_t^{p1}
1	01/01 18:20	—	—	—	—
2	01/01 20:45	—	—	—	—
3	01/01 22:40	19:27	19:27	15:45 - 23:10	True
4	02/01 00:50	20:28	20:28	17:54 - 23:03	False
5	02/01 19:18	16:44	22:44	18:51 - 00:17	True
6	02/01 23:45	16:19	21:07	15:21 - 02:52	True
7	03/01 06:00	18:43	22:43	17:19 - 01:46	False

Table 4.4: Example calculation of periodic features.

the current time of the transaction is within the confidence interval of the time of the previous transactions with a confidence of $\alpha = 0.9$. The example is shown in Table 4.4, where the arithmetic and periodic means differ, as for the last transaction both means are significantly different. Moreover, the new feature helps to get a better understanding of when a customer is expected to make transactions.

Finally, when calculating the periodic features, it is important to use longer time frames t_p , since if the distribution is calculated using only a couple of transactions it may not be as relevant of a customer behavior patterns, compared against using a full year of transactions. Evidently, if t_c is less than 24 hours, any transaction made afterwards will not be expected to be within the distribution of previous transactional times.

To avoid this, we recommend using at least the previous 7 days of transactional information, therefore, having a better understanding of its behavioral patterns. Lastly, this approach can also be used to estimate features such as the expected day of the week of transactions, as some customers may only use their credit cards during the weekend nights, or during working hours.

4.1.3 Database

For this thesis we use a dataset provided by a large European card processing company. The dataset consists of fraudulent and legitimate transactions made with credit and debit cards between January 2012 and June 2013. The total dataset contains 120,000,000 individual transactions, each one with 27 attributes, including a fraud label indicating whenever a transaction is identified as fraud. This label was created internally in the card processing company, and can be regarded as highly accurate. In the dataset only 40,000 transactions were labeled as fraud, leading to a fraud ratio of 0.025%.

4.2 CREDIT SCORING

In this section, we consider the second financial risk management application, namely, credit scoring. In order to mitigate the impact of credit risk and make more objective and accurate decisions, financial institutions use credit scores to predict and control their losses. The objective in credit scoring is to classify which potential customers are likely to default a contracted financial obligation based on the customer's past financial experience, and with that information decide whether to approve or decline a loan [Anderson, 2007]. This tool has become a standard practice among financial institutions around the world in order to predict and control their loans portfolios. When constructing credit scores, it is a common practice to use standard cost-insensitive binary classification algorithms such as logistic regression, neural networks, discriminant analysis, genetic programming, decision trees, among others [Hand and Henley, 1997; Correa Bahnsen and Gonzalez Montoya, 2011].

Formally, a credit score is a statistical model that allows the estimation of the probability $\hat{p}_i = P(y_i = 1|x_i)$ of a customer i defaulting a contracted debt. Additionally, since the objective of credit scoring is to estimate a classifier c_i to decide whether or not to grant a loan to a customer i , a threshold t is defined such that if $\hat{p}_i < t$, then the loan is granted, i.e., $c_i(t) = 0$, and denied otherwise, i.e., $c_i(t) = 1$.




Figure 4.4: Credit scoring sensitivity versus specificity thresholding procedure.

There exists different approaches for defining the probability threshold. The sensitivity versus specificity (SvsS) approach is the most widely used among financial institutions [Anderson, 2007], where specificity is the true positive rate $F_0(t)$ for a threshold t , and the sensitivity is one minus the false positive rate $F_1(t)$ given a threshold t [Hernandez-Orallo et al., 2012]. In this method the objective is to fix the threshold at the point where the sensitivity is equal to the specificity $F_0(t) = 1 - F_1(t)$, where $F_0(t)$ and $F_1(t)$ are calculated using:

$$F_a(t) = \frac{1}{N_a} |\{x_i | x_i \in \mathcal{S}_a \wedge \hat{p}_i \leq t\}|, \text{ for } a \in \{0, 1\}. \quad (4.13)$$

Lastly, the SvsS threshold t_{SvsS} is found by using

$$t_{\text{SvsS}} = \arg \min_t |F_0(t) - (1 - F_1(t))|. \quad (4.14)$$

This process is further clarified in Figure 4.4.

After the classifier c_i is estimated, there is a need to evaluate its performance. In practice, many statistical evaluation measures are used to assess the performance of a credit scoring model. Measures such as the area under the receiver operating characteristic curve (AUC), Brier score, Kolmogorov-Smirnoff (K-S) statistic, F_1 -Score, and misclassification are among the most common [Beling et al., 2005]. Nevertheless, none of these measures takes into account the business and economical realities that take place in credit scoring. Costs that the financial institution had incurred to acquire customers, or the expected profit due to

a particular client, are not considered in the evaluation of the different models. This is explored in detail in the next section.

4.2.1 *Financial evaluation of a credit scorecard*

Typically, a credit risk model is evaluated using standard cost-insensitive measures. However, in practice, the cost associated with approving what is known as a bad customer, i.e., a customer who defaulted his credit loan, is quite different from the cost associated with declining a good customer, i.e., a customer who successfully repaid his credit loan. Furthermore, the costs are not constant among customers. This is because loans have different credit line amounts, terms, and even interest rates. Some authors have proposed methods that include the misclassification costs in the credit scoring context [Verbraken et al., 2014; Alejo and Garc, 2013; Beling et al., 2005; Oliver and Thomas, 2009]. However, they assume a constant misclassification cost, which is not the case in credit scoring.

Initial approaches to include the different costs have been published in recent years, particularly the one proposed by Beling et al. [Beling et al., 2005; Oliver and Thomas, 2009], in which the costs of misclassification are assigned for each error. Specifically, setting the cost of a false positive C_{FP} to the loan's annual interest rate charged to the customer int_r , the cost of a false negative C_{FN} to the loss given default L_{gd} , which is the percentage of loss over the total credit line when the customer defaulted, and setting to zero the costs of true positive C_{TP} and true negative C_{TN} . Using that, they proposed the expected cost (EC) method to find the probability threshold that minimizes those costs, $t_{ec} = \frac{C_{FN}}{C_{FN} + C_{FP}} = \frac{L_{gd}}{L_{gd} + \text{int}_r}$. Nevertheless, this approach assumes a constant cost within examples, which is a strong assumption, since in practice each example carries a very different cost given by the different credit limits and conditions of each loan. Consequently, there is a need for an example-dependent cost matrix that takes into account the cost of misclassifying each example.

4.2.1.1 *Example-dependent cost-sensitive evaluation measure*

In order to take into account the varying costs that each example carries, we propose a cost matrix with example-dependent misclassification costs as given in Table 4.5. First, we assume that the costs of a correct classification, C_{TP_i} and C_{TN_i} , are zero for every customer i . We define C_{FN_i} the losses if the customer i defaults to be proportional to his credit line Cl_i . We define the cost of a false positive per customer

	Actual Positive $y_i = 1$	Actual Negative $y_i = 0$
Predicted Positive $c_i = 1$	$C_{TP_i} = 0$	$C_{FP_i} = r_i + C_{FP}^a$
Predicted Negative $c_i = 0$	$C_{FN_i} = Cl_i \cdot L_{gd}$	$C_{TN_i} = 0$

Table 4.5: Credit scoring example-dependent cost matrix

C_{FP_i} as the sum of two real financial costs r_i and C_{FP}^a , where r_i is the loss in profit by rejecting what would have been a good customer.

The profit per customer r_i is calculated as the present value of the difference between the financial institution gains and expenses, given the credit line Cl_i , the term l_i and the financial institution lending rate int_{r_i} for customer i , and the financial institution of cost funds int_{cf} .

$$r_i = PV(A(Cl_i, int_{r_i}, l_i), int_{cf}, l_i) - Cl_i, \quad (4.15)$$

with A being the customer monthly payment and PV the present value of the monthly payments, which are calculated using the time value of money equations [Lawrence and Solomon, 2012],

$$A(Cl_i, int_{r_i}, l_i) = Cl_i \frac{int_{r_i}(1 + int_{r_i})^{l_i}}{(1 + int_{r_i})^{l_i} - 1}, \quad (4.16)$$

and

$$PV(A, int_{cf}, l_i) = \frac{A}{int_{cf}} \left(1 - \frac{1}{(1 + int_{cf})^{l_i}} \right). \quad (4.17)$$

The second term C_{FP}^a is related to the assumption that the financial institution will not keep the money of the declined customer idle. It will instead give a loan to an alternative customer [Nayak and Turvey, 1997]. Since no further information is known about the alternative customer, it is assumed to have an average credit line \overline{Cl} and an average profit \bar{r} . The false positive cost for an alternative customer becomes

$$C_{FP}^a = -\bar{r} \cdot \pi_0 + \overline{Cl} \cdot L_{gd} \cdot \pi_1 \quad (4.18)$$

In other words, it is minus the profit of an average alternative customer plus the expected loss, taking into account that the alternative customer will pay his debt with a probability equal to the prior negative rate, and similarly will default with probability equal to the prior positive rate.

4.2.1.2 Calculation of the credit limit

One key parameter of our model is the credit limit. There exists several strategies to calculate the Cl_i depending on the type of loans, the

state of the economy, the current portfolio, among others [Anderson, 2007; Lawrence and Solomon, 2012]. Nevertheless, given the lack of information regarding the specific business environments of the considered datasets, we simply define Cl_i as

$$Cl_i = \min \left\{ q \cdot Inc_i, Cl_{\max}, Cl_{\max}(debt_i) \right\}, \quad (4.19)$$

where Inc_i and $debt_i$ are the monthly income and debt ratio of the customer i , respectively, q is a parameter that defines the maximum Cl_i as a function of the Inc_i , and Cl_{\max} the maximum overall credit line. Lastly, the maximum credit line given the current debt is calculated as the maximum credit limit such that the current debt ratio plus the new monthly payment does not surpass the customer monthly income. It is calculated as

$$Cl_{\max}(debt_i) = PV(Inc_i \cdot P_m(debt_i), int_{r_i}, l_i), \quad (4.20)$$

and

$$P_m(debt_i) = \min \left\{ \frac{A(q \cdot Inc_i, int_{r_i}, l_i)}{Inc_i}, (1 - debt_i) \right\}. \quad (4.21)$$

4.2.2 Databases

For this thesis we use two different publicly available credit scoring datasets. The first dataset is the **2011 Kaggle competition Give Me Some Credit¹**, in which the objective is to identify those customers of personal loans that will experience financial distress in the next two years. The second dataset is from the **2009 Pacific-Asia Knowledge Discovery and Data Mining conference (PAKDD) competition²**. Similarly, this competition had the objective of identifying which credit card applicants were likely to default and by doing so deciding whether or not to approve their applications. The Kaggle Credit and PAKDD Credit datasets contain information regarding the features, and more importantly about the income of each example, from which an estimated credit limit Cl_i can be calculated.

The Kaggle Credit dataset contains 112,915 examples, each one with 10 features and the class label. The proportion of default or positive examples is 6.74%. On the other hand, the PAKDD Credit dataset contains 38,969 examples, with 30 features and the class label, with a proportion of 19.88% positives. This database comes from a Brazilian financial institution, and as it can be inferred from the competition description, the data was obtained around 2004.

¹ <http://www.kaggle.com/c/GiveMeSomeCredit/>

² <http://sede.neurotech.com.br:443/PAKDD2009/>

Parameter	Kaggle Credit	PAKDD Credit
Interest rate (int_r)	4.79%	63.0%
Cost of funds (int_{cf})	2.94%	16.5%
Term (l) in months	24	24
Loss given default (L_{gd})	75%	75%
Times income (q)	3	3
Maximum credit line (Cl_{max})	25,000	25,000

Table 4.6: Credit scoring model parameters

Since no specific information regarding the datasets is provided, we assume that they belong to average European and Brazilian financial institutions. This enable us to find the different parameters needed to calculate the cost measure. In Table 4.6, the different parameters are shown. In particular, we obtained the average interest rates in Europe during 2013 from the European Central Bank [ECB, 2014], and the average interest and exchange rates in Brazil during 2004 from Trading Economics [Economics, 2014]. Because the income is not in the same currency on both datasets, we convert the PAKDD Credit dataset to Euros. Additionally, we use a fixed loan term l for both datasets, considering that in the Kaggle Credit dataset the class was constructed to predict two years of credit behavior, and because the PAKDD Credit dataset is related to credit cards the term is fixed to two years [Lawrence and Solomon, 2012]. Moreover, we set the loss given default L_{gd} using information from the Basel II standard³, q to 3 since it is the average personal loan requests related to monthly income, and the maximum credit limit Cl_{max} to 25,000 Euros.

4.3 SUMMARY OF THE DATASETS

For each dataset we used a pre-defined cost matrix as shown in Section 4.1 and Section 4.2. Additionally, for each database, three different datasets are extracted: training, validation and testing, each one containing 50%, 25% and 25% of the examples, respectively. Afterwards, because classification algorithms suffer when the label distribution is skewed towards one of the classes [Hastie et al., 2009], an under-sampling of the positive examples is made, in order to have a balanced class distribution. Moreover, we perform the cost-proportionate rejection-sampling and cost proportionate over-sampling procedures, that we previously described in Section 3.3.3. Table 4.7, summarizes the different datasets. It is important to note that the sampling procedures

³ <http://www.bis.org/publ/bcbsca.htm>.

were only applied to the training dataset since the validation and test datasets must reflect the real distribution.

Database	Set	N	π_1	Cost (Euros)
Fraud	Total	236,735	1.50	895,154
Detection	Training (t)	94,599	1.51	358,078
	Under-sampled (u)	2,828	50.42	358,078
	Rejection-sampled (r)	94,522	1.43	357,927
	Over-sampled (o)	189,115	1.46	716,006
	Validation	70,910	1.53	274,910
	Testing	71,226	1.45	262,167
Credit	Total	112,915	6.74	83,740,181
Scoring 1	Training (t)	45,264	6.75	33,360,130
	Under-sampled (u)	6,038	50.58	33,360,130
	Rejection-sampled (r)	5,271	43.81	29,009,564
	Over-sampled (o)	66,123	36.16	296,515,655
	Validation	33,919	6.68	24,786,997
	Testing	33,732	6.81	25,593,055
Credit	Total	38,969	19.88	3,117,960
Scoring 2	Training (t)	15,353	19.97	1,221,174
	Under-sampled (u)	6,188	49.56	1,221,174
	Rejection-sampled (r)	2,776	35.77	631,595
	Over-sampled (o)	33,805	33.93	6,798,282
	Validation	11,833	20.36	991,795
	Testing	11,783	19.30	904,991

Table 4.7: Summary of the financial datasets, where N is the number of examples and π_1 is the percentage of positive examples.

OUTLINE

In this chapter, we present two different marketing real-world example-dependent cost-sensitive problems, namely, churn modeling and direct marketing. Both problems deal with identifying those customers with certain characteristic with the objective to maximize the results of the different CRM strategies. First, we introduce the churn modeling problem and present the proposed financial evaluation measure for a churn campaign in Section 5.1. Lastly in Section 5.2, the direct marketing problem is presented including the proposed example-dependent cost-sensitive cost matrix for evaluating this problem.

5.1 CHURN MODELING

Customer churn predictive modeling deals with predicting the probability of a customer defecting using historical, behavioral and socio-economical information. This tool is of great benefit to subscription based companies allowing them to maximize the results of retention campaigns. The problem of churn predictive modeling has been widely studied by the data mining and machine learning communities. It is usually tackled by using classification algorithms in order to learn the different patterns of both the churners and non-churners. Nevertheless, current state-of-the-art classification algorithms are not well aligned with commercial goals, in the sense that, the models miss to include the real financial costs and benefits during the training and evaluation phases. In the case of churn, evaluating a model based on a traditional measure such as accuracy or predictive power, does not yield to the best results when measured by the actual financial cost, i.e., investment per subscriber on a loyalty campaign and the financial impact of failing to detect a real churner versus wrongly predicting a non-churner as a churner.

In this section, we propose a cost-sensitive framework for customer churn predictive modeling. First, in Section 5.1.1, we introduce the problem of churn modeling. Then in Section 5.1.2, we present a financial based measure for evaluating the effectiveness of a churn campaign taking into account the available portfolio of offers, their individual financial cost and probability of offer acceptance depending on the customer

profile. Finally, in Section 5.1.3, we describe the real-world churn modeling dataset that will be used during the experiments.

5.1.1 *Flow analysis of a churn campaign*

The two main objectives of subscription-based companies are to acquire new subscribers and retain those they already have, mainly because profits are directly linked with the number of subscribers. In order to maximize the profit, companies must increase the customer base by incrementing sales while decreasing the number of churners. Furthermore, it is common knowledge that retaining a customer is about five times less expensive than acquiring a new one [Farris et al., 2010], this creates pressure to have better and more effective churn campaigns.

A typical churn campaign consists in identifying from the current customer base which ones are more likely to leave the company, and make an offer in order to avoid that behavior. With this in mind the companies use intelligence to create and improve retention and collection strategies. In the first case, this usually implies an offer that can be either a discount or a free upgrade during certain span of time. In both cases the company has to assume a cost for that offer, therefore, accurate prediction of the churners becomes important. The logic of this flow is shown in Figure 5.1.

The churn campaign process starts with the sales that every month increase the customer base, however, monthly there is a group of customers that decide to leave the company for many reasons. Then the objective of a churn model is to identify those customers before they take the decision of defecting.

Using a churn model, those customers more likely to leave are predicted as churners and an offer is made in order to retain them. However, it is known that not all customers will accept the offer, in the case when a customer is planning to defect, it is possible that the offer is not good enough to retain him or that the reason for defecting can not be influenced by an offer. Using historical information, it is estimated that a customer will accept the offer with probability γ . On the other hand, there is the case in which the churn model misclassified a non-churner as churner, also known as false positives, in that case the customer will always accept the offer that means an additional cost to the company since those misclassified customers do not have the intentions of leaving.

In the case where the churn model predicts customers as non-churners, there is also the possibility of a misclassification, in this case an actual churner is predicted as non-churner, since these customers do not receive an offer and they will leave the company, these cases are known




Figure 5.1: Flow analysis of a churn campaign [Verbraken, 2012]

as false negatives. Lastly, there is the case where the customers are actually non-churners, then there is no need to make a retention offer to these customers since they will continue to be part of the customer base.

It can be seen that a churn campaign (or churn model) has three main points. First, avoid false positives since there is a financial cost of making an offer where it is not needed. Second, find the right offer to give to those customers identified as churners. And lastly, to decrease the number of false negatives.

From a machine learning perspective, a churn model is a classification algorithm. In the sense that using historical information, a prediction of which current customers are more likely to defect, is made. This model is normally created using one of a number of well established algorithms (Logistic regression, neural networks, random forests, among others) [Ngai et al., 2009; KhakAbi et al., 2010]. Then, the model is evaluated using measures such as misclassification error, receiver operating characteristic (ROC), Kolmogorov-Smirnov (KS) or F_1 Score statistics [Verbeke et al., 2012]. However, these measures may not be the most appropriate evaluation criteria when evaluating a churn model, because they tacitly assume that misclassification errors carry the same cost, similarly with the correct classified examples. This assumption does not hold in many real-world applications such as churn modeling, since when misidentifying a churner the financial losses are quite different than when misclassifying a non-churner as churner [Glady et al., 2009]. Furthermore, the accuracy measure also assumes that the class distribution among examples is constant and balanced [Provost et al., 1998], and typically the distributions of a churn data set are skewed [Verbeke et al., 2012].

In the next section, we propose a new financial based measure for evaluating the effectiveness of a voluntary churn campaign taking into account the available portfolio of offers, their individual financial cost and probability of acceptance depending on the customer profile.

5.1.2 Propose evaluation measure of a churn campaign

Different studies have proposed measures to deal with these cost-sensitivity related to evaluating a churn model. In [Neslin et al., 2006], a profit-based measure was proposed by starting with the confusion matrix and multiplying it with the expected profit of each case.

$$\text{Profit}_1 = (\text{TP} + \text{FP}) \left[(\gamma \text{CLV} + C_o(1 - \gamma)(-C_a)) \pi_1 \gamma - C_o - C_a \right] - N \cdot C_a, \quad (5.1)$$

with C_a being the fixed administrative cost of running the campaign, C_o the average cost of the retention offer, C_a the cost of contacting the customer, π_1 the prior churn rate and CLV the average customer lifetime value or the present value of the expected profit that a customer will generate. Moreover, as discussed in [Verbraken et al., 2013], if the average instead of the total profit is considered and the fixed cost $N \cdot C_a$ is discarded since it is irrelevant for the classifier selection, the profit can be expressed as:

$$\text{Profit}_2 = \text{TP} (\gamma(\text{CLV} - C_o - C_a) + (1 - \gamma)C_a) + \text{FP}(-C_o - C_a). \quad (5.2)$$

Nevertheless, equations (5.1) and (5.2) assume that every customer has the same CLV and C_o , whereas this is not true in practice. In fact, different customers have a very different CLV, and not all offers can be made to every customer, neither do they have the same impact across customers. In order to obtain a more business oriented measure, we first analyze the financial impact of the different decisions, i.e., false positives, false negatives, true positives and true negatives, for each customer.

In Figure 5.2, the financial impact of a churn model is shown. Note that we take into account the costs and not the profit in each case. When a customer is predicted to be a churning, an offer is made with the objective of avoiding the customer defecting. However, if a customer is actually a churning, he may or not accept the offer with a probability γ_i . If the customer accepts the offer, the financial impact is equal to the cost of the offer (C_{oi}) plus the administrative cost of contacting the customer (C_a). On the other hand, if the customer declines the offer, the cost is the expected income that the clients would otherwise generate, also called customer lifetime value (CLV_i), plus C_a . Lastly, if the customer is not actually a churning, he will be happy to accept the offer and the cost will be C_{oi} plus C_a .

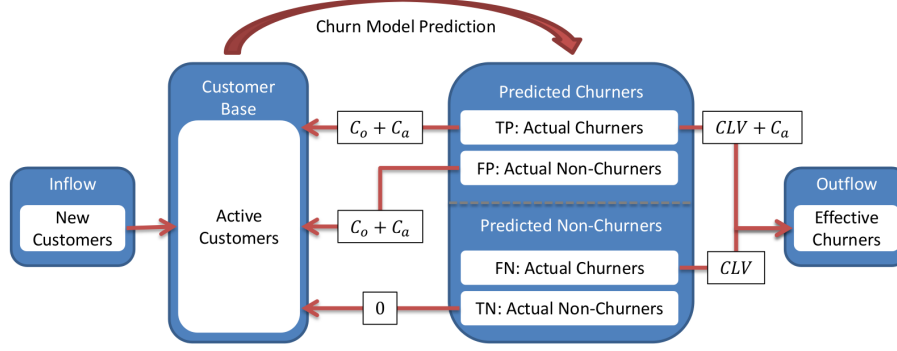


Figure 5.2: Financial impact of the different decisions, i.e., False positives, false negatives, true positives and true negatives

In the case that the customer is predicted as non-churner, there are two possible outcomes. Either the customer is not a churner, then the cost is zero, or the customer is a churner and the cost is CLV_i .

The different costs are summarized in Table 5.1. Then using the cost matrix, and the example-dependent cost-sensitive framework as described in Section 3.3.1, an example-dependent cost statistic is calculated as:

$$\begin{aligned}
 Cost_i &= y_i(c_i C_{TP_i} + (1 - c_i) C_{FN_i}) \\
 &\quad + (1 - y_i)(c_i C_{FP_i} + (1 - c_i) C_{TN_i}) \\
 &= y_i(c_i (\gamma_i (C_{o_i} - CLV_i - C_a) - C_{o_i}) + CLV_i) \\
 &\quad + c_i (C_{o_i} + C_a),
 \end{aligned} \tag{5.3}$$

leading to a total cost of:

$$Cost = \sum_{i=1}^N Cost_i. \tag{5.4}$$

Furthermore, using (3.7), the savings are calculated as:

$$Savings = \frac{Cost_l - Cost}{Cost_l}, \tag{5.5}$$

	Actual Positive $y_i = 1$	Actual Negative $y_i = 0$
Predicted Positive $c_i = 1$	$C_{TP_i} = \gamma_i C_{o_i} + (1 - \gamma_i)(CLV_i + C_a)$	$C_{FP_i} = C_{o_i} + C_a$
Predicted Negative $c_i = 0$	$C_{FN_i} = CLV_i$	$C_{TN_i} = 0$

Table 5.1: Proposed churn modeling example-dependent cost matrix

In almost all cases the costless class ($Cost_l$) will be the negative class, as typically the distribution of a churn dataset is skewed towards the non-churners [Verbeke et al., 2012]. Given that $Cost_l$ can be expressed as $Cost(f_0)$, or simply $Cost$ with $c_i = 0 \forall i$:

$$Cost_l = \sum_{i=1}^N y_i CLV_i. \quad (5.6)$$

This is consistent with the notion that if no model is used, the total cost would be the sum of the customer lifetime values of the actual churners, which gives the insight that the Savings measure consists in comparing the financial impact of the campaign of using a classification model against not using a model at all.

5.1.2.1 Customer lifetime value

Lastly, one of the key values to calculate the Savings is the customer lifetime value. Within marketing there exists a common misconception between customer profitability and customer lifetime value. The two terms are usually used in an interchangeable way, creating confusion of what the actual objective of a churn modeling campaign should be. Several studies have proposed models providing a unique definition of both terms [Neslin et al., 2006; Pfeifer et al., 2004; Milne and Boza, 1999; van Raaij et al., 2003]. Customer profitability indicates the difference between the income and the cost generated by a customer i during a financial period t . It is defined as:

$$CP_{i,t} = \mu \cdot s_{i,t}, \quad (5.7)$$

where $s_{i,t}$ refers to the consumption of customer i during time period t , and μ refers to the average marginal profit by unit product usage.

Moreover, we are interested to see what is the expected income that a particular customer will generate in the future, in other words, calculating the expected sum of discount future earnings [Neslin et al., 2006]. Therefore, the CLV_i is defined as:

$$CLV_i = \sum_{t=1}^T \frac{\mu \cdot s_{i,t}}{(1+r)^t}, \quad (5.8)$$

where r is the discount rate, and T the number of time period. Typically T should be considered large enough since without prior knowledge a customer is expected to keep being a customer for the foreseeable future. In practice T is set up to be infinity [Glady et al., 2009]. Also, for simplicity, it can be assumed that $s_{i,t+1} = s_{i,t} \cdot (1+g) \forall i, t$, which means

that there is a constant growth g in the customer consumption. Given that, the customer lifetime value can be re-written as

$$CLV_i = \sum_{t=1}^{\infty} \frac{(1+g)^t}{(1+r)^t} \cdot \mu \cdot s_{i,1}, \quad (5.9)$$

which in the case of $g < r$, this is a geometric series and can be expressed as

$$CLV_i = \frac{\mu \cdot s_{i,1}}{(r-g)}. \quad (5.10)$$

5.1.3 Churn modeling database

For our experiments we use a dataset provided by a TV cable provider. The dataset consists of active customers during the first semester of 2014. The total dataset contains 9,410 individual registries, each one with 45 attributes, including a churn label indicating whenever a customer is a churning. This label was created internally in the company, and can be regarded as highly accurate. In the dataset only 455 customers are churners, leading to a churn ratio of 4.83%.

5.1.3.1 Offer acceptance calculation

In practice companies have a set of offers to make to a customer as part of the retention campaign. They vary from discounts to upgrades, among others. In the particular case of a TV cable provider, the offers include adding a new set of channels, changing the TV receiver to one with new technology (i.e., high definition, video recording, 4K), or to offer a discount on the monthly bill. Unsurprisingly, not all offers apply to all clients. For instance, a customer that already has all the channels can not be offered a new set of channels. Moreover, an offer usually means an additional cost to the company and not all offers have the same cost or the same impact in reducing churn.

Taking into account the cost and the implication of the offers, the problem can be summarized in making each customer the offer that will maximize the acceptance rate and more importantly reduce the overall cost.

In order to calculate the acceptance probability γ_i a champion-challenger process was made. First, the customers were grouped into clusters according to their behavioral and socio-economical characteristics. In particular the K-means algorithm was used [Marlsan, 2009]. Then for a period of two months, randomly selected offers were made to the customers and their response was evaluated. Unfortunately, for confidentiality reasons we can not describe the different clusters, or the




Figure 5.3: Acceptance rate (γ) of the best offer for each customer profile. As expected, the higher the churn rate the lower the acceptance rate, as it is more difficult to make a good offer to a customer which is more likely to defect.

actual offer made to each customer. Nevertheless, in Figure 5.3, the average churn rate and acceptance rate γ_i per cluster is shown. As expected, the higher the churn rate the lower the acceptance rate, as it is more difficult to make a good offer to a customer that is more likely to defect.

5.2 DIRECT MARKETING

In direct marketing the objective is to classify those customers who are more likely to have a certain response to a marketing campaign [Ngai et al., 2009]. We used a direct marketing dataset [Moro et al., 2011] available on the UCI machine learning repository [Bache and Lichman, 2013]. The dataset contains 45,000 clients of a Portuguese bank who were contacted by phone between March 2008 and October 2010 and received an offer to open a long-term deposit account with attractive interest rates. The dataset contains features such as age, job, marital status, education, average yearly balance and current loan status and the label indicating whether or not the client accepted the offer.

This problem is example-dependent cost sensitive, since there are different costs of false positives and false negatives. Specifically, in direct marketing, false positives have the cost of contacting the client, and false negatives have the cost due to the loss of income by failing to contact a client that otherwise would have opened a long-term deposit.

We propose a direct marketing example-dependent cost matrix as shown in Table 5.2. Where C_a is the administrative cost of contacting the client, as is credit card fraud, and Int_i is the expected income when

	Actual Positive $y_i = 1$	Actual Negative $y_i = 0$
Predicted Positive $c_i = 1$	$C_{TP_i} = C_a$	$C_{FP_i} = C_a$
Predicted Negative $c_i = 0$	$C_{FN_i} = \text{Int}_i$	$C_{TN_i} = 0$

Table 5.2: Direct marketing example-dependent cost matrix

a client opens a long-term deposit. This last term is defined as the long-term deposit amount times the interest rate spread.

In order to estimate Int_i , first the long-term deposit amount is assumed to be a 20% of the average yearly balance, and lastly, the interest rate spread is estimated to be 2.463%, which is the average between 2008 and 2010 of the retail banking sector in Portugal as reported by the Portuguese central bank. Given that, the Int_i is equal to $(\text{balance} * 20\%) * 2.463\%$.

5.3 SUMMARY OF THE DATASETS

In this section we present the different marketing datasets. For each dataset we used a pre-define cost matrix as shown in Section 5.1 and Section 5.2. Moreover, the datasets are split in training, validation and testing, each one containing 50%, 25% and 25% of the examples, respectively. Afterwards, an under-sampling of the positive examples is made, and we perform the cost-proportionate rejection-sampling and cost proportionate over-sampling procedures, as described in Section 3.3.3. Table 5.3, summarizes the different datasets.

Database	Set	N	π_1	Cost (Euros)
Churn	Total	9,410	4.83	580,884
Modeling	Training (t)	3,758	5.05	244,542
	Under-sampled (u)	374	50.80	244,542
	Rejection-sampled (r)	428	41.35	431,428
	Over-sampled (o)	5,767	31.24	2,350,285
	Validation	2,824	4.77	174,171
	Testing	2,825	4.42	162,171
Direct	Total	37,931	12.62	59,507
Marketing	Training (t)	15,346	12.55	24,304
	Under-sampled (u)	3,806	50.60	24,304
	Rejection-sampled (r)	1,644	52.43	20,621
	Over-sampled (o)	22,625	40.69	207,978
	Validation	11,354	12.30	16,154
	Testing	11,231	13.04	19,048

Table 5.3: Summary of the marketing datasets, where N is the number of examples and π_1 is the percentage of positive examples.

Part III

PROPOSED EXAMPLE-DEPENDENT COST-SENSITIVE METHODS

OUTLINE

In this chapter, we propose the cost-sensitive Bayes minimum risk classifier. The method consists in quantifying tradeoffs between various decisions using probabilities and the costs that accompany such decisions. First, in Section 6.1, we present the Bayes minimum risk algorithm. Then, in Section 6.2, we discuss the impact of calibrating the probabilities in the model. Finally, in Section 6.3, using the five real-world cost-sensitive databases, we compare the results of the proposed algorithm, against state-of-the-art methods.

6.1 BAYES MINIMUM RISK MODEL

As defined in [Jayanta K. et al., 2006], the BMR classifier is a decision model based on quantifying tradeoffs between various decisions using probabilities and the costs that accompany such decisions. This is done in a way that for each example the expected losses are minimized. In what follows, we consider the probability estimates \hat{p}_i as known, regardless of the algorithm used to calculate them. The risk that accompanies each decision is calculated using the cost matrix as described in Table 3.1. In the specific framework of binary classification, the risk of predicting the example i as negative is

$$R(c_i = 0|x_i) = C_{TN_i}(1 - \hat{p}_i) + C_{FN_i} \cdot \hat{p}_i, \quad (6.1)$$

and

$$R(c_i = 1|x_i) = C_{TP_i} \cdot \hat{p}_i + C_{FP_i}(1 - \hat{p}_i), \quad (6.2)$$

is the risk when predicting the example as positive, where \hat{p}_i is the estimated positive probability for example i . Subsequently, if

$$R(c_i = 0|x_i) \leq R(c_i = 1|x_i), \quad (6.3)$$

then the example i is classified as negative. This means that the risk associated with the decision c_i is lower than the risk associated with classifying it as positive.

6.2 CALIBRATION OF PROBABILITIES

When using the output of a binary classifier as a basis for decision making, there is a need for a probability that not only separates well between positive and negative examples, but that also assesses the real probability of the event [Cohen and Goldszmidt, 2004].

In this section, two methods for calibrating probabilities are explained. First, the method proposed in [Elkan, 2001] to adjust the probabilities based on the difference in bad rates between the training and testing datasets. Then, the method proposed in Hernandez-Orallo et al. [2012], in which calibrated probabilities are extracted after modifying the ROC curve using the ROC convex hull methodology, is described.

6.2.1 Calibration due to a change in base rates

One of the reasons why a probability may not be calibrated is because the algorithm is trained using a dataset with a different base (or positive) rate than the one on the evaluation dataset. This is something common in machine learning since using under-sampling or over-sampling is a typical method to solve problems such as class imbalance and cost sensitivity [Hulse and Khoshgoftaar, 2007].

In order to solve this and find probabilities that are calibrated, in [Elkan, 2001] a formula that corrects the probabilities based on the difference of the base rates is proposed. The objective is using \hat{p} which was estimated using a population with base rate π_1 , to find \hat{p}' for the real population which has a base rate π'_1 . A solution for \hat{p}' is given as follows:

$$\hat{p}' = \pi'_1 \frac{\hat{p} - \hat{p}\pi_1}{\pi_1 - \pi_1\hat{p} + \pi'_1\hat{p} - \pi_1\pi'_1}. \quad (6.4)$$


6.2.2 Calibration using the ROC convex hull

In order to illustrate the ROC convex hull approach proposed in [Hernandez-Orallo et al., 2012], let us consider the set of probabilities given in Figure 6.1a. Their corresponding ROC curve is shown in Figure 6.1b. It can be seen that this set of probabilities is not calibrated, since when $\text{fpr} = 0.1$ there is a positive example followed by 2 negative examples. This inconsistency is represented in the ROC curve as a non convex segment over the curve.


In order to obtain a set of calibrated probabilities, first the ROC curve must be modified in order to be convex. The way to do that, is to find the convex hull [Hernandez-Orallo et al., 2012] to find the minimal convex set containing the different points of the ROC curve. In Figure 6.1c,

Probability	Label
0.0	0
0.1	1
0.2	0
0.3	0
0.4	1
0.5	0
0.6	1
0.7	1
0.8	0
0.9	1
1.0	1

(a) Set of probabilities and their respective class label



(b) ROC curve of the set of probabilities



(c) Convex hull of the ROC curve

Prob	Cal Prob
0.0	0
0.1	0.333
0.2	0.333
0.3	0.333
0.4	0.5
0.5	0.5
0.6	0.666
0.7	0.666
0.8	0.666
0.9	1
1.0	1

(d) Calibrated probabilities

Figure 6.1: Estimation of calibrated probabilities using the ROC convex hull.

the convex hull algorithm is applied to the previously evaluated ROC curve shown in Figure 6.1b. It is shown that the new curve is convex, and includes all the points of the previous ROC curve.

Now that there is a new convex ROC curve or ROCCH, the calibrated probabilities can be extracted as shown in Figure 6.1d. The procedure to extract the new probabilities is to first group the probabilities according to the points in the ROCCH curve, and then make the calibrated probabilities be the slope of the ROCCH for each group.

6.3 EXPERIMENTS

For the experiments we use five datasets from four different real world example-dependent cost-sensitive problems: Credit card fraud detection (see Section 4.1), credit scoring (see Section 4.2), churn modeling (see Section 5.1) and direct marketing (see Section 5.2). The different datasets are summarized in Table 4.7 and Table 5.3.

For the experiments, we first used three classification algorithms, decision tree (DT), logistic regression (LR) and random forest (RF). Using the implementation of *Scikit-learn* [Pedregosa et al., 2011], each algorithm is trained using the different training sets: training (t), under-sampling (u), cost-proportionate rejection-sampling (r) [Zadrozny et al., 2003] and cost-proportionate over-sampling (o) [Elkan, 2001]. Afterwards, we evaluate the results of the algorithms using BMR methods. In particular, we check the impact on the results of the different calibration methods. The implementation of the cost-sensitive algorithms is done using the *Cost-Cla* library, see Appendix A.

The results are shown in Table 6.1. First, when observing the results of the cost-insensitive methods (CI), that is, DT, LR and RF algorithms trained on the t and u sets, the RF algorithm produces the best result by savings in three out of the five sets, followed by the LR – u. It is also clear that the results on the t dataset are not as good as the ones on the u dataset, this is highly related to the unbalanced distribution of the positives and negatives in all the databases.

In the case of cost-proportionate sampling methods (CPS), specifically the cost-proportionate rejection sampling (r) and cost-proportionate over sampling (o), it is observed that in four cases the savings increase quite significantly. It is on the fraud detection database where these methods do not outperform the algorithms trained on the under-sampled set. This may be related to the fact that in this database the initial percentage of positives is 1.5% which is similar to the percentage in the r and o sets. However, it is 50.42% in the u set, which may help explain why this method performs much better as measured by savings.

Afterwards, in the case of the BMR algorithms, the results show that this method outperforms the previous ones in four cases and has almost the same result in the other set. In the fraud detection set, the results are better, since the savings of the three classification algorithms increase when using this methodology. Moreover, we found that by calibrating the probabilities, the results of the Bayes minimum risk increase. In Figure 6.2, we compare the different families of algorithms, by calculating in each dataset the performance of the methods compared with the best method. First, it is shown the huge difference of using only cost-insensitive methods, compared with any of the cost-sensitive fam-

Family	Algorithm	Fraud	Churn	Credit 1	Credit 2	Marketing
CI	DT-t	0.3176±0.0357	-0.0018±0.0194	0.1931±0.0087	-0.0616±0.0229	-0.2342±0.0609
	LR-t	0.0092±0.0002	-0.0001±0.0002	0.0177±0.0126	0.0039±0.0012	-0.2931±0.0602
	RF-t	0.3342±0.0156	-0.0026±0.0079	0.1471±0.0071	0.0303±0.0040	-0.2569±0.0637
	DT-u	0.5239±0.0118	-0.0389±0.0583	0.3287±0.0125	-0.1893±0.0314	-0.0278±0.0475
	LR-u	0.1243±0.0387	0.0039±0.0492	0.4118±0.0313	0.1850±0.0231	0.2200±0.0376
	RF-u	0.5684±0.0097	0.0433±0.0533	0.4981±0.0079	0.1237±0.0228	0.1227±0.0443
CPS	DT-r	0.3439±0.0453	0.0054±0.0568	0.3310±0.0126	0.0724±0.0212	0.1960±0.0527
	LR-r	0.3077±0.0301	0.0484±0.0375	0.3965±0.0263	0.2650±0.0115	0.4210±0.0267
	RF-r	0.3812±0.0264	0.1056±0.0412	0.4989±0.0080	0.3055±0.0106	0.3840±0.0360
	DT-o	0.3172±0.0274	0.0251±0.0195	0.1738±0.0092	0.0918±0.0225	-0.2598±0.0559
	LR-o	0.2793±0.0185	0.0316±0.0228	0.3301±0.0109	0.2554±0.0090	0.3129±0.0277
	RF-o	0.3612±0.0295	0.0205±0.0156	0.2128±0.0081	0.2242±0.0070	-0.1782±0.0618
BMR	DT-t-BMR	0.6045±0.0386	0.0226±0.0200	0.1931±0.0087	-0.0616±0.0230	-0.2185±0.0611
	LR-t-BMR	0.4552±0.0203	0.0872±0.0308	0.1973±0.0404	0.3119±0.0089	0.4915±0.0090
	RF-t-BMR	0.6175±0.0149	0.0435±0.0356	0.4878±0.0082	0.3027±0.0096	0.3722±0.0263
CAL	DT-t-BMR-cal	0.5936±0.0386	0.0298±0.0145	0.1054±0.0358	0.2740±0.0067	0.4598±0.0089
BMR	LR-t-BMR-cal	0.0897±0.0203	0.1082±0.0316	0.2189±0.0541	0.3148±0.0094	0.4973±0.0084
	RF-t-BMR-cal	0.6414±0.0154	0.0856±0.0354	0.4924±0.0087	0.3133±0.0094	0.4807±0.0093

(Models with the highest savings are marked in bold)

Table 6.1: Results of the algorithms measured by savings




Figure 6.2: **Comparison of the average savings of the algorithms versus the highest savings by family of classifiers.** When the probabilities are calibrated there is a significant increase in savings.

ilies. Furthermore, the Bayes minimum risk methods outperforms the cost-proportionate sampling methods.

Moreover, we evaluate the Brier score of the different algorithms. The objective of this, is because as mention in [Cohen and Goldszmidt, 2004], when using the output of a binary classifier as a basis for decision making, there is a need for a probability that not only separates well between positive and negative examples, but that also assesses the real probability of the event. The results are shown in Table 6.2. First of all, as expected, the logistic regression models are very well calibrated regardless of the dataset used for train them. This is because, the logistic regression algorithm is intended to estimate the most reliable probabilities as we will discuss in Chapter 7.

When comparing the results of the Bayes minimum risk with and without calibration, it is clear that the calibration of the probabilities lead to a lower Brier score as shown in Figure 6.3. Interestingly, the results of Brier score and the savings are highly correlated, confirming the intuition behind the need to calibrate the probabilities before using the Bayes minimum risk method.

Family	Algorithm	Fraud	Churn	Credit 1	Credit 2	Marketing
CI	DT-t	0.0134±0.0005	0.0932±0.0054	0.1054±0.0016	0.3103±0.0041	0.1924±0.0035
	LR-t	0.0440±0.0008	0.0449±0.0025	0.0598±0.0016	0.1511±0.0016	0.0960±0.0018
	RF-t	0.0064±0.0002	0.0532±0.0027	0.0519±0.0007	0.1525±0.0016	0.1090±0.0018
	DT-u	0.1424±0.0071	0.4152±0.0205	0.3141±0.0052	0.4526±0.0057	0.4125±0.0062
	LR-u	0.1872±0.0195	0.2439±0.0139	0.1718±0.0088	0.2317±0.0017	0.2016±0.0027
	RF-u	0.0724±0.0007	0.2442±0.0134	0.1565±0.0019	0.2325±0.0024	0.2256±0.0036
CPS	DT-r	0.2041±0.0111	0.3595±0.0185	0.2884±0.0054	0.5516±0.0070	0.3581±0.0111
	LR-r	0.1726±0.0098	0.1906±0.0112	0.1448±0.0063	0.2657±0.0051	0.2029±0.0055
	RF-r	0.1764±0.0127	0.1964±0.0101	0.1374±0.0028	0.3613±0.0055	0.2123±0.0054
	DT-o	0.1264±0.0102	0.0963±0.0048	0.1020±0.0014	0.4550±0.0039	0.1904±0.0028
	LR-o	0.2158±0.0098	0.1136±0.0025	0.1134±0.0014	0.2258±0.0023	0.1517±0.0013
	RF-o	0.1694±0.0146	0.0701±0.0031	0.0550±0.0008	0.3283±0.0030	0.1190±0.0016
BMR	DT-t-BMR	0.0134±0.0005	0.0932±0.0054	0.1054±0.0016	0.3103±0.0041	0.1924±0.0035
	LR-t-BMR	0.0440±0.0008	0.0449±0.0025	0.0598±0.0016	0.1511±0.0016	0.0960±0.0018
	RF-t-BMR	0.0064±0.0002	0.0532±0.0027	0.0519±0.0007	0.1525±0.0016	0.1090±0.0018
CAL	DT-t-BMR-cal	0.0098±0.0004	0.0448±0.0029	0.0604±0.0009	0.1586±0.0019	0.1075±0.0017
BMR	LR-t-BMR-cal	0.0148±0.0007	0.0444±0.0025	0.0590±0.0018	0.1505±0.0016	0.0952±0.0017
	RF-t-BMR-cal	0.0058±0.0003	0.0446±0.0026	0.0514±0.0008	0.1510±0.0016	0.0999±0.0019

(Models with the lowest Brier score are marked as bold)

Table 6.2: Results of the algorithms measured by Brier score




Figure 6.3: **Comparison of the average Brier score of the algorithms versus the lowest Brier score by family of classifiers.** Overall, the models that are calibrated are indeed the ones with the best Brier score.

OUTLINE

In this chapter, we propose a cost-sensitive logistic regression algorithm. The model consists in a new logistic regression cost function, one that takes into account the real costs due to misclassification and correct classification. First, in Section 7.1, we give the background behind logistic regression. Then, in Section 7.2, we describe the cost-sensitive logistic regression. For this, we carry a deep analysis of the logistic regression implicit misclassification costs in Section 7.2.1. Then in Section 7.2.2, we give a new version of the logistic regression cost function. Finally, in Section 7.3, we compare the results of the proposed algorithm against state-of-the-art methods, using the five real-world cost-sensitive databases presented in Part II.

7.1 LOGISTIC REGRESSION

Logistic regression is a classification model that, in the specific context of binary classification, estimates the posterior probability of the positive class, as the logistic sigmoid of a linear function of the feature vector [Bishop, 2006]. The estimated probability is evaluated as

$$\hat{p}_i = P(y = 1|x_i) = h_{\theta}(x_i) = g\left(\sum_{j=1}^k \theta^j x_i^j\right), \quad (7.1)$$

where $h_{\theta}(x_i)$ refers to the hypothesis of i given the parameters θ , the feature vector x_i for example i and $g(\cdot)$ is the logistic sigmoid function defined as

$$g(z) = \frac{1}{(1 + e^{-z})}. \quad (7.2)$$

In Figure 7.1, the logistic sigmoid function is shown.

The problem then becomes on finding the right parameters that minimize a given cost function. Usually, in the case of logistic regression, the cost function $J(\theta)$ refers to the negative logarithm of the likelihood, such that

$$J(\theta) = \frac{1}{N} \sum_{i=1}^N J_i(\theta), \quad (7.3)$$




Figure 7.1: Sigmoid function

where

$$J_i(\theta) = -y_i \log(h_\theta(\mathbf{x}_i)) - (1 - y_i) \log(1 - h_\theta(\mathbf{x}_i)). \quad (7.4)$$

Therefore, the parameters are estimated by minimizing (7.4) $\hat{\theta} = \arg \min_{\theta} J(\theta)$.

There are several methods used to estimate the logistic regression, in particular, maximum likelihood [Hastie et al., 2009], Newton, coordinate descent [Murphy, 2012] and dual coordinate descent [Yu et al., 2011]. Nevertheless, all these methods rely on the assumption of convexity of the logistic regression cost function $J(\theta)$.

7.2 COST-SENSITIVE LOGISTIC REGRESSION

In this section we present our cost-sensitive logistic regression algorithm. First, we motivate the need to modify the logistic regression, as we analyze the implicit costs that the logistic regression assigned to each misclassification error during the estimation of the parameters θ . Then, we show our proposed algorithm.

7.2.1 Implicit costs of the logistic regression

The logistic regression cost function, as described in (7.4), implicitly assume that false positives and false negatives have the same costs, i.e. $C_{FP_i} = C_{FN_i} \forall i \in \{1, \dots, N\}$. This can be easily shown by analyzing the logistic cost function for both values of y_i and the algorithm prediction $h_\theta(\mathbf{x}_i)$:

- If $y_i = 0$ and $h_\theta(\mathbf{x}_i) \approx 0$, then

$$\begin{aligned} J_i(\theta) &= -y_i \log(h_\theta(\mathbf{x}_i)) - (1 - y_i) \log(1 - h_\theta(\mathbf{x}_i)) \\ &\approx -(0) \log((0)) - (1 - (0)) \log(1 - (0)) \\ &\approx 0. \end{aligned}$$

- If $y_i = 0$ and $h_\theta(\mathbf{x}_i) \approx 1$, then

$$\begin{aligned} J_i(\theta) &\approx -(0) \log((1)) - (1 - (0)) \log(1 - (1)) \\ &\approx \infty. \end{aligned}$$

- If $y_i = 1$ and $h_\theta(\mathbf{x}_i) \approx 0$, then

$$\begin{aligned} J_i(\theta) &\approx -(1) \log((0)) - (1 - (1)) \log(1 - (0)) \\ &\approx \infty. \end{aligned}$$

- If $y_i = 1$ and $h_\theta(\mathbf{x}_i) \approx 1$, then

$$\begin{aligned} J_i(\theta) &\approx -(1) \log((1)) - (1 - (1)) \log(1 - (1)) \\ &\approx 0. \end{aligned}$$

Then, we collect the previous results in a cost matrix:

	Actual Positive $y_i = 1$	Actual Negative $y_i = 0$
Predicted Positive $c_i = 1$	$C_{TP_i} \approx 0$	$C_{FP_i} \approx \infty$
Predicted Negative $c_i = 0$	$C_{FN_i} \approx \infty$	$C_{TN_i} \approx 0$

Table 7.1: Logistic regression cost matrix

This confirms that the logistic regression cost function implicitly assume that false positives and false negatives have the same costs. However, as discussed in Chapter 4 and Chapter 5, this is not the case in several real-world applications.

7.2.2 Cost-sensitive logistic regression cost function

In order to incorporate the different real costs, as showed in Table 3.1, into the logistic regression, we start by analyzing the expected costs that a modified logistic regression cost function should make for each misclassification and correct classification case.

$$J_i^c(\theta) = \begin{cases} C_{TP_i} & \text{if } y_i = 1 \text{ and } h_\theta(\mathbf{x}_i) \approx 1 \\ C_{TN_i} & \text{if } y_i = 0 \text{ and } h_\theta(\mathbf{x}_i) \approx 0 \\ C_{FP_i} & \text{if } y_i = 0 \text{ and } h_\theta(\mathbf{x}_i) \approx 1 \\ C_{FN_i} & \text{if } y_i = 1 \text{ and } h_\theta(\mathbf{x}_i) \approx 0. \end{cases}$$

Then, as we already have the real costs, we create a new cost-sensitive logistic regression cost function, by including the different costs into the logistic function,

$$J^c(\theta) = \frac{1}{N} \sum_{i=1}^N \left(y_i(h_\theta(\mathbf{x}_i)C_{TP_i} + (1 - h_\theta(\mathbf{x}_i))C_{FN_i}) \right. \\ \left. + (1 - y_i)(h_\theta(\mathbf{x}_i)C_{FP_i} + (1 - h_\theta(\mathbf{x}_i))C_{TN_i}) \right). \quad (7.5)$$

Since this new cost function is not convex, we estimate the parameters θ with genetic algorithms, as this optimization heuristic does not require the underlying function to be differentiable or convex [Haupt and Haupt, 2004].

7.3 EXPERIMENTS

For the experiments we use five datasets from four different real world example-dependent cost-sensitive problems: Credit card fraud detection (see Section 4.1), credit scoring (see Section 4.2), churn modeling (see Section 5.1) and direct marketing (see Section 5.2). The different datasets are summarized in Table 4.7 and Table 5.3.

In particular, we are interested in comparing the results of the different logistic regression models. First we train a logistic regression (LR) using the training (t), under-sampled (u), cost-proportionate rejection-sampling (r) [Zadrozny et al., 2003] and cost-proportionate over-sampling (o) [Elkan, 2001] datasets. Afterwards, we evaluate the results of the algorithms using BMR, see Chapter 6. Lastly, we calculate the cost-sensitive logistic regression (CSLR). We use the logistic regression implementation of *Scikit-learn* [Pedregosa et al., 2011], and the *Cost-Cla* library, see Appendix A, for the cost-sensitive algorithms. Unless otherwise stated, the random selection of the training set was repeated 50 times, and in each time the models were trained and results collected, this allows us to measure the stability of the results.

In Table 7.2, we show the results of each algorithm in the different databases measured by savings. Firstly, the proposed CSLR has the




Figure 7.2: **Comparison of the average savings and F_1 Score of the algorithms versus the the best model.** The models that perform the best measured by F_1 Score are not the best in terms of savings.

highest savings in the fraud, churn and credit 1 databases, moreover, in the credit 2 and marketing databases, it is the second best model. It is interesting to see how different the results from a standard logistic regression and the cost-sensitive logistic regression are. Moreover, we also calculate the results of the F_1 Score for each model, as shown in Table 7.3. It is observed that the CSLR algorithm is not the one that gives the best results measured by F_1 Score. In fact, there is not a clear relation between the results measured by savings or F_1 Score.

Finally, we compute the perBest statistic for the savings and the F_1 Score. This statistic is calculated as the average result of each algorithm compared with the best in each set. The results are shown in Figure 7.2. On average, the CSLR is the best model, as it yield to 95.5% of the best model in the different databases. Nevertheless, when measured by F_1 Score, it only yield to 74.7% of the best model. This leads to the conclusion of the importance of using a cost-sensitive measure such as savings when evaluating real-world example-dependent cost-sensitive problems.

Moreover, it is observed that the standard logistic regression trained using the training set is the model that performs the worst measured by both savings and F_1 Score. This is related not only to the cost-sensitivity of the problems, but also to the highly unbalanced distribution of positive and negatives presented in all the databases. This is the reason why by using an under-sampling procedure, the results are improved by both measures.

Family	Algorithm	Fraud	Churn	Credit 1	Credit 2	Marketing
CI	LR-t	0.0092±0.0002	-0.0001±0.0002	0.0177±0.0126	0.0039±0.0012	-0.2931±0.0602
	LR-u	0.1243±0.0387	0.0039±0.0492	0.4118±0.0313	0.1850±0.0231	0.2200±0.0376
CPS	LR-r	0.3077±0.0301	0.0484±0.0375	0.3965±0.0263	0.2650±0.0115	0.4210±0.0267
	LR-o	0.2793±0.0185	0.0316±0.0228	0.3301±0.0109	0.2554±0.0090	0.3129±0.0277
BMR	LR-t-BMR	0.4552±0.0203	0.1082±0.0316	0.2189±0.0541	0.3148±0.0094	0.4973±0.0084
CST	CSLR-t	0.6113±0.0262	0.1118±0.0484	0.4554±0.1039	0.2748±0.0069	0.4484±0.0072

(Models with the highest savings are marked in bold)

Table 7.2: Results of the algorithms measured by savings

Family	Algorithm	Fraud	Churn	Credit 1	Credit 2	Marketing
CI	LR-t	0.1531±0.0045	0.0000±0.0000	0.0494±0.0277	0.0155±0.0037	0.2702±0.0125
	LR-u	0.0241±0.0163	0.1222±0.0098	0.3160±0.0314	0.3890±0.0053	0.3440±0.0083
CPS	LR-r	0.1846±0.0123	0.1258±0.0111	0.3597±0.0156	0.3793±0.0049	0.3374±0.0101
	LR-o	0.1776±0.0117	0.1085±0.0203	0.3769±0.0067	0.3804±0.0044	0.3568±0.0102
BMR	LR-t-BMR	0.1384±0.0044	0.1370±0.0150	0.1915±0.0340	0.3572±0.0045	0.2954±0.0079
CST	CSLR-t	0.2031±0.0065	0.1134±0.0151	0.1454±0.0517	0.3363±0.0045	0.2339±0.0051

(Models with the highest F1 Score are marked in bold)

Table 7.3: Results of the algorithms measured by F1 Score

OUTLINE

In this chapter, we propose the example-dependent cost-sensitive decision tree algorithm. This algorithm is based on incorporating the different example-dependent costs into a new cost-based impurity measure and a new cost-based pruning criteria. First, in Section 8.2, we give the background behind the decision tree algorithm. Then, in Section 8.3, we introduce the cost-sensitive decision tree method. For this, we first show the new cost-sensitive impurity measure that introduces the different costs when evaluating the performance of a split. Then, we introduce a new cost-sensitive pruning criteria. Finally, in Section 8.4, we compare the results of the proposed algorithm against state-of-the-art methods using the five real-world cost-sensitive databases.

8.1 INTRODUCTION

Decision trees are one of the most widely used machine learning algorithms [Maimon, 2008]. The technique is considered to be white box, in the sense that is easy to interpret, and has a very low computational cost, while maintaining a good performance as compared with more complex techniques [Hastie et al., 2009]. There are two types of decision tree depending on the objective of the model. They work either for classification or regression. In this section we focus on binary classification decision tree.

Standard decision tree algorithms focus on inducing trees that maximize accuracy. However this is not optimal when the misclassification costs are unequal [Elkan, 2001]. This has led to many studies that develop algorithms that aim to introduce the cost-sensitivity into the algorithms [Lomax and Vadera, 2013]. These studies have focused on introducing the class-dependent costs [Draper et al., 1994; Ting, 2002; Ling et al., 2004; Li et al., 2005; Kretowski and Grześ, 2006; Vadera, 2010], which is not optimal for some applications. For example in credit card fraud detection, it is true that false positives have a different cost than false negatives, nevertheless, false negatives may vary significantly, which makes class-dependent cost-sensitive methods not suitable for this problem.

In this chapter, we propose a cost-sensitive decision tree algorithm. The algorithm is based on incorporating the different example-

dependent costs into a new cost-based impurity measure and a new cost-based pruning criteria. We evaluate the proposed example-dependent cost-sensitive decision tree using three different databases. In particular, a credit card fraud detection, a credit scoring and a direct marketing databases. The results show that the proposed method outperforms state-of-the-art example-dependent cost-sensitive methods. Furthermore, when compared against a standard decision tree, our method builds significantly smaller trees in only a fifth of the time.

8.2 DECISION TREES

There are two types of decision tree depending on the objective of the model. They work either for classification or regression. In this section we focus on binary classification decision tree.

8.2.1 Construction of classification trees

Classification trees is one of the most common types of decision tree, in which the objective is to find the Tree that best discriminates between classes. In general the decision tree represents a set of splitting rules organized in levels in a flowchart structure.

8.2.1.1 Splitting criteria

In the Tree, each rule is shown as a node, and it is represented as (\mathbf{x}^j, \bar{v}) , meaning that the set \mathcal{S} is split in two sets \mathcal{S}^l and \mathcal{S}^r according to \mathbf{x}^j and \bar{v} :

$$\mathcal{S}^l = \{\mathbf{x}_i^* | \mathbf{x}_i^* \in \mathcal{S} \wedge x_i^j \leq \bar{v}\}, \quad (8.1)$$

and

$$\mathcal{S}^r = \{\mathbf{x}_i^* | \mathbf{x}_i^* \in \mathcal{S} \wedge x_i^j > \bar{v}\}, \quad (8.2)$$

where x_i^j is the j^{th} feature represented in the vector $\mathbf{x}^j = [x_1^j, x_2^j, \dots, x_N^j]$, and \bar{v} is a value such that $\min(\mathbf{x}^j) \leq \bar{v} < \max(\mathbf{x}^j)$. Moreover, $\mathcal{S} = \mathcal{S}^l \cup \mathcal{S}^r$.

After the training set have been split, the percentage of positives in the different sets is calculated. First, the number of positives in each set is estimated by

$$\mathcal{S}_1 = \{\mathbf{x}_i^* | \mathbf{x}_i^* \in \mathcal{S} \wedge y_i = 1\}, \quad (8.3)$$

then the percentage of positives is calculates as

$$\pi_1 = |\mathcal{S}_1|/|\mathcal{S}|. \quad (8.4)$$




Figure 8.1: Impurity measures for a binary classification, as a function of the proportion of positive examples in the set (Cross-entropy is scaled) [Hastie et al., 2009].

Then, the impurity of each leaf is calculated using either a misclassification error, entropy or Gini measures:

- Misclassification

$$I_m(\pi_1) = 1 - \max(\pi_1, 1 - \pi_1) \quad (8.5)$$

- Entropy

$$I_e(\pi_1) = -\pi_1 \log \pi_1 - (1 - \pi_1) \log(1 - \pi_1) \quad (8.6)$$

- Gini

$$I_g(\pi_1) = 2\pi_1(1 - \pi_1) \quad (8.7)$$

In Figure 8.1, the three measures are shown. All measures are similar, with the exception that the gini index and the cross-entropy are differentiable, which makes them more suitable for numerical optimization.

Finally the gain of the splitting criteria using the rule (\mathbf{x}^j, ν) is calculated as the impurity of \mathcal{S} minus the weighted impurity of each leaf:

$$\text{Gain}(\mathbf{x}^j, \nu) = I(\pi_1) - \frac{|\mathcal{S}^l|}{|\mathcal{S}|} I(\pi_1^l) - \frac{|\mathcal{S}^r|}{|\mathcal{S}|} I(\pi_1^r), \quad (8.8)$$

where $I(\pi_1)$ can be either of the impurity measures $I_e(\pi_1)$ or $I_g(\pi_1)$.

Subsequently, the gain of all possible splitting rules is calculated. The rule with maximal gain is selected

$$(\text{best}_x, \text{best}_l) = \arg \max_{(x^j, l^j)} \text{Gain}(x^j, l^j), \quad (8.9)$$

and the set S is split into S^l and S^r according to that rule.

8.2.1.2 Tree growing

In order to grow a tree typical algorithms use a top-down induction using a greedy search in each iteration [Rokach and Maimon, 2010]. In each iteration, the algorithms evaluates all possible splitting rules and pick the one that maximizes the splitting criteria. After the selection of a splitting rule, each leaf is further selected and it is subdivides into smaller leaves, until one of the stopping criteria is meet. In Algorithm 8.1, the pseudocode of the tree growing procedure is presented.

Algorithm 8.1. *Pseudocode of the tree growing procedure.*

```

1: function TREEGROW( $S$ )
2:   Tree = Empty
3:   if STOPPING(Tree,  $S$ ) then
4:     return Tree
5:   end if
6:   # For all features
7:   for  $j = 1$  to  $k$  do
8:     # For all possible thresholds
9:     for  $m = 1$  to  $N_k$  do
10:      Gains( $j, m$ ) = GAIN( $x^j, l_m^j, S$ )
11:    end for
12:  end for
13:  # Select the threshold with the highest gain
14:   $(j^*, l^*) = \arg \max_{(j, m)} \text{Gains}$ 
15:  # Split the set using the best rule
16:   $S^l = \{x_i^* | x_i^* \in S \wedge x_i^{j^*} \leq l^*\}$ 
17:   $S^r = \{x_i^* | x_i^* \in S \wedge x_i^{j^*} > l^*\}$ 
18:  # Recursively continue growing tree
19:  TREEGROW( $S^l$ )
20:  TREEGROW( $S^r$ )
21:  # Return the tree
22:  return Tree
23: end function

```

8.2.1.3 Stopping criteria

As the growing phase of the algorithm continue, at each iteration the stopping criteria are evaluated [Rokach and Maimon, 2010]. The most common stopping criteria are:

- All examples in the set belong to the same class. Meaning that either $\pi_1 = 1$ or $\pi_1 = 0$.
- The maximum number of iterations have been reached.
- The number of examples in S are less than the minimum number of examples defined for a split.
- The number of examples in S^l or S^r are less than the minimum number of examples defined for a leaf.
- The best splitting rule has a Gain lower than a defined threshold.

8.2.1.4 Pruning techniques

After a decision tree has been fully grown, there is the big chance that the algorithm generates a large tree that is probably over fitting the training data. In order to solve this in [Breiman et al., 1984] originally suggest the use of pruning techniques after the tree growing phase. The overall objective of pruning is to eliminate branches that are not contributing to the generalization accuracy of the tree [Rokach and Maimon, 2010].

In general pruning techniques start from a fully grown tree, and recursively check if by eliminating a Branch there is an Improvement in the error rate Err of the Tree. There are two main methods to calculate the Improvement of the error rate Err of the Tree, cost-complexity and error based pruning:

- Cost-complexity pruning:

Initially proposed by Breiman [Breiman et al., 1984], this method evaluate iterative if the removal of a Branch improve the error rate Err of a Tree, weighted by the difference of the number of leafs trees.

$$PC_{cc} = \frac{\text{Err}(\text{EB}(\text{Tree}, \text{branch}), S) - \text{Err}(\text{Tree}, S)}{|\text{Tree}| - |\text{EB}(\text{Tree}, \text{branch})|} \quad (8.10)$$

where $\text{Err}(\text{Tree}, S)$ calculate the error rate of the Tree in set S and $\text{EB}(\text{Tree}, \text{branch})$ is an auxiliary function that removes branch from Tree and return a new Tree. At each iteration, the current Tree is compared against all possible Branches.

- Error based pruning:

This method proposed by Quinlan [Quinlan, 1992] is more pessimistic and instead of using the error rate Err calculate the upper bound of the confidence interval of the error Err_{UB} calculated assuming a normal distribution with a defined significance level Z_α .

$$\text{PC}_{\text{eb}} = \frac{\text{Err}_{\text{UB}}(\text{EB}(\text{Tree}, \text{branch}), \mathcal{S}) - \text{Err}_{\text{UB}}(\text{Tree}, \mathcal{S})}{|\text{Tree}| - |\text{EB}(\text{Tree}, \text{branch})|} \quad (8.11)$$

where

$$\begin{aligned} \text{Err}_{\text{UB}}(\text{Tree}, \mathcal{S}) &= \text{Err}(\text{Tree}, \mathcal{S}) \\ &+ Z_\alpha \sqrt{\frac{\text{Err}(\text{Tree}, \mathcal{S})(1 - \text{Err}(\text{Tree}, \mathcal{S}))}{|\mathcal{S}|}} \end{aligned} \quad (8.12)$$

Moreover, as a difference from cost-complexity pruning, in the error based method, at each iteration not all the branches are compared, but only the ones on the same level. Meaning leafs are only compared against other leafs that are in the same level of iterations when the Tree was constructed. This allows for much faster pruning procedure.

In Algorithm 8.2, the general method of pruning is shown. The pruning procedure consists in evaluating the Improvement of eliminating all possible branches in a Tree, and then eliminate the one with the higher Improvement and repeat the process until a stopping criteria is meet.

Algorithm 8.2. *Pseudocode of the tree pruning procedure.*

```

1: function TREEPRUNING( $\mathcal{S}$ , Tree)
2:   # For all branches
3:   for branch in Tree.branches do
4:     # Evaluate the pruning criteria
5:     Improvements(branch) = PC( $\mathcal{S}$ , Tree, branch)
6:   end for
7:   # Select the branch to prune
8:   branch* =  $\text{argmax}_{\text{branch}} \text{Improvements}$ 
9:   # Prune selected branch
10:  Tree = EB(Tree, branch*)
11:  # Recursively continue pruning the tree
12:  TREEPRUNING( $\mathcal{S}$ , Tree)
13:  return Tree
14: end function

```


8.2.1.5 Categorical features

When using continuous features the splitting is done by testing all possible thresholds on the database and selecting the one that maximizes the desired measure. Nevertheless, this is not straight forward when using a categorical feature, since testing all possible ways of binning the feature becomes prohibitively time consuming the feature has many categories. Instead, a common method for doing this is to calculate π_1 of each category, then sort it, and applied the method as a continuous feature [Marslan, 2009].

8.2.2 Decision tree algorithms

Two main branches of decision trees has being studied in the last years. The main difference is the impurity measure used when splitting. First the CART algorithm that is based in the Gini index, and later the ID3 and C4.5 which uses the entropy measure.

- CART or classification and regression trees was introduced by Brieman et al. in 1984 [Breiman et al., 1984]. It is based on using the Gini index as the impurity measure and the tree is grow until all examples in each leaf belong to the same class. Afterwards, the tree is pruned using the cost-complexity method [Rokach and Maimon, 2010; Marslan, 2009].
- ID3 algorithm uses entropy as the impurity measure. The growing of the tree stop when all examples belong of each leaf belongs to the same class. In ID3 no pruning is applied [Quinlan, 1992].
- C4.5 the extension of ID3 both proposed by Quinlan [Quinlan, 1992]. Both are similar regarding the measure used, but C4.5 define the stopping criteria during the growth process to be when the number of examples in a set is less than a threshold. Moreover, after the tree is created a error based pruning is applied [Rokach and Maimon, 2010].

8.3 EXAMPLE-DEPENDENT COST-SENSITIVE DECISION TREES

In this section, we first propose a new method to introduce the costs into the decision tree induction stage, by creating new-cost based impurity measures. Afterwards, we propose a new pruning method based on minimizing the cost as pruning criteria.

8.3.1 Cost-sensitive impurity measures

Standard impurity measures such as misclassification, entropy or Gini, take into account the distribution of classes of each leaf to evaluate the predictive power of a splitting rule, leading to an impurity measure that is based on minimizing the misclassification rate. However, as has been previously shown in Chapter 6 and Chapter 7, minimizing misclassification does not lead to the same results than minimizing cost. Instead, we are interested in measuring how good is a splitting rule in terms of cost not only accuracy. For doing that, we propose a new example-dependent cost based impurity measure that takes into account the cost matrix of each example.

We define a new cost-based impurity measure taking into account the costs when all the examples in a leaf are classified both as negative using f_0 and positive using f_1

$$I_c(\mathcal{S}) = \min \left\{ \text{Cost}(f_0(\mathcal{S})), \text{Cost}(f_1(\mathcal{S})) \right\}. \quad (8.13)$$

The objective of this measure is to evaluate the lowest expected cost of a splitting rule. Following the same logic, the classification of each set is calculated as the prediction that leads to the lowest cost

$$f(\mathcal{S}) = \begin{cases} 0 & \text{if } \text{Cost}(f_0(\mathcal{S})) \leq \text{Cost}(f_1(\mathcal{S})) \\ 1 & \text{otherwise} \end{cases} \quad (8.14)$$

Finally, using the cost-based impurity, the splitting criteria cost based gain of using the splitting rule $(\mathbf{x}^j, \mathcal{V})$ is calculated with (8.8).

8.3.2 Cost-sensitive pruning

Most of the literature in class-dependent cost-sensitive decision tree focuses on using the misclassification costs during the construction of the algorithms [Lomax and Vadera, 2013]. Only few algorithms such as AUCSplit [Ferri et al., 2002] have included the costs both during and after the construction of the tree. However, this approach only used the class-dependent costs, and not the example-dependent costs.

We propose a new example-dependent cost-based impurity measure, by replacing the error rate Err in (8.10) with the cost of using the Tree on \mathcal{S} i.e., by replacing with $\text{Cost}(f(\mathcal{S}))$.

$$\text{PC}_c = \frac{\text{Cost}(f(\mathcal{S})) - \text{Cost}(f^*(\mathcal{S}))}{|\text{Tree}| - |\text{EB}(\text{Tree}, \text{node})|}, \quad (8.15)$$

where f^* is the classifier of the tree without the selected node $\text{EB}(\text{Tree}, \text{node})$.

Using the new pruning criteria, nodes of the tree that do not contribute to the minimization of the cost will be pruned, regardless of the impact of those nodes on the accuracy of the algorithm. This follows the same logic as in the proposed cost-based impurity measure, since minimizing the misclassification is different than minimizing the cost, and in several real-world applications the objectives align with the cost not with the misclassification error.

8.4 EXPERIMENTS

In this section we present the experimental results. First, we evaluate the performance of the proposed CSDT algorithm and compare it against a classical decision tree (DT). We evaluate the different trees using them without pruning (notp), with error based pruning (errp), and also with the proposed cost-sensitive pruning technique (costp). The different algorithms are trained using the training (t), under-sampling (u), cost-proportionate rejection-sampling (r), and cost-proportionate over-sampling (o) datasets. We use the decision trees implementation of *Scikit-learn* [Pedregosa et al., 2011], and the *CostCla* library, see Appendix A, for the cost-sensitive algorithms.

For the experiments we use three datasets from three different real world example-dependent cost-sensitive problems: Credit card fraud detection (see Section 4.1), credit scoring (see Section 4.2) and direct marketing (see Section 5.2). The different datasets are summarized in Table 4.7 and Table 5.3.

We evaluate a decision tree constructed using the Gini impurity measure, with and without the pruning defined in (8.10). We also apply the cost-based pruning procedure given in (8.15). Lastly, we compared against the proposed CSDT constructed using the cost-based impurity measure defined in (8.13), using the two pruning procedures.

In Figure 8.2, the results using the three databases are shown. In particular we first evaluate the impact of the algorithms when trained using the training set. There is a clear difference between the savings of the DT and the CSDT algorithms. However, that difference is not observable on the F_1 Score results. Since the CSDT is focused on maximizing the savings not the accuracy or F_1 Score. There is a small increase in savings when using the DT with cost-sensitive pruning. Nevertheless, in the case of the CSDT algorithm, there is no change when using any pruning procedure, neither in savings or F_1 Score.

In addition, we also evaluate the algorithms on the different sets, under-sampling, rejection-sampling and over-sampling. The results are shown in Table 8.1. Moreover, in Figure 8.3, the average results of the different algorithms measured by savings is shown. The best results are

set	Algorithm	Fraud Detection		Direct Marketing		Credit Scoring	
		%Sav	F ₁ Score	%Sav	F ₁ Score	%Sav	F ₁ Score
t	DT _{notp}	31.76	0.4458	19.11	0.2976	18.95	0.3062
	DT _{errp}	31.76	0.4458	19.70	0.3147	18.95	0.3062
	DT _{costp}	35.89	0.4590	28.08	0.3503	27.59	0.3743
	CSDT _{notp}	70.85	0.2529	69.00	0.2920	49.28	0.3669
	CSDT _{errp}	70.85	0.2529	68.97	0.3193	48.85	0.3669
	CSDT _{costp}	71.16	0.2522	69.105	0.2878	49.39	0.3684
u	DT _{notp}	52.39	0.1502	49.80	0.3374	48.91	0.2983
	DT _{errp}	52.39	0.1502	49.80	0.3374	48.91	0.2983
	DT _{costp}	70.26	0.2333	53.20	0.3565	49.77	0.3286
	CSDT _{notp}	12.46	0.0761	64.21	0.2830	30.68	0.2061
	CSDT _{errp}	14.98	0.0741	66.21	0.2822	41.49	0.2564
	CSDT _{costp}	15.01	0.0743	68.07	0.2649	44.89	0.2881
r	DT _{notp}	34.39	0.4321	68.59	0.3135	48.97	0.3931
	DT _{errp}	34.39	0.4321	68.79	0.3196	48.97	0.3931
	DT _{costp}	38.99	0.4478	69.27	0.3274	50.48	0.3501
	CSDT _{notp}	70.85	0.2529	66.87	0.2761	31.25	0.1940
	CSDT _{errp}	70.85	0.2529	67.47	0.2581	40.69	0.2529
	CSDT _{costp}	71.09	0.2515	68.08	0.2642	44.51	0.2869
o	DT _{notp}	31.72	0.4495	60.30	0.3674	47.39	0.3994
	DT _{errp}	31.72	0.4495	60.30	0.3674	47.39	0.3994
	DT _{costp}	37.35	0.4575	69.44	0.3108	50.92	0.3977
	CSDT _{notp}	70.84	0.2529	68.75	0.2935	41.37	0.2205
	CSDT _{errp}	70.84	0.2529	68.75	0.2935	41.38	0.3457
	CSDT _{costp}	71.09	0.2515	68.75	0.2935	41.65	0.2896

Table 8.1: Results on the three datasets of the cost-sensitive and standard decision tree, without pruning (notp), with error based pruning (errp), and with cost-sensitive pruning technique (costp). Estimated using the different training sets: training, under-sampling, cost-proportionate rejection-sampling and cost-proportionate over-sampling




Figure 8.2: Results of the DT and the CSDT. For both algorithms, the results are calculated with and without both types of pruning criteria.




Figure 8.3: Average savings on the three datasets of the different cost-sensitive and standard decision tree, estimated using the different training sets: training, under-sampling, cost-proportionate rejection-sampling and cost-proportionate over-sampling.

found when using the training set. When using the under-sampling set there is a decrease in savings of the CSDT algorithm. Lastly, in the case of the cost-proportionate sampling sets, there is a small increase in savings when using the CSDT algorithm.

Finally, we also analyze the different models taking into account the complexity and the training time. In particular we evaluate the size of each Tree. In Table 8.2, and Figure 8.4, the results are shown. The CSDT algorithm creates significantly smaller trees, which leads to a lower training time. In particular this is a result of using the non weighted gain, the CSDT only accepts splitting rules that contribute to the overall reduction of the cost, which is not the case if instead the weighted gain was used. Even that the DT with cost pruning, produce a good re-

set	Algorithm	Fraud Detection		Direct Marketing		Credit Scoring	
		Tree	Time	Tree	Time	Tree	Time
t	DT _{notp}	488	2.45	298	1.58	292	1.58
	DT _{errp}	488	3.90	298	2.13	292	2.13
	DT _{costp}	446	19.19	291	5.23	280	5.23
	CSDT _{notp}	89	1.47	51	0.40	69	0.40
	CSDT _{errp}	88	1.87	51	0.64	69	0.64
	CSDT _{costp}	89	1.74	51	0.45	69	0.45
u	DT _{notp}	308	1.10	198	1.00	167	1.00
	DT _{errp}	308	1.43	198	1.14	167	1.14
	DT _{costp}	153	2.59	190	1.34	142	1.34
	CSDT _{notp}	14	0.20	23	0.17	42	0.17
	CSDT _{errp}	14	0.23	23	0.19	42	0.19
	CSDT _{costp}	14	0.24	23	0.18	42	0.18
r	DT _{notp}	268	0.98	181	0.90	267	0.90
	DT _{errp}	268	1.24	181	0.95	267	0.95
	DT _{costp}	153	2.48	162	1.20	261	1.20
	CSDT _{notp}	18	0.22	10	0.07	70	0.07
	CSDT _{errp}	18	0.23	10	0.07	70	0.07
	CSDT _{costp}	18	0.23	10	0.07	70	0.07
o	DT _{notp}	425	2.30	340	1.80	277	1.80
	DT _{errp}	425	3.98	340	2.65	277	2.65
	DT _{costp}	364	10.15	288	5.99	273	5.99
	CSDT _{notp}	37	1.58	51	0.38	70	0.38
	CSDT _{errp}	37	1.90	51	0.45	70	0.45
	CSDT _{costp}	37	1.98	51	0.42	70	0.42

Table 8.2: Training time and tree size of the different cost-sensitive and standard decision tree, estimated using the different training sets: training, under-sampling, cost-proportionate rejection-sampling and cost-proportionate over-sampling, for the three databases.

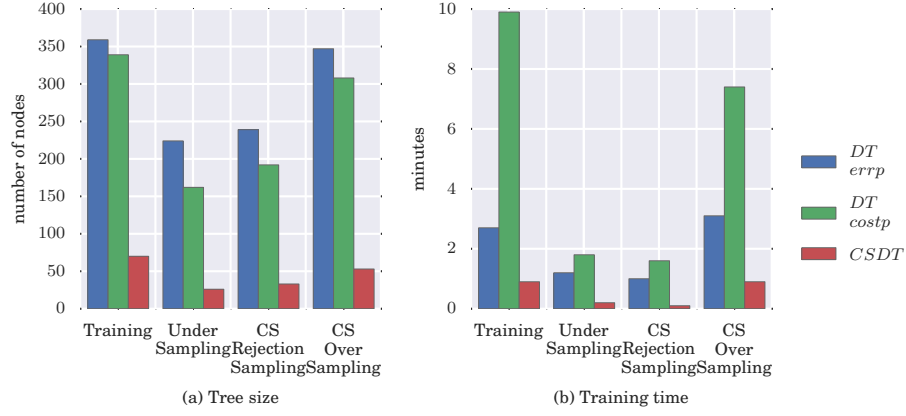


Figure 8.4: Average tree size (a) and training time (b), of the different cost-sensitive and standard decision tree, estimated using the different training sets.

sult measured by savings, it is the one that takes the longer to estimate. Since the algorithm first creates a big decision tree using the Gini impurity, and then attempt to find a smaller tree taking into account the cost. Measured by training time, the CSDT is by all means faster to train than the DT algorithm, leading to an algorithm that not only gives better results measured by savings but also one that can be trained much quicker than the standard DT.

OUTLINE

In this chapter, we propose a framework of ensembles of example-dependent cost-sensitive decision-trees. The algorithms are based in training example-dependent cost-sensitive decision trees using four different random inducer methods and then blending them using three different combination approaches. First, in Section 9.2, we give the background behind ensemble learning. Then, in Section 9.3, we present the ensembles of cost-sensitive decision-trees framework. Moreover, in Section 9.4, we prove theoretically that combining individual cost-sensitive classifiers achieves better results in the sense of higher financial savings. Finally, in Section 9.5, we compare the results of the proposed algorithm, against state-of-the-art methods, using the five real-world cost-sensitive databases.

9.1 INTRODUCTION

Ensemble learning is a widely studied topic in the machine learning community. The main idea behind the ensemble methodology is to combine several individual base classifiers in order to have a classifier that outperforms each of them [Rokach, 2009].

In the context of cost-sensitive classification, some authors have proposed methods for using ensemble techniques. In [Masnadi-shirazi and Vasconcelos, 2011], the authors proposed a framework for cost-sensitive boosting that is expected to minimize the losses by using optimal cost-sensitive decision rules. In [Street, 2008], a bagging algorithm with adaptive costs was proposed. In his doctoral thesis, Nesbitt [Nesbitt, 2010], proposed a method for cost-sensitive tree-stacking. In this method different decision trees are learned, and then combined in a way that a cost function is minimized. Lastly in [Lomax and Vadera, 2013], a survey of application of cost-sensitive learning with decision trees is shown, in particular including other methods that create cost-sensitive ensembles. However, in all these methods, the misclassification costs only depend on the class, therefore, assuming a constant cost across examples. As a consequence, these methods are not well suited for example-dependent cost-sensitive problems.

In this chapter, we propose a framework of ensembles of example-dependent cost-sensitive decision-trees. The algorithms are based in

training example-dependent cost-sensitive decision trees using four different random inducer methods and then blending them using three different combination approaches. Moreover, we propose two new cost-sensitive combination approaches, cost-sensitive weighted voting and cost-sensitive stacking. The latter being an extension of our previously proposed cost-sensitive logistic regression. We evaluate the proposed framework using five different databases from four real-world problems. In particular, credit card fraud detection, churn modeling, credit scoring and direct marketing. The results show that the proposed method outperforms state-of-the-art example-dependent cost-sensitive methods in three databases, and have a similar result in the other two.

9.2 ENSEMBLE METHODS

Nowadays, ensemble methods are one of the most popular and well studied machine learning techniques [Zhou, 2012], and it can be noted that since 2009 all the first-place and second-place winners of the KDD-Cup competition¹ used ensemble methods. The core principle in ensemble learning, is to induce random perturbations into the learning procedure in order to produce several different base classifiers from a single training set, then combining the base classifiers in order to make the final prediction. In order to induce the random perturbations and therefore create the different base classifiers, several methods have been proposed, in particular: bagging [Breiman, 1996], pasting [Breiman, 1999], random forests [Breiman, 2001] and random patches [Louppe and Geurts, 2012]. Finally, after the base classifiers are trained, they are typically combined using either majority voting, weighted voting or stacking [Zhou, 2012].

As shown in Figure 9.1, there are three main reasons regarding why ensemble methods perform better than single models: statistical, computational and representational [Dietterich, 2000]. First, from a statistical point of view, when the learning set is too small, an algorithm can find several good models within the search space, that arise to the same performance on the training set S . Nevertheless, without a validation set, there is a risk of choosing the wrong model. The second reason is computational; in general, algorithms rely on some local search optimization and may get stuck in a local optima. Then, an ensemble may solve this by focusing different algorithms to different spaces across the training set. The last reason is representational. In most cases, for a learning set of finite size, the true function f cannot be represented by any of the candidate models. By combining several models in an ensem-

¹ <https://www.sigkdd.org/kddcup/>




Figure 9.1: Main reasons regarding why ensemble methods perform better than single models: statistical, computational and representational [Dietterich, 2000].

ble, it may be possible to obtain a model with a larger coverage across the space of representable functions.

The most typical form of an ensemble is made by combining T different base classifiers. Each base classifier $M(\mathcal{S}_j)$ is trained by applying algorithm M to a random subset \mathcal{S}_j of the training set \mathcal{S} . For simplicity we define $M_j \equiv M(\mathcal{S}_j)$ for $j = 1, \dots, T$, and $\mathcal{M} = \{M_j\}_{j=1}^T$ a set of base classifiers. Then, these models are combined using majority voting to create the ensemble H as follows

$$f_{mv}(\mathcal{S}, \mathcal{M}) = \arg \max_{c \in \{0,1\}} \sum_{j=1}^T \mathbf{1}_c(M_j(\mathcal{S})). \quad (9.1)$$

9.2.1 Theoretical performance of an ensemble

If we assume that each one of the T base classifiers has a probability ρ of being correct, the probability of an ensemble making the correct de-

cision, assuming independence, denoted by P_c , can be calculated using the binomial distribution [Hansen and Salamon, 1990]

$$P_c = \sum_{j \geq T/2}^T \binom{T}{j} \rho^j (1 - \rho)^{T-j}. \quad (9.2)$$

Furthermore, as shown in Lam and Suen [1997], if $T \geq 3$ then:

$$\lim_{T \rightarrow \infty} P_c = \begin{cases} 1 & \text{if } \rho > 0.5 \\ 0 & \text{if } \rho < 0.5 \\ 0.5 & \text{if } \rho = 0.5, \end{cases} \quad (9.3)$$

leading to the conclusion that

$$\rho \geq 0.5 \quad \text{and} \quad T \geq 3 \quad \Rightarrow \quad P_c \geq \rho. \quad (9.4)$$

9.3 ENSEMBLES OF COST-SENSITIVE DECISION TREES

In this section we present our framework of ensembles of example-dependent cost-sensitive decision-trees. The algorithms are based in training example-dependent cost-sensitive decision trees, see Chapter 8, using four different random inducer methods and then blending them using three different combination approaches. Moreover, we propose two new cost-sensitive combination approaches, cost-sensitive weighted voting and cost-sensitive stacking. The latter being an extension of our proposed cost-sensitive logistic regression, see Chapter 7.

The remainder of the section is organized as follows: First, we introduce the different random inducers used to create the base classifiers. Then, we present the combination methods. Finally, we define our proposed algorithms.

9.3.1 Random inducers

With the objective of creating an ensemble of example-dependent cost-sensitive decision trees, we first create T different random subsamples S_j for $j = 1, \dots, T$, of the training set S , and train a CSDT algorithm on each one. In particular we create the different subsets using four different methods: bagging [Breiman, 1996], pasting [Breiman, 1999], random forests [Breiman, 2001] and random patches [Louppe and Geurts, 2012].

In bagging [Breiman, 1996], base classifiers are built on randomly drawn bootstrap subsets of the original data, hence producing different base classifiers. Similarly, in pasting [Breiman, 1999], the base classifiers are built on random samples without replacement from the training




Figure 9.2: Visual representation of the random inducers algorithms.

set. In random forests [Breiman, 2001], using decision trees as the base learner, bagging is extended and combined with a randomization of the input features that are used when considering candidates to split internal nodes. In particular, instead of looking for the best split among all features, the algorithm selects, at each node, a random subset of features and then determines the best split only over these features. In the random patches algorithm [Louppe and Geurts, 2012], base classifiers are created by randomly drawn bootstrap subsets of both examples and features. To further clarify the difference between the random inducer methods, in Figure 9.2, we show a visual representation of the random inducers algorithms.

9.3.2 Combination methods

Lastly, the base classifiers are combined using either majority voting, cost-sensitive weighted voting and cost-sensitive stacking. Majority voting consists in collecting the predictions of each base classifier and selecting the decision with the highest number of votes, see (9.1).

9.3.2.1 Cost-sensitive weighted voting

This method is an extension of weighted voting. First, in the traditional approach, a similar comparison of the votes of the base classifiers is

made, but giving a weight α_j to each classifier M_j during the voting phase [Zhou, 2012]

$$f_{\text{vv}}(\mathcal{S}, \mathcal{M}, \alpha) = \arg \max_{c \in \{0,1\}} \sum_{j=1}^T \alpha_j \mathbf{1}_c(M_j(\mathcal{S})), \quad (9.5)$$

where $\alpha = \{\alpha_j\}_{j=1}^T$. The calculation of α_j is related to the performance of each classifier M_j . It is usually defined as the normalized misclassification error ϵ of the base classifier M_j in the out of bag set $\mathcal{S}_j^{\text{ob}} = \mathcal{S} - \mathcal{S}_j$

$$\alpha_j = \frac{1 - \epsilon(M_j(\mathcal{S}_j^{\text{ob}}))}{\sum_{j=1}^T 1 - \epsilon(M_{j_1}(\mathcal{S}_{j_1}^{\text{ob}}))}. \quad (9.6)$$

However, as previously discussed, the misclassification measure is not suitable in many real-world classification problems. We herein propose a method to calculate the weights α_j taking into account the actual savings of the classifiers. Therefore using (3.7), we define

$$\alpha_j = \frac{\text{Savings}(M_j(\mathcal{S}_j^{\text{ob}}))}{\sum_{j=1}^T \text{Savings}(M_{j_1}(\mathcal{S}_{j_1}^{\text{ob}}))}. \quad (9.7)$$

This method guaranties that the base classifiers that contribute to a higher increase in savings have more importance in the ensemble.

9.3.2.2 Cost-sensitive stacking

The staking method consists in combining the different base classifiers by learning a second level algorithm on top of them [Wolpert, 1992]. In this framework, once the base classifiers are constructed using the training set \mathcal{S} , a new set is constructed where the output of the base classifiers are now considered as the features while keeping the class labels.

Even though there is no restriction on which algorithm can be used as a second level learner, it is common to use a linear model [Zhou, 2012], such as

$$f_s(\mathcal{S}, \mathcal{M}, \beta) = g \left(\sum_{j=1}^T \beta_j M_j(\mathcal{S}) \right), \quad (9.8)$$

where $\beta = \{\beta_j\}_{j=1}^T$, and $g(\cdot)$ is the sign function $g(z) = \text{sign}(z)$ in the case of a linear regression or the sigmoid function, defined as $g(z) = 1/(1 + e^{-z})$, in the case of a logistic regression.

Moreover, following the logic used in [Nesbitt, 2010], we propose learning the set of parameters β using our cost-sensitive logistic regression (CSLR), see Chapter 7. The CSLR algorithm consists in introducing example-dependent costs into a logistic regression, by changing the objective function of the model to one that is cost-sensitive. For the specific case of cost-sensitive stacking, we define the cost function as:

$$J(\mathcal{S}, \mathcal{M}, \beta) = \sum_{i=1}^N \left[y_i \left(f_s(\mathbf{x}_i, \mathcal{M}, \beta) \cdot (C_{TP_i} - C_{FN_i}) + C_{FN_i} \right) + (1 - y_i) \left(f_s(\mathbf{x}_i, \mathcal{M}, \beta) \cdot (C_{FP_i} - C_{TN_i}) + C_{TN_i} \right) \right]. \quad (9.9)$$

Then, the parameters β that minimize the logistic cost function are used in order to combine the different base classifiers. However, as discussed in Chapter 7, this cost function is not convex for all possible cost matrices, therefore, we use genetic algorithms to minimize it.

Similarly to cost-sensitive weighting, this method guarantees that the base classifiers that contribute to a higher increase in savings have more importance in the ensemble. Furthermore, by learning an additional second level cost-sensitive method, the combination is made such that the overall savings measure is maximized.

9.3.3 Algorithms

Finally, Algorithm 9.1 summarizes the ECSDT methods. In total, we evaluate 12 different algorithms, as four different random inducers (bagging, pasting, random forest and random patches) and three different combinators (majority voting, cost-sensitive weighted voting and cost-sensitive stacking) can be selected in order to construct the cost-sensitive ensemble.

Algorithm 9.1. *The proposed ECSDT algorithms.*

Require: CSDT (an example-dependent cost-sensitive decision tree algorithm), T the number of iterations, \mathcal{S} the training set, inducer, N_e number of examples for each base classifier, N_f number of examples for each base classifier, combinator.

- 1: **function** ECSDT($\mathcal{S}, T, \text{inducer}, N_e, N_f, \text{combinator}$)
- 2: **Step 1:** Create the set of base classifiers
- 3: **for** $j \leftarrow 1$ to T **do**
- 4: **switch**(inducer)
- 5: **case** Bagging:
- 6: $\mathcal{S}_j \leftarrow$ Sample N_e examples from \mathcal{S} with replacement.
- 7: **case** Pasting:
- 8: $\mathcal{S}_j \leftarrow$ Sample N_e examples from \mathcal{S} without replacement.

```

9:      case Random forests:
10:       $\mathcal{S}_j \leftarrow$  Sample  $N_e$  examples from  $\mathcal{S}$  with replacement.
11:      case Random patches:
12:       $\mathcal{S}_j \leftarrow$  Sample  $N_e$  examples and  $N_f$  features from  $\mathcal{S}$  with re-
placement.
13:      end switch
14:       $M_j \leftarrow \text{CSDT}(\mathcal{S}_j)$ 
15:       $\mathcal{S}_j^{\text{ob}} \leftarrow \mathcal{S} - \mathcal{S}_j$ 
16:       $\alpha_j \leftarrow \text{Savings}(M_j(\mathcal{S}_j^{\text{ob}}))$ 
17:      end for
18:      Step 2: Combine the different base classifiers
19:      switch(combinator)
20:      case Majority voting:
21:       $H \leftarrow f_{\text{mv}}(\mathcal{S}, \mathcal{M})$ 
22:      case Cost-sensitive weighted voting:
23:       $H \leftarrow f_{\text{wv}}(\mathcal{S}, \mathcal{M}, \alpha)$ 
24:      case Cost-sensitive stacking:
25:       $\beta \leftarrow \arg \min_{\beta \in \mathbb{R}^T} J(\mathcal{S}, \mathcal{M}, \beta)$ 
26:       $H \leftarrow f_s(\mathcal{S}, \mathcal{M}, \beta)$ 
27:      end switch
28:      return  $H$ 
29: end function

```

9.4 THEORETICAL ANALYSIS OF THE ECSDT

Although the above proposed algorithm is simple, there is no work that has formally investigated ensemble performance in terms other than accuracy. In this section, our aim is to prove theoretically that combining individual cost-sensitive classifiers achieves better results in the sense of higher savings.

We denote \mathcal{S}_a $a \in \{0, 1\}$, as the subset of \mathcal{S} where the examples belong to the class a :

$$\mathcal{S}_a = \{\mathbf{x}_i^* | y_i = a, i \in 1, \dots, N\}, \quad (9.10)$$

where $\mathcal{S} = \mathcal{S}_0 \cup \mathcal{S}_1$, $\mathcal{S}_0 \cap \mathcal{S}_1 = \emptyset$, and $N_a = |\mathcal{S}_a|$. Also, we define the average cost of the base classifiers as:

$$\overline{\text{Cost}}(\mathcal{M}(\mathcal{S})) = \frac{1}{T} \sum_{j=1}^T \text{Cost}(M_j(\mathcal{S})). \quad (9.11)$$

Firstly, we prove the following lemma that states the cost of an ensemble H on the subset \mathcal{S}_a is lower than the average cost of the base classifiers on the same set for $a \in \{0, 1\}$.

Lemma 9.1. *Let H be an ensemble of $T \geq 3$ classifiers $\mathcal{M} = \{M_1, M_2, \dots, M_T\}$, and \mathcal{S} a testing set of size N . If each one of the base classifiers has a probability of being correct higher or equal than one half, $\rho \geq \frac{1}{2}$, and the reasonableness conditions of the cost matrix are satisfied, then the following holds true*

$$\text{Cost}(H(\mathcal{S}_a)) \leq \overline{\text{Cost}}(\mathcal{M}(\mathcal{S}_a)), \quad a \in \{0, 1\}, \quad (9.12)$$

Proof. First, we decompose the total cost of the ensemble by applying equations (3.2) and (3.3). Additionally, we separate the analysis for $a = 0$ and $a = 1$:

- $a = 0$:

$$\begin{aligned} \text{Cost}(H(\mathcal{S}_0)) &= \sum_{i=1}^{N_0} y_i (c_i C_{TP_i} + (1 - c_i) C_{FN_i}) + \\ &\quad (1 - y_i) (c_i C_{FP_i} + (1 - c_i) C_{TN_i}). \end{aligned} \quad (9.13)$$

Moreover, we know from (9.2) that the probability of an ensemble making the right decision, i.e., $y_i = c_i$, for any given example, is equal to P_c . Therefore, we can use this probability to estimate the expected savings of an ensemble:

$$\text{Cost}(H(\mathcal{S}_0)) = \sum_{i=1}^{N_0} P_c C_{TN_i} + (1 - P_c) C_{FP_i}. \quad (9.14)$$

- $a = 1$:

In the case of \mathcal{S}_1 , and following the same logic as when $a = 0$, the cost of an ensemble is:

$$\text{Cost}(H(\mathcal{S}_1)) = \sum_{i=1}^{N_1} P_c C_{TP_i} + (1 - P_c) C_{FN_i}. \quad (9.15)$$

The second part of the proof consists in analyzing the right hand side of (9.12), specifically, the average cost of the base classifiers on set \mathcal{S}_a . To do that, with the help of (3.2) and (9.11), we may express the average cost of the base classifiers as:

$$\overline{\text{Cost}}(\mathcal{M}(\mathcal{S}_a)) = \frac{1}{T} \sum_{j=1}^T \sum_{i=1}^{N_a} \text{Cost}(M_j(\mathbf{x}_i^*)). \quad (9.16)$$

We define the set of base classifiers that make a negative prediction as

$$\mathcal{T}_{i0} = \{M_j(\mathbf{x}_i^*) | M_j(\mathbf{x}_i^*) = 0, j \in 1, \dots, T\}, \quad (9.17)$$

similarly, the set of classifiers that make a positive prediction as

$$\mathcal{T}_{i1} = \{M_j(\mathbf{x}_i^*) | M_j(\mathbf{x}_i^*) = 1, j \in 1, \dots, T\}. \quad (9.18)$$

Then, by taking the cost of negative and positive predictions from (3.6), the average cost of the base learners becomes:

$$\begin{aligned} \overline{\text{Cost}}(\mathcal{M}(\mathcal{S}_a)) &= \frac{1}{T} \sum_{i=1}^{N_a} \left(|\mathcal{T}_{i0}| \cdot \text{Cost}(f_0(\mathbf{x}_i^*)) \right. \\ &\quad \left. + |\mathcal{T}_{i1}| \cdot \text{Cost}(f_1(\mathbf{x}_i^*)) \right). \end{aligned} \quad (9.19)$$

We separate the analysis for $a = 0$ and $a = 1$:

- $a = 0$:

$$\overline{\text{Cost}}(\mathcal{M}(\mathcal{S}_0)) = \sum_{i=1}^{N_0} \left(\frac{|\mathcal{T}_{i0}|}{T} \cdot C_{\text{TN}_i} + \frac{|\mathcal{T}_{i1}|}{T} \cdot C_{\text{FP}_i} \right). \quad (9.20)$$

Furthermore, we know from (9.2) that an average base classifier will have a correct classification probability of ρ , then $\frac{|\mathcal{T}_{i0}|}{T} = \rho$, leading to:

$$\overline{\text{Cost}}(\mathcal{M}(\mathcal{S}_0)) = \sum_{i=1}^{N_0} \left(\rho \cdot C_{\text{TN}_i} + (1 - \rho) \cdot C_{\text{FP}_i} \right). \quad (9.21)$$

- $a = 1$:

Similarly, for the set \mathcal{S}_1 , the average classifier will have a correct classification probability of ρ , then $\frac{|\mathcal{T}_{i1}|}{T} = \rho$.

Therefore,

$$\overline{\text{Cost}}(\mathcal{M}(\mathcal{S}_1)) = \sum_{i=1}^{N_1} \left(\rho \cdot C_{\text{TP}_i} + (1 - \rho) \cdot C_{\text{FN}_i} \right). \quad (9.22)$$

Finally, by replacing in (9.12) the expected savings of an ensemble with (9.14) for $a = 0$ and (9.15) for $a = 1$, and the average cost of the base learners with (9.21) for $a = 0$ and (9.22) for $a = 1$, (9.12) is rewritten as: for $a = 0$:

$$\sum_{i=1}^{N_0} \left(P_c C_{\text{TN}_i} + (1 - P_c) C_{\text{FP}_i} \right) \leq \sum_{i=1}^{N_0} \left(\rho \cdot C_{\text{TN}_i} + (1 - \rho) \cdot C_{\text{FP}_i} \right), \quad (9.23)$$

for $a = 1$:

$$\sum_{i=1}^{N_1} \left(P_c C_{\text{TP}_i} + (1 - P_c) C_{\text{FN}_i} \right) \leq \sum_{i=1}^{N_1} \left(\rho \cdot C_{\text{TP}_i} + (1 - \rho) \cdot C_{\text{FN}_i} \right). \quad (9.24)$$

Since $\rho \geq \frac{1}{2}$, then $P_c \geq \rho$ from (9.4), and using the *reasonableness conditions* described in Section 3.3, i.e, $C_{FP_i} > C_{TN_i}$ and $C_{FN_i} > C_{TP_i}$, we find that (9.23) and (9.24) are True. \square

Lemma 9.1 separates the costs on sets S_0 and S_1 . We are interested in analyzing the overall savings of an ensemble. In this direction, we demonstrate in the following theorem, that the expected savings of an ensemble of classifiers are higher than the expected average savings of the base learners.

Theorem 9.2. *Let H be an ensemble of $T \geq 3$ classifiers $\mathcal{M} = \{M_1, \dots, M_T\}$, and S a testing set of size N , then the expected savings of using H in S are lower than the average expected savings of the base classifiers, in other words,*

$$\text{Savings}(H(S)) \geq \overline{\text{Savings}}(\mathcal{M}(S)). \quad (9.25)$$

Proof. Given (3.7), (9.25) is equivalent to

$$\text{Cost}(H(S)) \leq \overline{\text{Cost}}(\mathcal{M}(S)). \quad (9.26)$$

Afterwards, by applying the cost definition of (3.3), and grouping the sets of negative and positive examples using (9.10), (9.26) becomes

$$\sum_{a \in \{0,1\}} \text{Cost}(H(S_a)) \leq \sum_{a \in \{0,1\}} \overline{\text{Cost}}(\mathcal{M}(S_a)), \quad (9.27)$$

which can be easily proved using Lemma 9.1, since, if the cost of an ensemble H is lower than the average cost of the base classifiers on both S_0 and S_1 , implies that it is also lower on the sum of the cost on both sets, therefore, proving Theorem 9.2. \square

9.5 EXPERIMENTS

For the experiments we use five datasets from four different real world example-dependent cost-sensitive problems: Credit card fraud detection (see Section 4.1), credit scoring (see Section 4.2), churn modeling (see Section 5.1) and direct marketing (see Section 5.2). The different datasets are summarized in Table 4.7 and Table 5.3.

We first use three classification algorithms, decision tree (DT), logistic regression (LR) and random forest (RF). Using the implementation of *Scikit-learn* [Pedregosa et al., 2011], each algorithm is trained using the different training sets: training (t), under-sampling (u), cost-proportionate rejection-sampling (r) [Zadrozny et al., 2003] and cost-proportionate over-sampling (o) [Elkan, 2001]. Afterwards, we evaluate

Family	Algorithm	Fraud	Churn	Credit 1	Credit 2	Marketing
CI	DT-t	0.3176±0.0357	-0.0018±0.0194	0.1931±0.0087	-0.0616±0.0229	-0.2342±0.0609
	LR-t	0.0092±0.0002	-0.0001±0.0002	0.0177±0.0126	0.0039±0.0012	-0.2931±0.0602
	RF-t	0.3342±0.0156	-0.0026±0.0079	0.1471±0.0071	0.0303±0.0040	-0.2569±0.0637
	DT-u	0.5239±0.0118	-0.0389±0.0583	0.3287±0.0125	-0.1893±0.0314	-0.0278±0.0475
	LR-u	0.1243±0.0387	0.0039±0.0492	0.4118±0.0313	0.1850±0.0231	0.2200±0.0376
	RF-u	0.5684±0.0097	0.0433±0.0533	0.4981±0.0079	0.1237±0.0228	0.1227±0.0443
CPS	DT-r	0.3439±0.0453	0.0054±0.0568	0.3310±0.0126	0.0724±0.0212	0.1960±0.0527
	LR-r	0.3077±0.0301	0.0484±0.0375	0.3965±0.0263	0.2650±0.0115	0.4210±0.0267
	RF-r	0.3812±0.0264	0.1056±0.0412	0.4989±0.0080	0.3055±0.0106	0.3840±0.0360
	DT-o	0.3172±0.0274	0.0251±0.0195	0.1738±0.0092	0.0918±0.0225	-0.2598±0.0559
	LR-o	0.2793±0.0185	0.0316±0.0228	0.3301±0.0109	0.2554±0.0090	0.3129±0.0277
	RF-o	0.3612±0.0295	0.0205±0.0156	0.2128±0.0081	0.2242±0.0070	-0.1782±0.0618
BMR	DT-t-BMR	0.6045±0.0386	0.0298±0.0145	0.1054±0.0358	0.2740±0.0067	0.4598±0.0089
	LR-t-BMR	0.4552±0.0203	0.1082±0.0316	0.2189±0.0541	0.3148±0.0094	0.4973±0.0084
	RF-t-BMR	0.6414±0.0154	0.0856±0.0354	0.4924±0.0087	0.3133±0.0094	0.4807±0.0093

(Models with the highest savings are marked in bold)

Table 9.1: Results of the CI and BMR algorithms measured by savings

the results of the algorithms using BMR, see Chapter 6. Then, the cost-sensitive logistic regression (CSLR), see Chapter 7, and cost-sensitive decision tree (CSDT), see Chapter 8, are also evaluated. Lastly, we calculate the propose ensembles of cost-sensitive decision trees algorithms. In particular, using each of the random inducer methods, bagging (CSB), pasting (CSP), random forests (CSRF) and random patches (CSRP), and then blending the base classifiers using each one of the combination methods; majority voting (mv), cost-sensitive weighted voting (wv) and

Family	Algorithm	Fraud	Churn	Credit 1	Credit 2	Marketing
CST	CSLR-t	0.6113±0.0262	0.1118±0.0484	0.4554±0.1039	0.2748±0.0069	0.4484±0.0072
	CSDT-t	0.7116±0.2557	0.1115±0.0378	0.4829±0.0098	0.2835±0.0078	0.4741±0.0063
ECSDT	CSB-mv-t	0.7124±0.0162	0.1237±0.0368	0.4862±0.0102	0.2945±0.0105	0.4837±0.0078
	CSB-wv-t	0.7276±0.0116	0.1539±0.0255	0.4862±0.0102	0.2948±0.0106	0.4838±0.0079
	CSB-s-t	0.7181±0.0109	0.1441±0.0364	0.4847±0.0096	0.2856±0.0088	0.4769±0.0078
	CSP-mv-t	0.7106±0.0113	0.1227±0.0399	0.4853±0.0104	0.2919±0.0097	0.4831±0.0081
	CSP-wv-t	0.7244±0.0202	0.1501±0.0302	0.4854±0.0105	0.2921±0.0098	0.4832±0.0082
	CSP-s-t	0.7212±0.0067	0.1488±0.0272	0.4848±0.0084	0.2870±0.0084	0.4752±0.0089
	CSRF-mv-t	0.6498±0.0598	0.0300±0.0488	0.4980±0.0120	0.2274±0.0520	0.3929±0.0655
	CSRF-wv-t	0.7249±0.0742	0.0624±0.0477	0.4979±0.0124	0.2948±0.0079	0.4728±0.0125
	CSRF-s-t	0.6731±0.0931	0.0586±0.0507	0.4839±0.0160	0.2518±0.0281	0.3854±0.0899
	CSRP-mv-t	0.7220±0.0082	0.1321±0.0280	0.5154±0.0077	0.3053±0.0087	0.4960±0.0075
	CSRP-wv-t	0.7348±0.0131	0.1615±0.0252	0.5152±0.0083	0.3015±0.0086	0.4885±0.0076
	CSRP-s-t	0.7336±0.0108	0.1652±0.0264	0.4989±0.0088	0.2956±0.0078	0.4878±0.0080

(Models with the highest savings are marked in bold)

Table 9.2: Results of the CST and ECSDT algorithms measured by savings

cost-sensitive stacking (s). Unless otherwise stated, the random selection of the training set was repeated 50 times, and in each time the models were trained and results collected, this allows us to measure the stability of the results. The implementation of the cost-sensitive algorithms is done using the *CostCla* library, see Appendix A.

The results are shown in Table 9.1 and Table 9.2. First, when observing the results of the cost-insensitive methods (CI), that is, DT, LR and RF algorithms trained on the t and u sets, the RF algorithm produces the best result by savings in three out of the five sets, followed by the LR – u.

Family	Algorithm	F-Rank	perBest
ECSDT	CSRP-wv-t	2.6	98.35
ECSDT	CSRP-s-t	3.4	97.72
ECSDT	CSRP-mv-t	4.0	94.99
ECSDT	CSB-wv-t	5.6	95.49
ECSDT	CSP-wv-t	7.4	94.72
ECSDT	CSB-mv-t	8.2	91.39
ECSDT	CSRF-wv-t	9.4	84.35
BMR	RF-t-BMR	9.4	86.16
ECSDT	CSP-s-t	9.6	93.80
ECSDT	CSP-mv-t	10.2	91.00
ECSDT	CSB-s-t	10.2	93.12
BMR	LR-t-BMR	11.2	73.98
CPS	RF-r	11.6	77.37
CST	CSDT-t	12.6	88.69
CST	CSLR-t	14.4	83.34
ECSDT	CSRF-mv-t	15.2	70.88
ECSDT	CSRF-s-t	16.0	75.68
CI	RF-u	17.2	52.83
CPS	LR-r	19.0	63.39
BMR	DT-t-BMR	19.0	60.05
CPS	LR-o	21.0	53.05
CPS	DT-r	22.6	35.33
CI	LR-u	22.8	40.43
CPS	RF-o	22.8	34.81
CI	DT-u	24.4	27.01
CPS	DT-o	25.0	24.25
CI	DT-t	26.0	16.14
CI	RF-t	26.2	16.73
CI	LR-t	28.0	1.19

Table 9.3: Savings Friedman ranking and average percentage of best result

It is also clear that the results on the t dataset are not as good as the ones on the u, this is highly related to the unbalanced distribution of the positives and negatives in all the databases.

In the case of cost-proportionate sampling methods (CPS), specifically the cost-proportionate rejection sampling (r) and cost-proportionate over sampling (o). It is observed that in four cases the savings increases quite significantly. It is on the fraud detection database where these methods do not outperform the algorithms trained on the under-sampled set. This may be related to the fact that in this database the initial percentage of positives is 1.5% which is similar to the percentage

Algorithm	Fraud	Churn	Credit1	Credit2	Marketing
RF-u	17	18	5	23	23
RF-r	20	13	3	3	19
RF-t-BMR	14	14	8	2	9
CSDT-t	10	11	16	14	12
CSRP-wv-t	1	2	2	5	3

Table 9.4: Savings ranks of best algorithm of each family by database





Figure 9.3: **Comparison of the savings of the algorithms versus the highest savings in each database.** The CSRP – wt is very close to the best result in all the databases. Additionally, even though the LR – BMR is the best algorithm in two databases, the performance in the other three is very poor.


in the r and o sets. However it is 50.42% in the u set, which may help explain why this method performs much better as measured by savings.

Afterwards, in the case of the BMR algorithms, the results show that this method outperforms the previous ones in four cases and has almost the same result in the other set. In the fraud detection set, the results are quite better, since the savings of the three classification algorithms increase when using this methodology. The next family of algorithms is the cost-sensitive training, which includes the CSLR and CSDT techniques. In this case, only in two databases the results are improved. Lastly, we evaluate the proposed ECSDT algorithms. The results show that these methods arise to the best overall results in three sets, while being quite competitive in the others.

Subsequently, in order to statistically sort the classifiers we computed the Friedman ranking (F-Rank) statistic [Demšar, 2006]. This rank increases with the cost of the algorithms. We also calculate the average savings of each algorithm compared with the highest savings in each




(a) Comparison of the Friedman ranking.




(b) Comparison of the average savings of the algorithms versus the highest savings.

Figure 9.4: **Comparison of the results by family of classifiers.** The ECSDT family has the best performance measured either by Friedman ranking or average percentage of best model.

set (perBest), as a measure of how close are the savings of an algorithm to the best result. In Table 9.3, the results are shown. It is observed that the first six algorithms, according to the F-Rank, belong to the ECSDT family. In particular, the best three classifiers is the ensemble of cost-sensitive decision trees using the random patches approach. Giving the best result the one that blends the base classifiers using weighted voting method. Moreover as shown in Table 9.4, this method ranks on each dataset 1st, 2nd, 2nd, 5th and 3rd, respectively. For comparison the best method from an other family is the RF with BMR, which ranks 14th, 14th, 8th, 2nd and 9th.



(a) Comparison by inducer methodology.



(b) Comparison by combination of base classifiers approach.

Figure 9.5: **Comparison of the Friedman ranking within the ECSDT family.**

Overall, the random inducer method that provides the best results is the CSRP. Moreover, the best combination method compared by Friedman ranking is the cost-sensitive weighted voting.

Moreover, when analyzing the perBest statistic, it is observed that it follows almost the same order as the F-Rank. Notwithstanding, there are cases in which algorithms ranks are different in the two statistics, for example the CSDT – t algorithm has a lower F-Rank than the RF – BMR, but the perBest is better. This happens because, the F-Rank does not take into account the difference in savings within algorithms. This can be further investigated in Figure 9.3. Even that ranks of the BMR models are better than the CSDT, the latter is on average closer to the best performance method in each set. Moreover, it is confirmed that the CSRP – wt is very close to the best result in all cases. Lastly, it is shown




Figure 9.6: **Comparison of the Friedman ranking of the savings and F1Score sorted by F1Score ranking.** The best two algorithms according to their Friedman rank of F1Score are indeed the best ones measured by the Friedman rank of the savings. However, this relation does not consistently hold for the other algorithms as the correlation between the rankings is just 65.10%.

why the F-Rank of the LR – BMR is high, given the fact that is the best model in two databases. The reason for that, is because the performance on the other sets is very poor.

Furthermore Figure 9.4a, shows the Friedman ranking of each family of classifiers. The ECSDT methods are overall better, followed by the BMR and the CST methods. As expected, the CI family is the one that performs the worst. Nevertheless, there is a significant variance within the ranks in the ECSDT family, as the best one has a Friedman ranking of 2.6 and the worst 16. Similar results are found when observing the perBest shown in Figure 9.4b. However, in the case of the perBest, the CST methods perform better than the BMR. It is important, in both cases it is confirmed that the ECSDT family of methods is the one that arise to the best results as measured by savings.

We further investigate the different methods that compose the ECSDT family, first by inducer methods and by the combination approach. In Figure 9.5a, the Friedman ranking of the ECSDT methods grouped by inducer algorithm are shown. It is observed that the worst method is the random forest methodology. This may be related to the fact that within the random inducer methods, this is the only one that also modified the learner algorithm in the sense that it randomly select features for each step during the decision tree growing. Moreover, as expected the bagging and pasting methods perform quite similar, after all the only difference is that in bagging the sampling is done with replacement,

Family	Algorithm	Fraud	Churn	Credit 1	Credit 2	Marketing
CI	DT-t	0.4458±0.0133	0.0733±0.0198	0.2593±0.0068	0.2614±0.0083	0.2647±0.0079
	LR-t	0.1531±0.0045	0.0000±0.0000	0.0494±0.0277	0.0155±0.0037	0.2702±0.0125
	RF-t	0.2061±0.0041	0.0249±0.0146	0.2668±0.0085	0.0887±0.0061	0.2884±0.0116
	DT-u	0.1502±0.0066	0.1175±0.0103	0.2276±0.0044	0.3235±0.0055	0.2659±0.0061
	LR-u	0.0241±0.0163	0.1222±0.0098	0.3160±0.0314	0.3890±0.0053	0.3440±0.0083
	RF-u	0.0359±0.0065	0.1346±0.0112	0.3193±0.0053	0.3815±0.0051	0.3088±0.0065
CPS	DT-r	0.4321±0.0086	0.1206±0.0132	0.2310±0.0049	0.3409±0.0046	0.2739±0.0076
	LR-r	0.1846±0.0123	0.1258±0.0111	0.3597±0.0156	0.3793±0.0049	0.3374±0.0101
	RF-r	0.2171±0.0100	0.1450±0.0131	0.3361±0.0067	0.3570±0.0048	0.3103±0.0075
	DT-o	0.4495±0.0063	0.1022±0.0180	0.2459±0.0081	0.3258±0.0056	0.2634±0.0089
	LR-o	0.1776±0.0117	0.1085±0.0203	0.3769±0.0067	0.3804±0.0044	0.3568±0.0102
	RF-o	0.2129±0.0080	0.0841±0.0201	0.3281±0.0078	0.3212±0.0054	0.3083±0.0093
BMR	DT-t-BMR	0.2139±0.0215	0.0941±0.0157	0.1514±0.0390	0.3338±0.0052	0.2433±0.0071
	LR-t-BMR	0.1384±0.0044	0.1370±0.0150	0.1915±0.0340	0.3572±0.0045	0.2954±0.0079
	RF-t-BMR	0.2052±0.0183	0.1264±0.0156	0.3186±0.0072	0.3551±0.0053	0.2744±0.0070

(Models with the highest F₁ Score are marked in bold)

Table 9.5: Results of the CI and BMR algorithms measured by F₁ Score

while it is not the case in pasting. In general the best methodology is random patches. Additionally, in Figure 9.5b, a similar analysis is made taking into account the combination of base classifiers approach. In this case, the best combination method is weighted voting, while majority voting and staking have a similar performance.

Finally, in Table 9.5 and Table 9.6 the results of the algorithms measured by F₁Score are shown. It is observed that the model with the highest savings is not the same as the one with the highest F₁Score in all of the databases as selecting a method by a traditional statistic does not

Family	Algorithm	Fraud	Churn	Credit 1	Credit 2	Marketing
CST	CSLR-t	0.2031±0.0065	0.1134±0.0151	0.1454±0.0517	0.3363±0.0045	0.2339±0.0051
	CSDT-t	0.2522±0.0980	0.1288±0.0194	0.2754±0.0059	0.3483±0.0046	0.2680±0.0060
	ECSDT					
	CSB-mv-t	0.2112±0.0125	0.1481±0.0122	0.2927±0.0108	0.3503±0.0046	0.2758±0.0072
	CSB-wv-t	0.2112±0.0091	0.1686±0.0125	0.2926±0.0108	0.3503±0.0046	0.2757±0.0072
	CSB-s-t	0.2072±0.0103	0.1554±0.0121	0.2818±0.0075	0.3533±0.0046	0.2799±0.0074
	CSP-mv-t	0.2098±0.0126	0.1480±0.0136	0.2903±0.0108	0.3499±0.0044	0.2749±0.0064
	CSP-wv-t	0.2099±0.0054	0.1651±0.0120	0.2905±0.0110	0.3498±0.0045	0.2749±0.0064
	CSP-s-t	0.2064±0.0069	0.1590±0.0099	0.2809±0.0062	0.3524±0.0046	0.2778±0.0104
	CSRF-mv-t	0.2208±0.0022	0.1081±0.0224	0.2994±0.0226	0.3560±0.0118	0.2780±0.0106
	CSRF-wv-t	0.2175±0.0019	0.1220±0.0234	0.2992±0.0236	0.3549±0.0069	0.2552±0.0187
	CSRF-s-t	0.2169±0.0045	0.1304±0.0162	0.2916±0.0236	0.3540±0.0071	0.2578±0.0186
	CSRP-mv-t	0.2691±0.0054	0.1511±0.0110	0.4031±0.0079	0.3743±0.0050	0.2916±0.0062
	CSRP-wv-t	0.2780±0.0041	0.1710±0.0140	0.4049±0.0066	0.3721±0.0051	0.2781±0.0063
	CSRP-s-t	0.2735±0.0148	0.1622±0.0103	0.3953±0.0141	0.3720±0.0049	0.2922±0.0137

(Models with the highest F₁ Score are marked in bold)

Table 9.6: Results of the CST and ECSDT algorithms measured by F₁ Score

give the same result as selecting it using a business oriented measure such as financial savings. This can be further examined in Figure 9.6, where the ranking of the F1Score and savings are compared. It is observed that the best two algorithms according to their Friedman rank of F1Score are indeed the best ones measured by the Friedman rank of the savings. However, this relation does not consistently hold for the other algorithms as the correlation between the rankings is just 65.10%.

CONCLUSIONS

This thesis deals with the problem of example-dependent cost-sensitive classification. Several real-world business applications of classification models are example-dependent cost-sensitive, in the sense that the objective of using an algorithm is related to maximizing the profit of the company. Moreover, the different costs due to misclassification vary among examples. Particularly, we focused on four different real-world applications: credit card fraud detection, credit scoring, churn modeling and direct marketing. In all cases, evaluating a classification algorithm using traditional statistics such as misclassification rate or F_1 Score, do not accurately represent the business oriented goals. Moreover, we found, significant differences in the results, when using traditionally cost-insensitive methods against example-dependent cost-sensitive methods.

First in Part I, we laid out the background for general classification methods, and their respective evaluation measures. Moreover, we introduced the problem of cost-sensitive learning. In particular, we defined the differences between class-dependent and example-dependent problems. In addition, we presented an evaluation measure that takes into account the real financial gains and losses of practical applications.

Then in Part II, we analyzed and discussed four real-world classification problems, that were the focus of this thesis, in particular, credit card fraud detection, credit scoring, churn modeling and direct marketing. In general, we showed why each of the applications is example-dependent cost-sensitive, and we elaborated a framework for the analysis of each problem.

Finally in Part III, we introduced our proposed example-dependent cost-sensitive methods. First, we presented a cost-sensitive Bayes minimum risk classifier. This method worked by comparing the expected financial losses of different outcomes when classifying the examples in the different classes. Then, we focused on introducing the costs to the algorithms during the training phase. In particular, we presented the cost-sensitive logistic regression and cost-sensitive decision trees algorithms. The cost-sensitive decision tree algorithm proved to be highly effective while maintaining the simplicity and interpretability of decision trees. However, this method suffers from high variance. To overcome this limitation, we proposed a framework for ensembles of cost-sensitive decision trees. We have shown theoretically and experimentally that the

method of ensembles of cost-sensitive decision trees ranks the best and outperforms state-of-the-art example-dependent cost-sensitive methodologies, when measured by financial savings.

This thesis showed the importance of using the real example-dependent financial costs associated with real-world applications. In particular, we found significant differences in the results when evaluating a model using a traditional cost-insensitive measure such as the accuracy or F1Score, than when using the savings, leading to the conclusion of the importance of using the real practical financial costs of each context. Lastly, in Appendix A, we presented the library *CostCla* that we developed as part of the thesis. This library is an open-source implementation of all the algorithms covered in this manuscript.


10.1 FUTURE RESEARCH DIRECTIONS

We foresee that the framework we developed through this thesis should open the door to developing more business focused algorithms, and that ultimately, the use of the actual financial costs during training will become a common practice. There is a list of points related to this work that can be further investigated. In the following, we address some issues.

- **Multi-class example-dependent cost-sensitive classification.** This thesis focused on binary cost-sensitive classification problems. Nevertheless, not all cost-sensitive applications are two-class problems. Some studies have started to look into Multi-class class-dependent cost-sensitive classification [Zhou and Liu, 2010]. Therefore, we expect that an interesting line of future work should include expanding our framework to multi-class problems.
- **Cost-sensitive calibration.** When evaluating how well a set of probabilities are calibrated, for example using the Brier score, the measure does not take into account the cost of each example. Therefore, the calibration methods that attempt to improve those kinds of measures, as the ones presented in Section 6.2, also fail to take into account the real cost-sensitive costs related to each application. Future work should include a deep analysis of the impact of the costs during the calibration of probabilities.
- **Staking cost-sensitive decision trees.** Even though we explore the method of staking in Section 9.3.2.2, we foresee that a deeper analysis of the impact of having several layers of example-dependent cost-sensitive methods could enhance the performance of the system.

- **Example-dependent cost-sensitive boosting.** In Chapter 9, we only focused our attention to independent ensemble methods, in particular bagging algorithms. However, there is an other branch of ensemble methods, those that are dependent, such as boosting [Schapire, 1990], adaboost [Freund and Schapire, 1996] or gradient boosting [Friedman, 2001, 2002]. For some applications these methods have proved to outperform the bagging algorithms [Zhou, 2012], therefore, we think that future research should focus on creating a framework for example-dependent cost-sensitive boosting algorithms.
- **Online example-dependent cost-sensitive classification.** The methods covered in this thesis are all batch methods, in the sense that the batch algorithms keep the system weights constant while calculating the evaluation measures. However, in some applications such as fraud detection, the evolving patterns due to change in the fraudsters behavior are not captured by using batch methods. Therefore, the need to investigate this problem from an online-learning perspective [Pozzolo et al., 2014b].

COSTCLA: A COST-SENSITIVE CLASSIFICATION LIBRARY IN PYTHON



CostCla is a **Python** open source cost-sensitive classification library built on top of *Scikit-learn*, *Pandas* and *Numpy*. *CostCla* provides a wide range of state-of-the-art example-dependent cost-sensitive methods for binary classification tasks. Source code, binaries and documentation are distributed under 3-Clause BSD license in the website <http://albahnsen.com/CostSensitiveClassification/>.

A.1 LIBRARY OVERVIEW

CostCla is a easy to use **Python** library for example-dependent cost-sensitive classification problems. It includes many example-dependent cost-sensitive algorithms. Since it is part of the scientific **Python** ecosystem, it can be easily integrated with other machine learning libraries.

Python offers several machine learning packages, some including classification algorithms like *Scikit-learn* [Pedregosa et al., 2011], *PyBrain* [Schaul and Felder, 2010], *PyMC* [Patil et al., 2010], *mlpy* [Albanese et al., 2012] and *Orange* [Demšar et al., 2013]. Nevertheless, all these packages are based on a cost-insensitive classification framework, in which all misclassification errors carry the same cost. As has been highlighted throughout this thesis, this is not the case in many real-world applications.

The core of the library consists of a number of state-of-the-art example-dependent cost-sensitive classification algorithms, such as cost-proportionate rejection-sampling [Zadrozny et al., 2003], cost-proportionate over-sampling [Elkan, 2001], Bayes minimum risk (see Chapter 6), cost-sensitive logistic regression (see Chapter 7), cost-sensitive decision trees (see Chapter 8), and the cost-sensitive ensemble

methods: cost-sensitive bagging, cost-sensitive pasting, cost-sensitive random forest and cost-sensitive random patches (see Chapter 9).

Moreover, the library also includes three different example-dependent cost-sensitive databases. In particular two credit scoring databases (see Section 4.2), and one direct marketing database (see Section 5.2).

A.2 USAGE

In this section we provided a quick example of the usage of the *CostCla* library, and compare the results of different algorithms using a credit scoring database.

- Prepare dataset and load libraries

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.cross_validation import train_test_split
from costcla.datasets import load_creditscoring2
from costcla.sampling import cost_sampling
from costcla.metrics import savings_score
from costcla import models
data = load_creditscoring2()
X_train, X_test, y_train, y_test,
cost_mat_train, cost_mat_test = \
train_test_split(data.data, data.target, data.cost_mat)
```

- Random forest

```
f_RF = RandomForestClassifier()
f_RF.fit(X_train, y_train)
y_pred = f_RF.predict(X_test)
print savings_score(y_test, y_pred, cost_mat_test)
0.042197359989
```

- Cost-proportionate rejection sampling

```
X_cps_r, y_cps_r, cost_mat_cps_r = \
cost_sampling(X_train, y_train, cost_mat_train,
              method='RejectionSampling')
f_RF.fit(X_train, y_train)
y_pred = f_RF.predict(X_test)
print savings_score(y_test, y_pred, cost_mat_test)
0.280743761779
```

- Bayes minimum risk

```
f_RF.fit(X_train, y_train)
y_prob_test = f_RF.predict_proba(X_test)
f_BMR = models.BayesMinimumRiskClassifier()
f_BMR.fit(y_test, y_prob_test)
y_pred = f_BMR.predict(y_prob_test, cost_mat_test)
print savings_score(y_test, y_pred, cost_mat_test)
0.285102564249
```

- Cost-sensitive decision tree

```
f_CSDT = models.CSDecisionTreeClassifier()
f_CSDT.fit(X_train, y_train, cost_mat_train)
y_pred = f_CSDT.predict(X_test)
print savings_score(y_test, y_pred, cost_mat_test)
0.289489571352
```

- Cost-sensitive random patches

```
f_CSRP = models.CSRandomPatchesClassifier()
f_CSRP.fit(X_train, y_train, cost_mat_train)
y_pred = f_CSRP.predict(X_test)
print savings_score(y_test, y_pred, cost_mat_test)
0.306607400467
```

In Table A.1 we summarize the results of the different algorithms. It is observed that the library follows the same API structure of *Scikit-learn*, therefore, making *CostCla* highly compatible with the most widely used machine learning library in **Python**.

Algorithm	Savings
Random forest	0.0422
Cost-proportionate rejection sampling	0.2807
Bayes minimum risk	0.2851
Cost-sensitive decision tree	0.2895
Cost-sensitive random patches	0.3066

Table A.1: Results of the different algorithms using the *CostCla* library

A.3 INSTALLATION

CostCla requires some prerequisite packages to be previously installed. Nevertheless, all are cross-platform and freely available online:

- **Python** version ≥ 2.7
- *Numpy* version $\geq 1.8.0$
- *Pandas* version $\geq 0.14.0$
- *Scikit-learn* version $\geq 0.15.0b2$
- *pyea* version ≥ 0.1

The easiest way to install *CostCla* is with *pip*:

```
pip install costcla
```

BIBLIOGRAPHY

- D. Albanese, R. Visintainer, S. Merler, S. Riccadonna, G. Jurman, and C. Furlanello. *mlpy: Machine Learning Python*. pages 1–4, Feb. 2012.
- R. Alejo and V. Garc. Making Accurate Credit Risk Predictions with Cost-Sensitive MLP Neural Networks. In J. Casillas, F. J. Martínez-López, R. Vicari, and F. De la Prieta, editors, *Advances in Intelligent Systems and Computing*, volume 220 of *Advances in Intelligent Systems and Computing*, pages 1–8. Springer International Publishing, Heidelberg, 2013.
- American Institute of CPAs. International Financial Reporting Standards (IFRS). Technical report, American Institute of CPAs, 2011.
- R. Anderson. *The Credit Scoring Toolkit : Theory and Practice for Retail Credit Risk Management and Decision Automation*. Oxford University Press, 2007.
- O. M. Aodha and G. J. Brostow. Revisiting Example Dependent Cost-Sensitive Learning with Decision Trees. In *2013 IEEE International Conference on Computer Vision*, pages 193–200, Washington, DC, USA, Dec. 2013.
- K. Bache and M. Lichman. UCI Machine Learning Repository. *School of Information and Computer Science, University of California*, 2013.
- S. Bachmayer. Artificial Immune Systems. *Artificial Immune Systems*, 5132:119–131, 2008.
- B. Baesens. *Analytics in a Big Data World: The Essential Guide to Data Science and its Applications*. Wiley, 2014.
- P. Beling, Z. Covaliu, and R. M. Oliver. Optimal scoring cutoff policies and efficient frontiers. *Journal of the Operational Research Society*, 56(9): 1016–1029, July 2005.
- S. Bhattacharyya, S. Jha, K. Tharakunnel, and J. C. Westland. Data mining for credit card fraud: A comparative study. *Decision Support Systems*, 50(3):602–613, Feb. 2011.
- C. M. Bishop. *Pattern Recognition and Machine Learning*, volume 4 of *Information science and statistics*. Springer, 2006.

- R. Bolton and D. J. Hand. Unsupervised profiling methods for fraud detection. In *Credit Scoring and Credit Control VII*, 2001.
- R. J. Bolton, D. J. Hand, F. Provost, and L. Breiman. Statistical Fraud Detection: A Review. *Statistical Science*, 17(3):235–255, 2002.
- R. Brause, T. Langsdorf, and M. Hepp. Neural data mining for credit card fraud detection. *Proceedings 11th International Conference on Tools with Artificial Intelligence*, pages 103–106, 1999.
- L. Breiman. Bagging predictors. *Machine Learning*, 24(2):123–140, Aug. 1996.
- L. Breiman. Pasting small votes for classification in large databases and on-line. *Machine Learning*, 103:85–103, 1999.
- L. Breiman. Random Forests. *Machine Learning*, 45:5–32, 2001.
- L. Breiman, J. Friedman, C. J. Stone, and R. Olshen. *Classification and Regression Trees*. Chapman and Hall/CRC, 1984.
- G. W. Brier. Monthly Weather Review. *Monthly weather review*, 78(1):1–3, Apr. 1950.
- I. Cohen and M. Goldszmidt. Properties and Benefits of Calibrated Classifiers. In *Knowledge Discovery in Databases: PKDD 2004*, volume 3202, pages 125–136. Springer Berlin Heidelberg, 2004.
- A. Correa Bahnsen and A. Gonzalez Montoya. Evolutionary Algorithms for Selecting the Architecture of a MLP Neural Network: A Credit Scoring Case. In *IEEE 11th International Conference on Data Mining Workshops*, pages 725–732. Ieee, Dec. 2011.
- A. Correa Bahnsen, A. Stojanovic, D. Aouada, and B. Ottersten. Cost Sensitive Credit Card Fraud Detection Using Bayes Minimum Risk. In *2013 12th International Conference on Machine Learning and Applications*, pages 333–338, Miami, USA, Dec. 2013. IEEE.
- A. Correa Bahnsen, D. Aouada, and B. Ottersten. Example-Dependent Cost-Sensitive Logistic Regression for Credit Scoring. In *2014 13th International Conference on Machine Learning and Applications*, pages 263–269, Detroit, USA, 2014a. IEEE.
- A. Correa Bahnsen, A. Stojanovic, D. Aouada, and B. Ottersten. Improving Credit Card Fraud Detection with Calibrated Probabilities. In *Proceedings of the fourteenth SIAM International Conference on Data Mining*, pages 677 – 685, Philadelphia, USA, 2014b.

- A. Correa Bahnsen, D. Aouada, and B. Ottersten. Example-Dependent Cost-Sensitive Decision Trees. *Expert Systems with Applications*, 42(19): 6609–6619, 2015a.
- A. Correa Bahnsen, D. Aouada, and B. Ottersten. A novel cost-sensitive framework for customer churn predictive modeling. *Decision Analytics*, 2(1):5, 2015b.
- A. Correa Bahnsen, D. Aouada, and B. Ottersten. Ensemble of Example-Dependent Cost-Sensitive Decision Trees. 2015c.
- A. Correa Bahnsen, D. Aouada, and B. Ottersten. Feature Engineering in Credit Card Fraud Detection. *Expert Systems with Applications*, submitted:1–29, 2015d.
- A. Correa Bahnsen, D. Aouada, and B. Ottersten. Detecting Credit Card Fraud using Periodic Features. In *International Conference on Machine Learning and Applications*, Miami, USA, 2015e. IEEE.
- M. H. DeGroot and S. E. Fienberg. The Comparison and Evaluation of Forecasters. *The Statistician*, 32(1):12–22, 1983.
- J. Demšar. Statistical Comparisons of Classifiers over Multiple Data Sets. *Journal of Machine Learning Research*, 7:1–30, 2006.
- J. Demšar, T. Curk, and A. Erjavec. Orange: data mining toolbox in Python. *The Journal of Machine Learning Research*, 14:2349–2353, 2013.
- T. Dietterich. Ensemble methods in machine learning. In *Multiple classifier systems*, pages 1–15. Springer, 2000.
- B. Draper, C. Brodley, and P. Utgoff. Goal-directed classification using linear machine decision trees. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 16(9):888–893, 1994.
- C. Drummond and R. Holte. C4.5, class imbalance, and cost sensitivity: why under-sampling beats over-sampling. In *Workshop on Learning from Imbalanced Datasets II, ICML*, Washington, DC, USA, 2003.
- ECB. European Central Bank, 2014.
- T. Economics. Brazil statistics, 2014.
- C. Elkan. The Foundations of Cost-Sensitive Learning. In *Seventeenth International Joint Conference on Artificial Intelligence*, pages 973–978, 2001.
- European Central Bank. Third report on card fraud. Technical report, European Central Bank, 2014.

- W. Fan, S. J. Stolfo, J. Zhang, and P. K. Chan. AdaCost: misclassification cost-sensitive boosting. In *In Proc. 16th International Conference on Machine Learning*, pages 97–105. Morgan Kaufmann Publishers, 1999.
- P. W. Farris, N. T. Bendle, P. E. Pfeifer, and D. J. Reibstein. *Marketing Metrics: The Definitive Guide to Measuring Marketing Performance*. Pearson FT Press, New Jersey, USA, 2nd edition, 2010.
- C. Ferri, P. Flach, and J. Hernández-Orallo. Learning decision trees using the area under the ROC curve. In *ICML '02 Proceedings of the Nineteenth International Conference on Machine Learning*, pages 139–146, San Francisco, CA, USA, 2002.
- N. I. Fisher. *Statistical Analysis of Circular Data*, volume 9. 1996.
- P. Flach, J. Hernandez-Orallo, and C. Ferri. Brier Curves : A New Cost-Based Visualisation of Classifier Performance. In *International Conference on Machine Learning*, pages 1–27, 2011.
- Y. Freund and R. R. E. Schapire. Experiments with a New Boosting Algorithm. In *International Conference on Machine Learning*, pages 148–156, 1996.
- J. Friedman. Greedy function approximation: a gradient boosting machine. *Annals of Statistics*, 2001.
- J. H. Friedman. Stochastic gradient boosting. *Computational Statistics and Data Analysis*, 38:367–378, 2002.
- N. Gladly, B. Baesens, and C. Croux. Modeling churn using customer lifetime value. *European Journal of Operational Research*, 197(1):402–411, Aug. 2009.
- D. J. Hand and W. E. Henley. Statistical Classification Methods in Consumer Credit Scoring: A Review. *Journal of the Royal Statistical Society. Series A (Statistics in Society)*, 160(3):523–541, 1997.
- D. J. Hand, C. Whitrow, N. M. Adams, P. Juszczak, and D. J. Weston. Performance criteria for plastic card fraud detection tools. *Journal of the Operational Research Society*, 59(7):956–962, May 2007.
- L. K. Hansen and P. Salamon. Neural network ensembles. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 12(October):993–1001, 1990.
- T. Hastie, R. Tibshirani, and J. Friedman. *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*. Springer, Stanford, CA, 2nd edition, 2009.

- R. L. Haupt and S. E. Haupt. *Practical genetic algorithms*. John Wiley & Sons, Inc., New Jersey, second edi edition, 2004.
- M. Herland, T. M. Khoshgoftaar, and R. Wald. A review of data mining using big data in health informatics. *Journal Of Big Data*, 1:2, 2014.
- J. Hernandez-Orallo, P. Flach, and C. Ferri. A Unified View of Performance Metrics : Translating Threshold Choice into Expected Classification Loss. *Journal of Machine Learning Research*, 13(July):2813–2869, 2012.
- J. V. Hulse and T. M. Khoshgoftaar. Experimental Perspectives on Learning from Imbalanced Data. In *International Conference on Machine Learning*, 2007.
- G. Jayanta K., D. Mohan, and S. Tapas. Bayesian Inference and Decision Theory. In *An Introduction to Bayesian Analysis*, volume 13, pages 26–63. Springer New York, Apr. 2006.
- S. KhakAbi, M. R. Gholamian, and M. Namvar. Data Mining Applications in Customer Churn Management. *2010 International Conference on Intelligent Systems, Modelling and Simulation*, pages 220–225, Jan. 2010.
- J. Kim, K. Choi, G. Kim, and Y. Suh. Classification cost: An empirical comparison among traditional classifier, Cost-Sensitive Classifier, and MetaCost. *Expert Systems with Applications*, 39(4):4013–4019, Mar. 2012.
- M. Kretowski and M. Grześ. Evolutionary induction of cost-sensitive decision trees. In *Foundations of Intelligent Systems*, pages 121–126. Springer Berlin Heidelberg, 2006.
- H. P. Kriegel, K. M. Borgwardt, P. Kröger, a. Pryakhin, M. Schubert, and a. Zimek. Future trends in data mining. *Data Mining and Knowledge Discovery*, 15:87–97, 2007.
- M. Krivko. A hybrid model for plastic card fraud detection systems. *Expert Systems with Applications*, 37(8):6070–6076, Aug. 2010.
- L. Lam and S. Y. Suen. Application of majority voting to pattern recognition: an analysis of its behavior and performance. *IEEE Transactions on Systems Man and Cybernetics Part A Systems and Humans*, 27(5):553–568, 1997.
- D. Lawrence and A. Solomon. *Managing a Consumer Lending Business*. Solomon Lawrence Partners, 2012.

- J. Li, X. Li, and X. Yao. Cost-Sensitive Classification with Genetic Programming. In *2005 IEEE Congress on Evolutionary Computation*, volume 3, pages 2114–2121. IEEE, 2005.
- Y. Li, J. Kwok, and Z. Zhou. Cost-Sensitive Semi-Supervised Support Vector Machine. In *Proceedings of the 24th AAAI*, pages 500–505, 2010.
- C. X. Ling, Q. Yang, J. Wang, and S. Zhang. Decision trees with minimal costs. In *Twenty-first international conference on Machine learning - ICML '04*, number Icml, page 69, New York, New York, USA, 2004. ACM Press.
- S. Lomax and S. Vadera. A survey of cost-sensitive decision tree induction algorithms. *ACM Computing Surveys*, 45(2):1–35, Feb. 2013.
- G. Louppe and P. Geurts. Ensembles on random patches. In *ECML PKDD'12 Proceedings of the 2012 European conference on Machine Learning and Knowledge Discovery in Databases*, pages 346–361. Springer Berlin Heidelberg, 2012.
- J. Ma, L. K. Saul, S. Savage, and G. M. Voelker. Learning to detect malicious URLs. *ACM Transactions on Intelligent Systems and Technology*, 2(3):1–24, Apr. 2011.
- S. Maes, K. Tuyls, B. Vanschoenwinkel, and B. Manderick. Credit card fraud detection using Bayesian and neural networks. In *Proceedings of NF2002*, 2002.
- N. Mahmoudi and E. Duman. Detecting credit card fraud by Modified Fisher Discriminant Analysis. *Expert Systems with Applications*, 42(5):2510–2516, 2015.
- L. R. O. Maimon. *Data Mining with Decision Trees: Theory and Applications*. World Scientific Publishing Company, 2008.
- S. Marslan. *Machine Learning: An Algorithmic Perspective*. CRC Press, New Jersey, USA, 2009.
- H. Masnadi-shirazi. Risk minimization, probability elicitation, and cost-sensitive SVMs. In *International Conference on Machine Learning*, 2010.
- H. Masnadi-shirazi and N. Vasconcelos. Cost-Sensitive Boosting. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 33(2):294–309, 2011.
- G. R. Milne and M.-E. Boza. Trust and Concern in Consumers' Perception of Marketing Information Management Practices. *Journal of Interactive Marketing*, 13(1):5–24, 1999.

- T. Minegishi and A. Niimi. Proposal of Credit Card Fraudulent Use Detection by Online-type Decision Tree Construction and Verification of Generality. *International Journal for Information Security Research (IJISR)*, 1(4):229–235, 2011.
- S. Moro, R. Laureano, and P. Cortez. Using data mining for bank direct marketing: An application of the crisp-dm methodology. In *European Simulation and Modelling Conference*, number Figure 1, pages 117–121, Guimares, Portugal, 2011.
- K. P. Murphy. *Machine Learning A Probabilistic Perspective*. MIT Press, 2012.
- G. N. Nayak and C. G. Turvey. Credit Risk Assessment and the Opportunity Costs of Loan Misclassification. *Canadian Journal of Agricultural Economics*, 45(3):285–299, 1997.
- T. A. Nesbitt. *Cost-Sensitive Tree-Stacking : Learning with Variable Prediction Error Costs*. PhD thesis, University of California, Los Angeles, 2010.
- S. A. Neslin, S. Gupta, W. Kamakura, J. Lu, and C. H. Mason. Defection Detection : Measuring and Understanding the Predictive Accuracy of Customer Churn Models. *Journal of Marketing Research*, 43(2):204–211, 2006.
- E. Ngai, L. Xiu, and D. Chau. Application of data mining techniques in customer relationship management: A literature review and classification. *Expert Systems with Applications*, 36(2):2592–2602, Mar. 2009.
- E. Ngai, Y. Hu, Y. Wong, Y. Chen, and X. Sun. The application of data mining techniques in financial fraud detection: A classification framework and an academic review of literature. *Decision Support Systems*, 50(3):559–569, Feb. 2011.
- R. Oliver and L. Thomas. Optimal score cutoffs and pricing in regulatory capital in retail credit portfolios. 2009.
- S. Panigrahi, A. Kundu, S. Sural, and A. Majumdar. Credit card fraud detection: A fusion approach using Dempster’s Shafer theory and Bayesian learning. *Information Fusion*, 10(4):354–363, Oct. 2009.
- A. Patil, D. Huard, and C. J. Fonnesbeck. PyMC: Bayesian Stochastic Modelling in Python. *Journal of statistical software*, 35(4):1–81, 2010.

- F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and E. Duchesnay. Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12:2825–2830, 2011.
- P. E. Pfeifer, M. E. Haskins, and R. M. Conroy. Customer lifetime value, customer profitability, and the treatment of acquisition spending. *Journal of Managerial Issues*, 17(1):11–25, 2004.
- A. D. Pozzolo, O. Caelen, Y.-A. Le Borgne, S. Waterschoot, and G. Bontempi. Learned lessons in credit card fraud detection from a practitioner perspective. *Expert Systems with Applications*, 41(10):4915–4928, Aug. 2014a.
- A. D. Pozzolo, R. Johnson, O. Caelen, S. Waterschoot, N. V. Chawla, and G. Bontempi. Using HDDT to avoid instance propagation in unbalanced and evolving data streams. In *The 2014 International Joint Conference on Neural Networks (IJCNN 2014)*, 2014b.
- F. Provost, T. Fawcett, and R. Kohavi. The case against accuracy estimation for comparing induction algorithms. In *Proceedings of the Fifteenth International Conference on Machine Learning*, pages 445–453. Morgan Kaufmann, 1998.
- J. R. Quinlan. *C4.5: Programs for Machine Learning*. Morgan Kaufmann, 1992.
- L. Rokach. Ensemble-based classifiers. *Artificial Intelligence Review*, 33(1-2):1–39, Nov. 2009.
- L. Rokach and O. Maimon. Decision Trees. In L. Rokach and O. Maimon, editors, *Data Mining and Knowledge Discovery Handbook*, chapter 9, pages 149–174. Springer US, 2nd edition, 2010.
- Y. Sahin, S. Bulkan, and E. Duman. A cost-sensitive decision tree approach for fraud detection. *Expert Systems with Applications*, 40(15):5916–5923, Nov. 2013.
- D. Sánchez, M. Vila, L. Cerda, and J. Serrano. Association rules applied to credit card fraud detection. *Expert Systems with Applications*, 36(2):3630–3640, Mar. 2009.
- R. E. Schapire. The strength of weak learnability. *Machine Learning*, 5(2):197–227, 1990.
- T. Schaul and M. Felder. PyBrain. *Journal of Machine Learning Research*, 11:743–746, 2010.

- V. Sheng and C. Ling. Thresholding for making classifiers cost-sensitive. In *Proceedings of the National Conference on Artificial Intelligence*, 2006.
- W. Street. Bagging with Adaptive Costs. *IEEE Transactions on Knowledge and Data Engineering*, 20(5):577–588, May 2008.
- D. Tasoulis and N. Adams. Mining information from plastic card transaction streams. In *Proceedings in 18th International Conference on Computational Statistics*, 2008.
- K. Ting. An instance-weighting method to induce cost-sensitive trees. *IEEE Transactions on Knowledge and Data Engineering*, 14(3):659–665, 2002.
- S. Vadera. CSNL: A cost-sensitive non-linear decision tree algorithm. *ACM Transactions on Knowledge Discovery from Data*, 4(2):1–25, 2010.
- E. M. van Raaij, M. J. a. Vernooij, and S. van Triest. The implementation of customer profitability analysis: A case study. *Industrial Marketing Management*, 32:573–583, 2003.
- V. Van Vlasselaer, C. Bravo, O. Caelen, T. Eliassi-Rad, L. Akoglu, M. Snoeck, and B. Baesens. APATE: A Novel Approach for Automated Credit Card Transaction Fraud Detection using Network-Based Extensions. *Decision Support Systems*, 75:38–48, 2015.
- W. Verbeke, K. Dejaeger, D. Martens, J. Hur, and B. Baesens. New insights into churn prediction in the telecommunication sector: A profit driven data mining approach. *European Journal of Operational Research*, 218(1):211–229, Apr. 2012.
- T. Verbraken. Toward profit-driven churn modeling with predictive marketing analytics. In *Cloud computing and analytics: innovations in e-business services. Workshop on E-Business (WEB2012)*, Orlando, US, 2012.
- T. Verbraken, W. Verbeke, and B. Baesens. A novel profit maximizing metric for measuring classification performance of customer churn prediction models. *IEEE Transactions on Knowledge and Data Engineering*, 25(5):961–973, 2013.
- T. Verbraken, C. Bravo, R. Weber, and B. Baesens. Development and application of consumer credit scoring models using profit-based classification measures. *European Journal of Operational Research*, 238(2): 505–513, Oct. 2014.

- J. Wang, P. Zhao, and S. C. H. Hoi. Cost-Sensitive Online Classification. *IEEE Transactions on Knowledge and Data Engineering*, 26(10):2425–2438, Oct. 2014.
- T. Wang. *Efficient Techniques for Cost-Sensitive Learning with Multiple Cost Considerations*. PhD thesis, University of Technology, Sydney, 2013.
- D. J. Weston, D. J. Hand, N. M. Adams, C. Whitrow, and P. Juszczak. Plastic card fraud detection using peer group analysis. *Advances in Data Analysis and Classification*, 2(1):45–62, Mar. 2008.
- C. Whitrow, D. J. Hand, P. Juszczak, D. J. Weston, and N. M. Adams. Transaction aggregation as a strategy for credit card fraud detection. *Data Mining and Knowledge Discovery*, 18(1):30–55, July 2008.
- D. H. Wolpert. Stacked generalization. *Neural Networks*, 5:241–259, 1992.
- H. F. Yu, F. L. Huang, and C. J. Lin. Dual coordinate descent methods for logistic regression and maximum entropy models. *Machine Learning*, 85(1-2):41–75, 2011.
- B. Zadrozny, J. Langford, and N. Abe. Cost-sensitive learning by cost-proportionate example weighting. In *Third IEEE International Conference on Data Mining*, pages 435–442. IEEE Comput. Soc, 2003.
- Z.-H. Zhou. *Ensemble Methods Foundations and Algorithms*. CRC Press, Boca Raton, FL, US, 2012.
- Z. H. Zhou and X. Y. Liu. On multi-class cost-sensitive learning. *Computational Intelligence*, 26(3):232–257, 2010.