

Fancy Title

Business Analytics and Data Science Group Project

submitted to

First Prof
Second Prof

Humboldt-Universität zu Berlin
School of Business and Economics

Chair

by

Claudia Guenther, Phi Nguyen, Julian Winkel
Immatriculation Numbers



Anything else we want to say

Berlin, Date

Abstract

Insert abstract here

List of Tables

1	Constructed features ordered by category (excerpt)	6
2	Accuracy comparison	13
3	Log-loss error comparison	13
4	Total cost comparison	13
5	Neural Network Hyperparameter Tuning	14
6	Random Forest Hyperparameter Tuning	14
7	Gradient Boost Hyperparameter Tuning	14
8	Cost Improvement from Ensembling	16

List of Figures

1	Threshold graphs for each classifier. Clockwise from top left: logistic regression, neural network, random forest, gradient boost	10
2	Reliability plots for each classifier. Clockwise from top left: logistic regression, neural network, random forest, gradient boost	12
3	Cross-Validation Accuracy	15
4	Cross-Validation Total Cost	16
5	Error cost and price ratios for both misclassification types (left); Ratio of total True Negatives and False Negatives per price bin (right).	17

Abbreviations

ANN Artificial Neural Network

LM Linear Model

Contents

1	Introduction	2
2	Previous Literature	3
3	Methodology	4
3.1	Heterogeneous ensemble models	4
3.2	Base classifier models	4
4	Data	5
5	Feature engineering	5
5.1	Feature construction	5
5.2	Feature selection	7
6	Model Building	8
6.1	Performance Measurement	8
6.2	Experimental Design	9
6.2.1	Nested Resampling	9
6.2.2	Model Calibration	11
6.3	Empirical Results	13
6.3.1	Ensembling	15
6.3.2	Post-hoc Cost-Sensitive Prediction Adjustment	16
7	Results	18
8	Conclusion	18
9	References	19

1 Introduction

2 Previous Literature

1 Page

Divide previous research in subsections that will be presented in the following.

This is how we cite Badea (2014). The reference is automatically pasted in the according section. You can also cite indirectly at the end of a sentence (Badea 2014). In this format, it is possible to insert pages, too (Badea 2014, 10–14).

3 Methodology

3.1 Heterogeneous ensemble models

The idea behind ensemble models (EM) is to combine multiple learning algorithms that taken together improve prediction accuracy over the performance of the individual models. The success of EM in predictive analytics is due to their ability of reducing both bias and variance by combining multiple model forecasts (Dietterich 2002). EM can be divided into homogeneous and heterogeneous approaches. Homogeneous EM use a single base-learning algorithm (e.g. decision tree) on distinct training sets. In contrast, heterogeneous EM train and combine learning algorithms of different types. In the context of classification problems, heterogeneous EM are also referred to as multi-classifier systems (MCS) and have gained increasing popularity (Ranawana and Palade 2006). Two key design features for the success of MCS is the diversity among the selected classifiers as well as the good individual performances (strength) of the baseline classifiers (Ali, Sremath Tirumala, and Sarrafzadeh 2015). However, it has been shown that neither diversity nor individual performance of baseline models can directly predict the ensemble success (Ranawana and Palade 2006).

MCS have been shown to be useful in predicting customer e-commerce behavior, outperforming single classifiers in prediction accuracy (Tsoumakas, Partalas, and Vlahavas (2008), Kim, Kim, and Lee (2003)). Specifically, Heilig et al. (2016) demonstrate the suitability of an ensemble selection approach (ES) for predicting product returns in online retailing. This three-step ensemble technique builds upon the creation of a library of heterogeneous candidate base models. Each classifier method can be used to build up multiple models with different meta-parameter settings. The ensemble model is created by combining base model from the library through an extensive search strategy. Though this approach has been demonstrated to be very effective for accurate predictions, it nevertheless requires extensive computational time. We therefore adopt a standard two-step ensemble approach. First, we develop and train four base classifiers models. Second, we combine these individual models with a simple non-linear combination method, namely majority voting. In case of a voting tie, the base model with the highest single accuracy will act as a tiebreaker.

3.2 Base classifier models

- baseline models can be seen in table 1
- in how far are they diverse? why did we select those?
- which one can capture linear, non-linear effects?

4 Data

The two data sets available to us contain a total of 150,000 order records from an online apparel retailer from a yearlong selling period. For 50,000 of these records, it is unknown whether an ordered item has been sent back by the customer or not. This second data set is the target of our binary predictions on customer’s returning behavior (return/not return). Both data sets include a total of 13 continuous/discrete and categorical variables. These covariates give information on customer demographics (e.g. user state, date of birth, title), order details (e.g. order date, delivery date), and item characteristics (e.g. item size, price or color). The majority of customers are middle-aged female shoppers, with a mean order volume of 70 Euro over the entire record period and an average product return rate of 48%.

To prepare the data sets for our analysis, we apply a set of standard pre-processing actions. Following the careful inspection of each variable, we remove all implausible values (e.g. extreme outliers). Approximately 20% of all records have missing values in the variables *delivery date* and/or *date of birth*. For better comprehensibility, we transform these covariates to *delivery time* and *age* respectively. Before imputing missing values, we flag them using dummy variables to extract their predictive power (Lessmann 2016, 18). Since *age* seems to be missing (completely) at random (MCAR) according to our data inspection, imputing it using mean substitution gives us an unbiased estimates (Schafer and Graham 2002)¹. Missing values in delivery time, caused by missing delivery dates, are clearly not missing not at random (MNAR) as they have a zero mean return rate and therefore are a perfect predictor. Possible reasons for these missing values are manifold. Without knowing the process generating these MNAR values, we cannot find unbiased substitutes for them (Schafer and Graham 2002, 171). We adopt three single substitution methods, namely case dropping, mean and median imputation, and chose the latter one based on model performance. We standardize the continuous variables only after the feature creation step to maintain their interpretability. Likewise, we do not directly drop zero (almost) zero variance predictors (e.g. date of birth) since we use them for feature extraction.

5 Feature engineering

5.1 Feature construction

Comprehensive feature engineering, i.e. creating and selecting useful and informative features, is the key to successful e-commerce prediction models (Liu et al. 2016). Based on the data structure available we separate our indicators into three e-commerce-related categories, namely product, customer and basket related features. Based on the data structure available we separate

¹Additionally, we carry out a Maximum Likelihood imputation of age in case the missing values are only missing at random (MAR), yielding the same model performance.

our indicators into three e-commerce-related categories, namely product, customer and basket related features as can be seen in table 2. Features in the product category are all item-specific indicators (e.g. item size, item price). Likewise, customer-related features contain variables with relational or demographic information on each customer (e.g. customer age, age of user account). We define the basket-related features as those that are common within one order basket (e.g. number of items within a basket). We consider all items belonging to one basket that were ordered by the same user on the same day.

Table 1: Constructed features ordered by category (excerpt)

Feature category	Feature	Type	Description
Product	Item price	Numeric	Price at which item is purchased
	Item category	Factor	Category (eg. clothing, shoes) of item
	Item WOE	Numeric	WOE based on item id
Customer	Customer Age	Numeric	Age in years at purchase
	Account Age	Numeric	Days since registration at purchase
	User’s income	Numeric	Proxy for costumers’s income
Order/Basket	Large basket	Factor	No. of items in same basket exceeds one
	Discount pc.	Numeric	Discount on item price in %
	Same category	Numeric	No. of items belonging to the same item category within same basket

In order to extract as much information as possible from the 13 original variables, we apply a variety of different extraction and projections methods, including unsupervised and supervised methods. Our extraction methods can be broadly summarized in four groups: discretization, projection, transformation/combination of variables and data augmentation by using external information. Based on these methods we construct an enriched data set with a total of 45 features. A small excerpt of these features can be seen in table 2. We apply discretization to several continuous variables, as this can be useful for capturing non-linear effects, constructing a more comprehensible representation of variables or improving predictive accuracy, especially in the context of tree-based decision algorithms (Liu et al. 2002). For example, we construct price item bins using equal frequency discretization as we suspect non-linearity regarding the

target variable. This approach is superior to equal width discretization for our purpose since it is outlier-resistant and gives us a higher interval resolution for item prices with a high density (Cichosz 2014, 531). A special challenge of this data set, which applies to many predictive modeling settings on customer behavior, is the presence of high-cardinality attributes (e.g. user id, item id). We replace each category with its Weight of evidence (WOE), a numerical projection method for categorical variables. The WOE is a measure of predictive power of a category with respect to the target variable, which has gained popularity in credit scoring (Lessmann 2016, 19–20). For a given customer, it shows a customers tendency in returning behavior, a valuable predictor variable. Replacing high-cardinality attributes with their respective WOE also reduces the dimensionality of our data set drastically. Since WOE has been shown to work well across the different types of our baseline models and requires low computational effort, we apply this projection methods to all categorical variables (Lessmann 2016, 32–33). For categories where WOE is not computable due to the existence of empty classes, we adapt WOE calculation by adding artificial data points as suggested by Kohavi and John (1997).

Another useful approach we take is the transformation and combination of variables based on domain knowledge about consumer’s shopping and returns behavior (Arnold and Reynolds (2003)). For example, we expect that discounted items will have a lower return rate due to lower item quality expectation as proposed by Petersen and Kumar (2009). Thus, we construct several discount variables (absolute discount, percentage discount, discount dummy) by taking the maximum item price for each item size as proxy for the regular item price and considering all deviations as a discount. Furthermore, a common defect of online shopping is that customers might be unsure about which size or color they should order based on the item description and thus place multiple orders of the same item (Foscht et al. 2013). We therefore construct informative variables on the order basket of consumers, e.g. the basket size in terms of items or the number of similar items with distinct sizes within one basket. Furthermore, we augment and enrich the data sets by including external sources to construct new features. For example, we use statistics on the mean income for different age groups and different regions within Germany. By combining these indicators with each customer’s demographical data (age and user state) we are able construct a proxy for each customer’s income. Moreover, we tackle the lack of information on item categories (e.g. pants, shoes, clothing, accessories). Based on size tables for different apparel categories we gathered we are able to categorize over 90% of items by finding sizes unique to a category. As expected, the mean return rate of shoes is almost twice as high as mean returns in accessories.

5.2 Feature selection

Finding an optimal feature subset is important for various reasons: First, learning algorithms like neural networks or tree-based models have been shown to degrade in performance when confronted with datasets containing redundant or irrelevant variables (Zdravevski et al. 2014).

Second, feature selection reduces data dimensionality, reducing the needed computational power for model building process and improving model comprehensibility. We adopt a hybrid feature selection strategy by combining filters and wrappers. As a first preprocessing step as proposed by Zdravevski et al. (2014), we remove very poor variables based on model agnostic information criteria. This has the advantage of reducing space dimensionality quickly and address overfitting (Guyon and Elisseeff 2003, 1170). For numerical variables we compute the F-Score as proposed by Chen and Lin (2006), whereas the average information value (obtained from WOE) is computed for categorical covariates.

The second step in our feature selection is a wrapper approach as described in Kohavi and John (1997). This Greedy selection strategy is advantageous since it has been shown to be robust against overfitting and is designed to optimize model performance (Guyon and Elisseeff 2003, 1167). We use a sequential floating backward selection (SFBS) algorithm as proposed by Pudil, Novovičová, and Kittler (1994). SFBS is able to capture interaction effects between variables better than non-floating approaches and has been shown to provide good results in medium-scale problems like this (Kudo and Sklansky 2000). Since the single baseline models from our ensemble model might perform best on distinct feature subsets, the optimal approach would be to conduct an individual wrapper for each model. Due to computational constraints we are nevertheless limited to conducting a single wrapper for all models. As feature subsetting is especially important for tree-based models and neural networks, we design a wrapper based on a simple random forest. First, we create a hold-out test set by creating a random stratified test and training data set. After standardization of both data sets, the WOE for all categorical variables is calculated on the training set and then passed on to the test set to prevent overfitting. Finally, we remove features by means of sequential floating backward selection based on maximizing AUC, which leaves us with a final subset of 17 features.

6 Model Building

6.1 Performance Measurement

To evaluate predictive accuracy, we first consider two performance measures - classification accuracy and logarithmic loss. Classification accuracy is the number of correctly classified observations divided by the total number of observations. We use this initial measure due to its ease of interpretability. However, this measure does not differentiate between types of error (which we is important due to asymmetric costs) nor does it penalize overconfidence in classification. Therefore we look additionally at the logarithmic loss, a form of cross-entropy that is used when the desired outcome is the probability that an example is positive. The formula for logarithmic loss is shown below, where N is the total number of observations, y_i is the true outcome, and p_i is our predicted probability outcome. This measure is designed to penalize

overly confident but misclassified observations and is often preferred for binary classification problems like ours (Caruana and Niculescu-Mizil 2004).

$$\text{logloss} = -\frac{1}{N} \sum_{i=1}^N [y_i \log(p_i) + (1 - y_i) \log(1 - p_i)]$$

Finally, we consider predictive accuracy within the framework of our specific prediction task, which is to minimize the cost associated with incorrect classifications. Specifically, our problem deals with example-dependent misclassification costs because the penalty for false classification is associated with the item’s corresponding cost due to restocking fees or return fees. Namely, in the event of false positive classification (we incorrectly predict the customer will return), there will be an opportunity cost of a lost potential purchase of $-0.5 * \text{itemprice}$. In the event of a false negative classification (we incorrectly predict the customer will not return), we incur an actual financial penalty of $-0.5 * 5 * (3 + 0.1 * \text{itemprice})$. For our purposes, we consider both actual and opportunity costs as impacting our cost calculations.

To solve this problem, we use a thresholding approach. The thresholding approach selects the probability cutoff point that minimizes the total misclassification cost from the training data set and uses it to predict the class label of the test instances. Test instances that have a predicted probability below or equal to this threshold are classified as a negative instance, whereas test instances with a probability higher than the threshold are classified as a positive (i.e. return) instance. Despite its simplicity, thresholding often empirically outperforms other cost-sensitive learning methods (Sheng and Ling 2006).

Our specific approach mirrors a similar approach as suggested by Fawcett (2006), in which he generates a ROC curve that is sensitive to instance-varying costs. In our approach, we loop through a sequence of possible threshold values between 0 and 1, plot the corresponding accuracy and cost, then select the threshold that minimizes the overall cost. Since costs are shown as negative, we seek to maximize the value. We plot these graphs showing the accuracy and costs side by side given different threshold points for each of the four classifiers in [FIGURE XXX]. Notice how the threshold that optimizes costs is different from the one that maximizes accuracy.

6.2 Experimental Design

6.2.1 Nested Resampling

In order to obtain reliable out-of-sample performance estimates for each classifier and the overall ensemble model, we undertake a nested resampling approach. This approach employs two cross-validation loops. We use the inner loop to tune the hyperparameters, while we use the outer loop to estimate how the model will generalize to a new, unknown data set by averaging the test results (Hastie, Tibshirani, and Friedman 2001, 241).

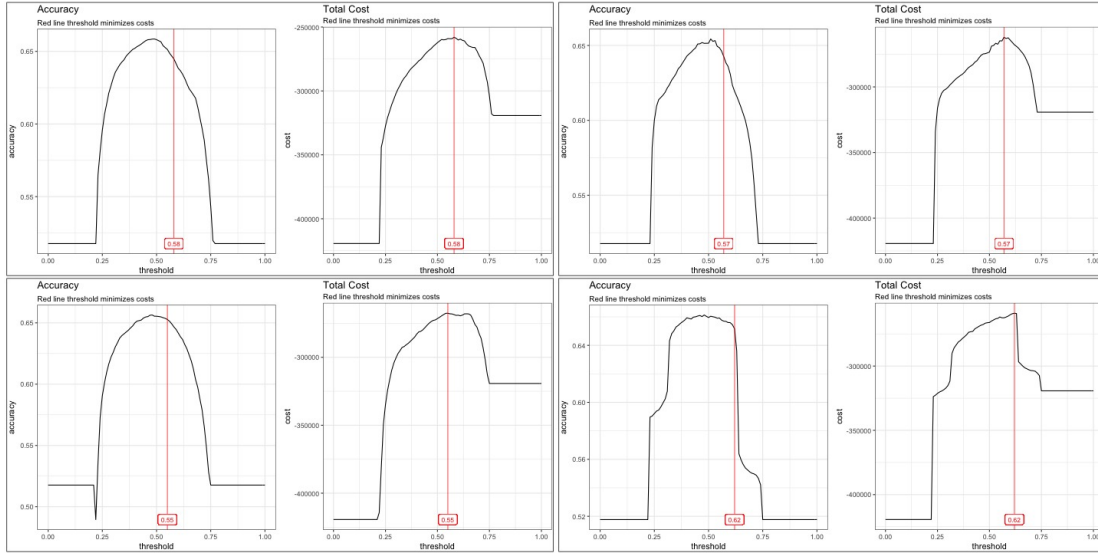


Figure 1: Threshold graphs for each classifier. Clockwise from top left: logistic regression, neural network, random forest, gradient boost

In the outer loop, we use 4-fold cross-validation. In this approach, we partition the initial test set containing 100,000 observations into four equally sized groups. For each iteration of the loop, we set aside the i -th partition for testing, while we use the remaining 4 partitions to train the individual classifiers (logistic regression, neural network, gradient boosting, random forest).

In the inner resampling loop, we train each classifier using a fixed holdout instance to find the optimal hyperparameters on each of the individual classifiers. We use a fixed holdout set to enable us to simply standardize numeric variables and create our Weight of Evidence features (by calculating off the training set and applying those values to the holdout set) within the inner loop. This preprocessing step is necessary for all variable transformations that are dependent on the data to prevent overfitting the model.

The inner resampling loop assesses the quality of the model given the chosen hyperparameters. With the consideration that a grid search across the entire hyperparameter search space would be too costly, we choose to instead do a random search across the search space. The random search method will randomly choose from our specified values, with a maximum of 20 iterations. Additionally, we train each classifier to optimize for AUC (area under the ROC curve), since such classifiers tend to produce better accuracy relative to those that only optimize for accuracy (Huang and Ling 2005). The chosen hyperparameters are then passed to the outer loop where a new model is trained with the corresponding hyperparameters. We then predict the results on the outer loop holdout set and aggregate results.

6.2.2 Model Calibration

Additionally, in the outer loop, we calibrate the probability estimates generated from each classifier. We use calibration for two reasons: Firstly, it ensures that the probability estimates better reflect their true posterior probabilities (Caruana, Niculescu-Mizil, and Munson 2006). For example, in a well-calibrated model, if the probability estimate for a given instance is 70%, that means we can assume that that event has a 70% chance of occurring. Secondly, calibration transforms predictions from all models to place on them equivalent scale. Combining calibrated models in ensembling not only provides improvement in probability measures like logarithmic loss, but failure to use calibrated models may hurt ensemble performance (Caruana and Niculescu-Mizil 2005b).

We use a method known as Platt scaling to obtain calibrated predictions. In Platt scaling, we pass the output probabilities through a sigmoid (which in our case is a logistic regression) (Platt 1999). We apply the specific method below as described in (Caruana and Niculescu-Mizil 2005a) in order to prevent overfitting and unwanted bias:

1. Split the data set into a training and a test set
2. Split the training data set further into a calibration set
3. Train the model on the training set
4. Predict the model on the calibration set and the test set
5. Use the predictions from calibration set to train a logistic regression model
6. Pass the predictions from the test set through the logistic regression model in step 5 to obtain the final predictions

We visualize the effect of calibration on each of our classifiers using reliability diagrams from the first fold of cross-validation. Reliability diagrams split the prediction space into 10 bins, then plot the mean prediction value for each bin against the true fraction of positive cases. The diagonal line indicates a perfectly calibrated model (Caruana and Niculescu-Mizil 2005b). Additionally, in each diagram, we plot the histogram of predicted values before and after calibration. The results of calibration for each classifier can be seen in [FIGURE xxxx].

From the reliability plots, we can see a few major effects of calibration. Firstly the results after calibration (the line in blue) shift away from a sigmoidal shape (the line in red) to to line up closer to the diagonal line indicating a perfectly calibrated model. Secondly, calibration pushes the probabilities closer to the center and away from 0 and 1, to better align with the true fraction of positive cases and prevent overconfidence in classifications. One benefit that we can recognize from the reliability plots is that overall, each histogram post-calibration appear similar to one another because calibration reduces the differences between predicted probabilities for each model, which will provide usefulness in ensembling (Caruana and Niculescu-Mizil 2005b).

Overall, we see benefits in calibration when comparing the accuracy, log-loss, and cost metrics. In [FIGURE xxx], we compare the average across all validation runs for the aforementioned




Figure 2: Reliability plots for each classifier. Clockwise from top left: logistic regression, neural network, random forest, gradient boost

metrics. We see that calibration yields significant improvements in log-loss error for all models without negatively impacting accuracy and cost measurements.

Table 2: Accuracy comparison

Classifier	Uncalibrated	Calibrated
Logistic Regression	0.657	0.657
Neural Network	0.662	0.661
Random Forest	0.656	0.656
Gradient Boost	0.653	0.654

Table 3: Log-loss error comparison

Classifier	Uncalibrated	Calibrated
Logistic Regression	0.709	0.618
Neural Network	0.774	0.621
Random Forest	0.774	0.625
Gradient Boost	0.817	0.628

Table 4: Total cost comparison

Classifier	Uncalibrated	Calibrated
Logistic Regression	-257376.4	-257644.0
Neural Network	-257108.2	-262476.1
Random Forest	-262464.2	-257532.9
Gradient Boost	-258767.0	-258816.6

6.3 Empirical Results

The output of the nested resampling results in a total of five candidate models and sets of predictions for each of the four classifiers. We use the calibrated results of these classifiers for the following:

1. Assess the stability of the hyperparameters
2. Estimate the prediction accuracy and total cost of the model
3. Get the average optimal decision threshold to apply to our final classifiers
4. Use the model with the lowest average cost as a tiebreaker in our ensemble

We can examine the stability of our hyperparameters in [TABLES x,y,z]. Since the hyperparameters do not vary wildly between iterations, we can make the assumption the inner loop optimization works as designed. The distribution of performance values (in terms of accuracy and total cost, respectively) across resampling iterations and for all classifiers is shown in [FIGURES x,y]. The dot in the center represents the average value, whereas the spread is indicated by 1 standard error above and below the mean. On average, the neural network consistently performs best both on prediction accuracy and total cost metrics. In comparison, gradient boosting tends to perform the worst in terms of prediction accuracy but random forests perform the worst in terms of total cost minimization. Logistic regression, as expected, shows the smallest variation between iterations. However, although neural networks and logistic regression tends to vary little between iterations, gradient boosting and random forest approaches have far larger variations between iterations in terms of total cost minimization. The best performance of any one run came from a gradient boost iteration.

Table 5: Neural Network Hyperparameter Tuning

Iteration	size	decay
1	4	1e-03
2	5	1e-05
3	8	1e-02
4	6	1e-03

Table 6: Random Forest Hyperparameter Tuning

Iteration	ntree	mtry	nodesize
1	300	3	15
2	350	3	30
3	250	3	25
4	400	3	15

Table 7: Gradient Boost Hyperparameter Tuning

Iteration	gamma	eta	nrounds	lambda	max_depth
1	4	1e-03	100	0.3	3
2	5	1e-05	200	0.1	3
3	8	1e-02	100	0.4	3
4	6	1e-03	100	0	4

- Note from Claudia : Cite paper by Prof Lessman Coussement, Lessmann, and Verstraeten (2017): We explain good performance of log reg with existence of linear effects and our meticulous data preparation. As Coussement, Lessmann, and Verstraeten (2017) have demonstrated, logistic regression can compete with more complex algorithms if exposed to well prepared data set




Figure 3: Cross-Validation Accuracy

We note that nested resampling is a computationally intensive process. We limit the inner loop cross-validation process to 4 iterations for this reason. Additionally, we employ relatively small search spaces for hyperparameters (usually not exceeding 30 possible combinations) that are informed by theoretical considerations (citation???). In order to accelerate the overall model fitting process, we employ a parallelization wrapper around the tuning process.

6.3.1 Ensembling

Finally, we ensemble models from each of the k runs and compare the final combined result to the average of the individual candidate classifiers. Our ensemble approach uses majority voting as described in the Methodology section of the paper. Given that we are ensembling four models, we use the best model in terms of cost minimization as a tie breaker. In our case, our best performing model is the neural network.

For each of the 4 iterations of the cross-validation process, we use our ensembling method to




Figure 4: Cross-Validation Total Cost

assign final predictions for each instance. We then calculate the total cost of misclassification and compare that to the average cost for each individual classifier. The results can be seen in [TABLE XXXX].

Table 8: Cost Improvement from Ensembling

Iteration	Average Classifier	Ensembled	% Improvement
1	-261651.96	-256435.62	0.0199
2	-256735.30	-251296.48	0.0212
3	-258906.22	-255041.46	0.0149
4	-259175.97	-257744.25	0.0055

The results of ensembling show an average improvement of 1.5% compared to the average base classifier. Although this figure may seem small, it can result in significant cost savings for the retailer over the long run. Additionally, we note that in some cases the ensemble performs slightly worse than the best individual classifier. However, we believe this approach remains valid because of how high the variation is between individual runs, depending on the observations in the data set. The ensembling process helps reduce this variation by averaging the performance of each model.

6.3.2 Post-hoc Cost-Sensitive Prediction Adjustment

The final class memberships of our model predictions are the results of our cost-minimizing threshold strategy. Despite the suitability of threshold tuning in general, the outcome might not be optimal due to the presence of cost asymmetry. This holds especially true for the items with a low item price, where disparity between the distinct classification errors is the largest. This can

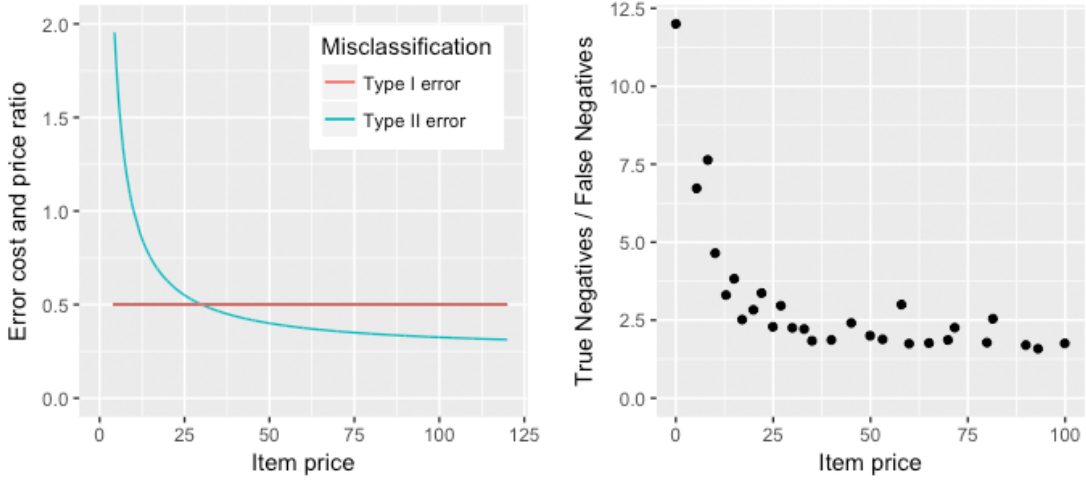


Figure 5: Error cost and price ratios for both misclassification types (left); Ratio of total True Negatives and False Negatives per price bin (right).

it must hold that $0 < \frac{C_2}{C_1} - \frac{TN}{FN}$ where C_1 and C_2 refer to the misclassifications errors of type 1 and 2 at a given item price. Since some item price levels have sparse class membership, we create equal frequency price bins and approximate C_1 and C_2 with mean errors costs. We test our strategy based on the baseline prediction of our logistic regression model on a stratified test set. In our sample, the pure return strategy only saves costs for a item price of zero. This is because the high spread between costs is outweighed by the high ratio between TN and FN. This can be seen in figure 1 (right), where this ratio is displayed for all price bins. As very few cheap items are returned by customers, the number of FN is small in low price bins, leading to a high ratio.

7 Results

8 Conclusion

9 References

- Ali, Shahid, Sreenivas Sremath Tirumala, and Abdolhossein Sarrafzadeh. 2015. “Ensemble Learning Methods for Decision Making: Status and Future Prospects.”
- Arnold, Mark J, and Kristy E Reynolds. 2003. “Hedonic Shopping Motivations.” *Journal of Retailing* 79 (2). Elsevier: 77–95.
- Badea, Laura Maria. 2014. “Predicting Consumer Behavior with Artificial Neural Networks.” *Procedia Economics and Finance* 15: 238–46. doi:[https://doi.org/10.1016/S2212-5671\(14\)00492-4](https://doi.org/10.1016/S2212-5671(14)00492-4).
- Caruana, Rich, and Alexandru Niculescu-Mizil. 2004. “Data Mining in Metric Space: An Empirical Analysis of Supervised Learning Performance Criteria.” In, 69–78. ACM Press.
- . 2005a. “Obtaining Calibrated Probabilities from Boosting.” In *In: Proc. 21st Conference on Uncertainty in Artificial Intelligence (Uai '05)*, Auai Press. AUAI Press.
- . 2005b. “Predicting Good Probabilities with Supervised Learning.” In *Proceedings of the 22Nd International Conference on Machine Learning*, 625–32. ICML '05. New York, NY, USA: ACM.
- Caruana, Rich, Alexandru Niculescu-Mizil, and A. Munson. 2006. “Getting the Most Out of Ensemble Selection.” In *Sixth International Conference on Data Mining (Icdm'06)*, 828–33.
- Chen, Yi-Wei, and Chih-Jen Lin. 2006. “Combining Svms with Various Feature Selection Strategies.” In *Feature Extraction*, 315–24. Springer.
- Cichosz, Pawel. 2014. *Data Mining Algorithms: Explained Using R*. John Wiley & Sons.
- Coussement, Kristof, Stefan Lessmann, and Geert Verstraeten. 2017. “A Comparative Analysis of Data Preparation Algorithms for Customer Churn Prediction: A Case Study in the Telecommunication Industry.” *Decision Support Systems* 95. Elsevier: 27–36.
- Dietterich, Thomas G. 2002. “Ensemble Learning.” *The Handbook of Brain Theory and Neural Networks* 2. MIT Press: Cambridge, MA: 110–25.
- Fawcett, Tom. 2006. “ROC Graphs with Instance-Varying Costs.” *Pattern Recogn. Lett.* 27 (8). New York, NY, USA: Elsevier Science Inc.: 882–91.
- Foscht, Thomas, Karin Ernstreiter, Cesar Maloles III, Indrajit Sinha, and Bernhard Swoboda. 2013. “Retaining or Returning? Some Insights for a Better Understanding of Return Behaviour.” *International Journal of Retail & Distribution Management* 41 (2). Emerald Group Publishing Limited: 113–34.
- Guyon, Isabelle, and André Elisseeff. 2003. “An Introduction to Variable and Feature Selection.”

Journal of Machine Learning Research 3 (Mar): 1157–82.

Hastie, Trevor, Robert Tibshirani, and Jerome Friedman. 2001. *The Elements of Statistical Learning*. Springer Series in Statistics. New York, NY, USA: Springer New York Inc.

Heilig, Leonard, Julien Hofer, Stefan Lessmann, and Stefan Voß. 2016. “Data-Driven Product Returns Prediction: A Cloud-Based Ensemble Selection Approach.” In *ECIS*, Research-in.

Huang, Jin, and Charles X. Ling. 2005. “Using Auc and Accuracy in Evaluating Learning Algorithms.” *IEEE Transactions on Knowledge and Data Engineering* 17: 299–310.

Kim, Eunju, Wooju Kim, and Yillbyung Lee. 2003. “Combination of Multiple Classifiers for the Customer’s Purchase Behavior Prediction.” *Decision Support Systems* 34 (2). Elsevier: 167–75.

Kohavi, Ron, and George H John. 1997. “Wrappers for Feature Subset Selection.” *Artificial Intelligence* 97 (1-2). Elsevier: 273–324.

Kudo, Mineichi, and Jack Sklansky. 2000. “Comparison of Algorithms That Select Features for Pattern Classifiers.” *Pattern Recognition* 33 (1). Elsevier: 25–41.

Lessmann, Stefan. 2016. “Data Preprocessing in Database Marketing: Tasks, Techniques, and Why.” In *Advanced Database Marketing: Innovative Methodologies and Applications for Managing Customer Relationships*. Routledge.

Liu, Guimei, Tam T Nguyen, Gang Zhao, Wei Zha, Jianbo Yang, Jianneng Cao, Min Wu, Peilin Zhao, and Wei Chen. 2016. “Repeat Buyer Prediction for E-Commerce.” In *Proceedings of the 22nd Acm Sigkdd International Conference on Knowledge Discovery and Data Mining*, 155–64. ACM.

Liu, Huan, Farhad Hussain, Chew Lim Tan, and Manoranjan Dash. 2002. “Discretization: An Enabling Technique.” *Data Mining and Knowledge Discovery* 6 (4). Springer: 393–423.

Petersen, J Andrew, and V Kumar. 2009. “Are Product Returns a Necessary Evil? Antecedents and Consequences.” *Journal of Marketing* 73 (3). American Marketing Association: 35–51.

Platt, John C. 1999. “Probabilistic Outputs for Support Vector Machines and Comparisons to Regularized Likelihood Methods.” In *ADVANCES in Large Margin Classifiers*, 61–74. MIT Press.

Pudil, Pavel, Jana Novovičová, and Josef Kittler. 1994. “Floating Search Methods in Feature Selection.” *Pattern Recognition Letters* 15 (11). Elsevier: 1119–25.

Ranawana, Romesh, and Vasile Palade. 2006. “Multi-Classifer Systems: Review and a Roadmap for Developers.” *International Journal of Hybrid Intelligent Systems* 3 (1). IOS Press: 35–61.

Schafer, Joseph L, and John W Graham. 2002. “Missing Data: Our View of the State of the Art.” *Psychological Methods* 7 (2). American Psychological Association: 147.

Sheng, Victor S., and Charles X. Ling. 2006. “Thresholding for Making Classifiers Cost-Sensitive.”

In *AAAI*, 476–81. AAAI Press.

Tsoumakas, Grigorios, Ioannis Partalas, and Ioannis Vlahavas. 2008. “A Taxonomy and Short Review of Ensemble Selection.” In *Workshop on Supervised and Unsupervised Ensemble Methods and Their Applications*.

Zdravevski, Eftim, Petre Lameski, Andrea Kulakov, and Dejan Gjorgjevikj. 2014. “Feature Selection and Allocation to Diverse Subsets for Multi-Label Learning Problems with Large Datasets.” In *Computer Science and Information Systems (Fedcsis), 2014 Federated Conference on*, 387–94. IEEE.

Declaration of Authorship

TEXT

15.01.2018